

Graphs-2

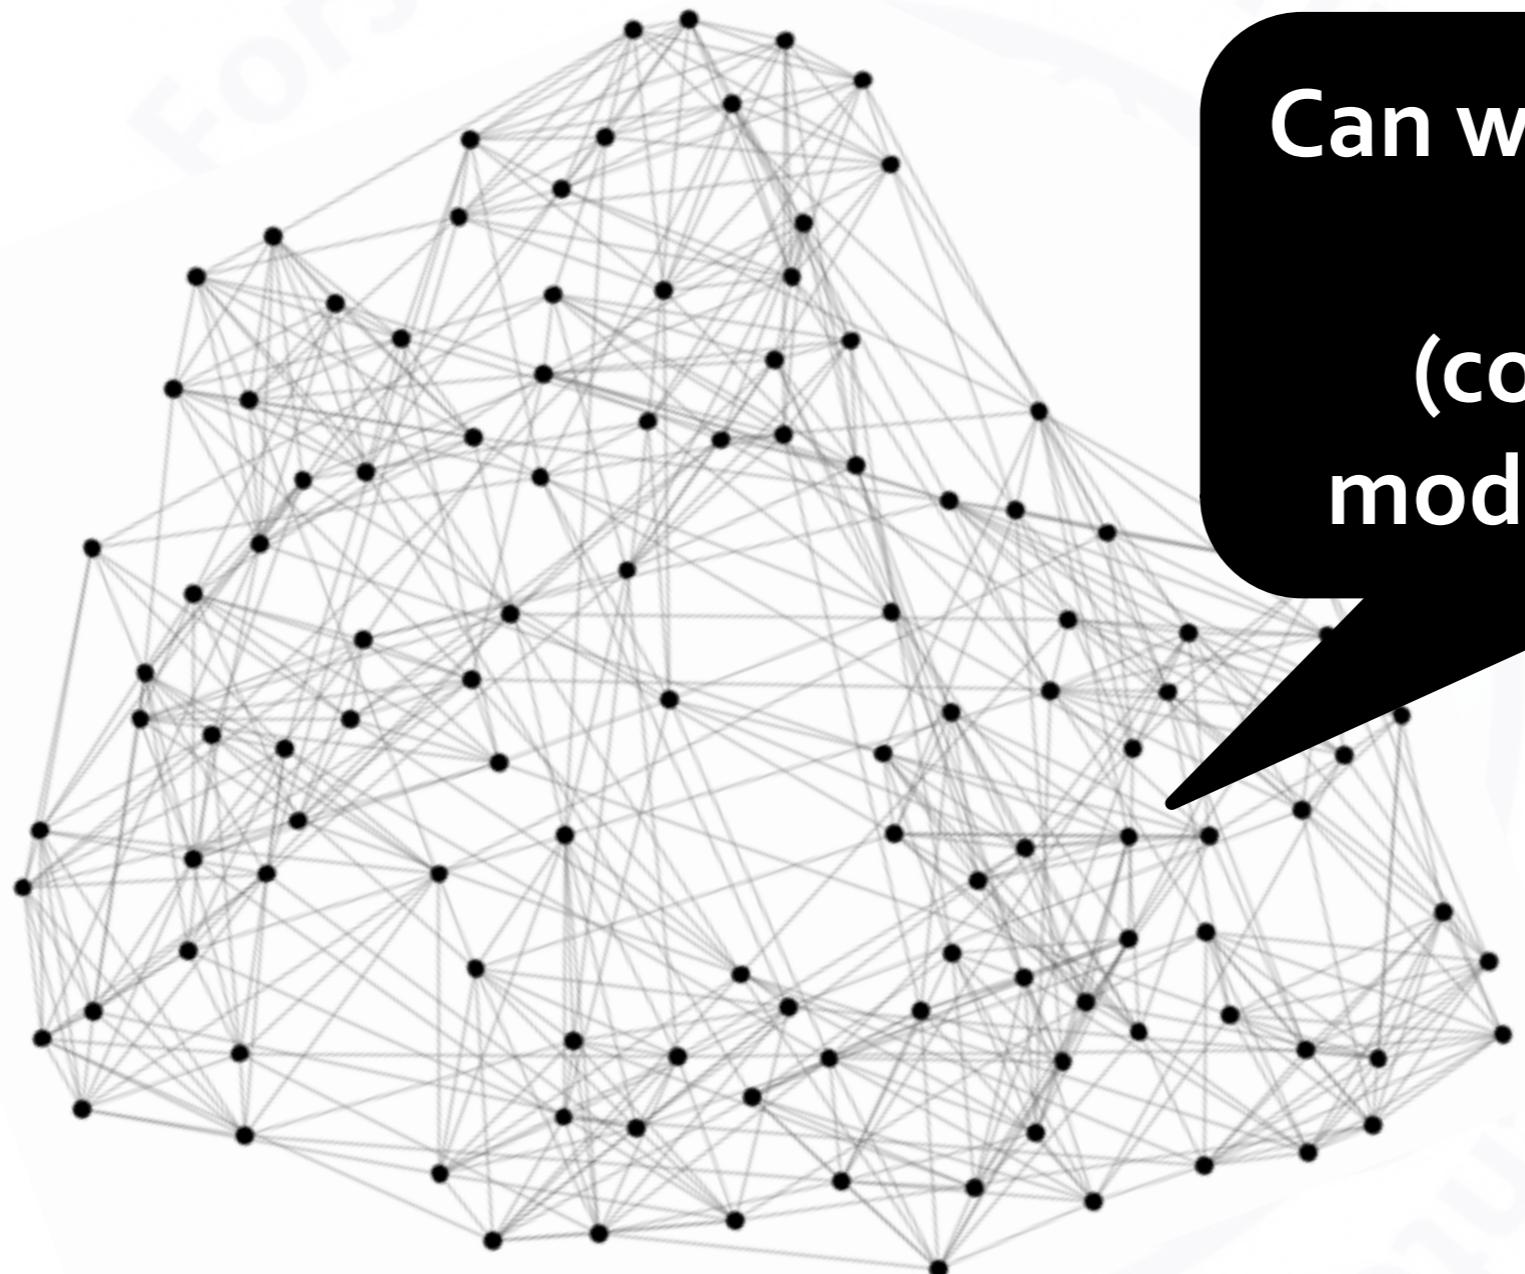
AGM, MLE on Graphs, Counting triangles

Identifying Communities



**Nodes: Football Teams
Edges: Games played**

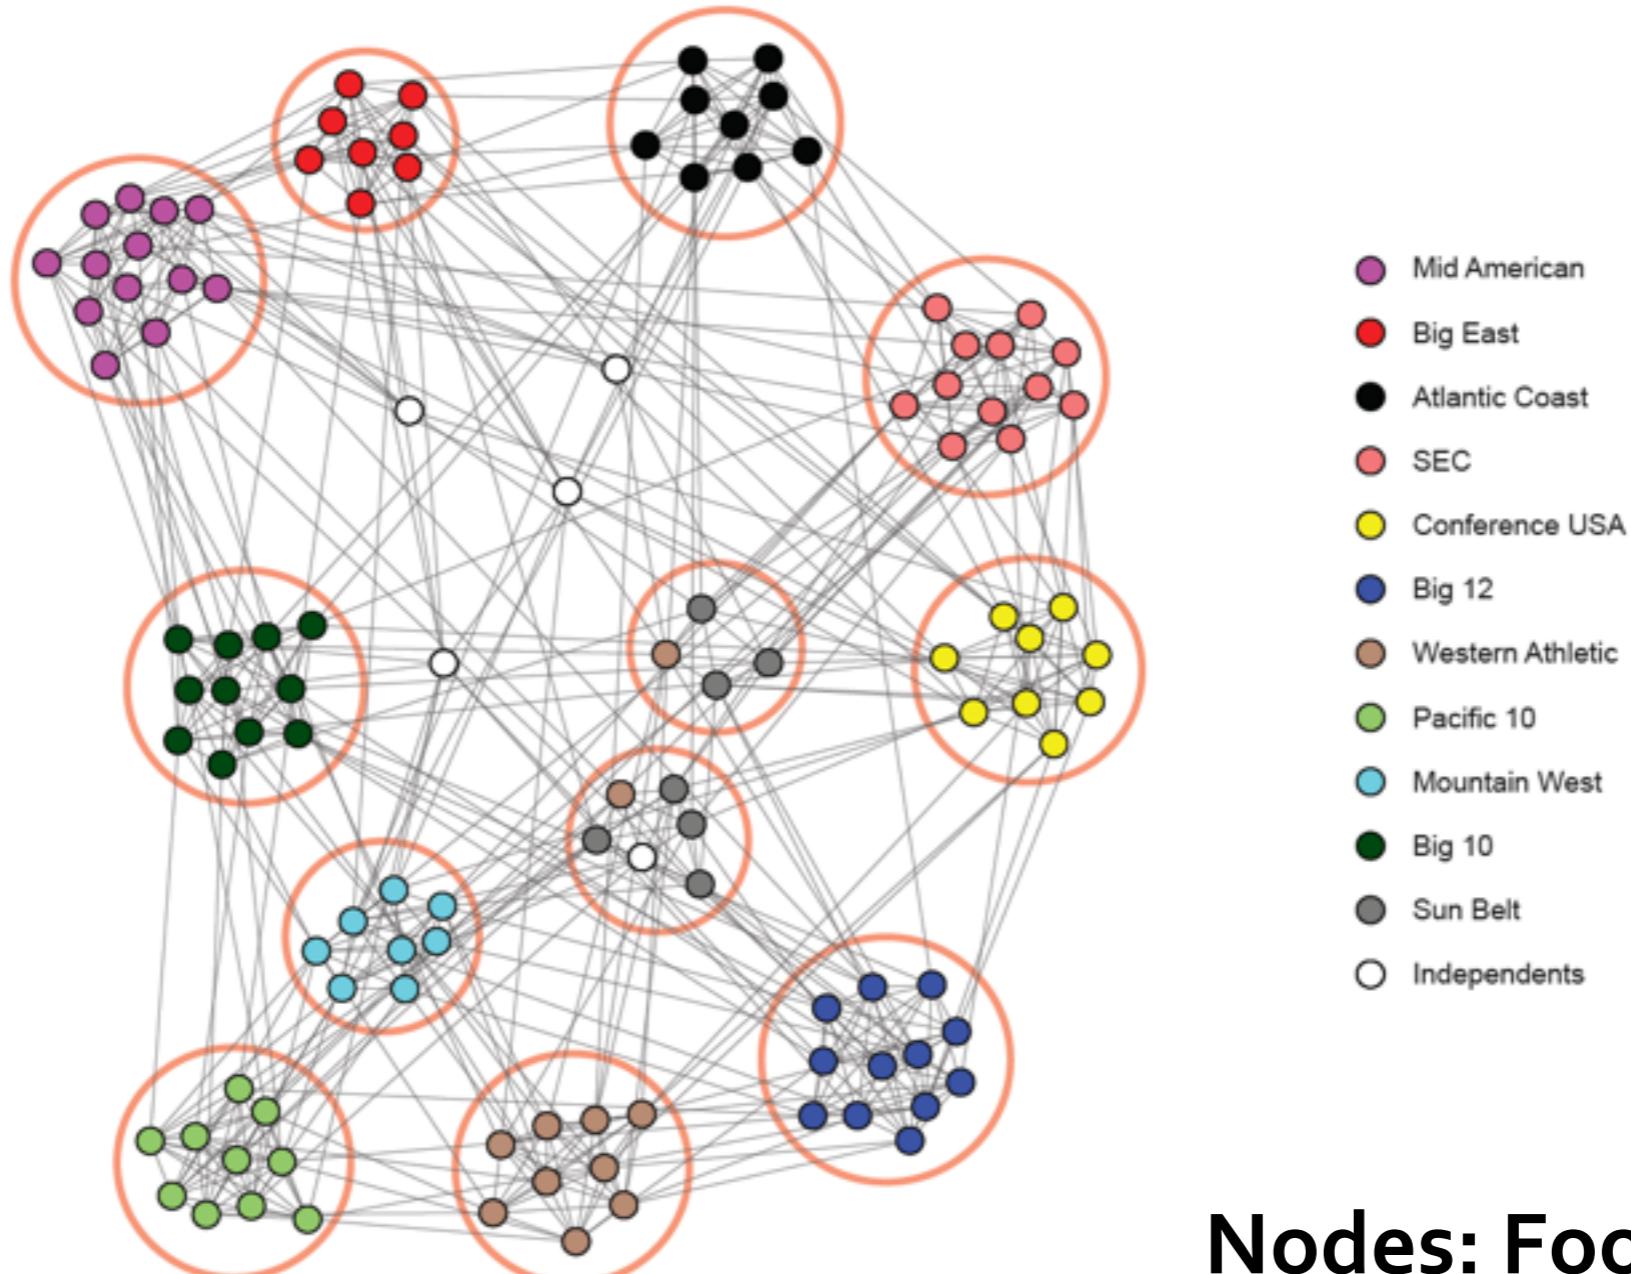
Identifying Communities



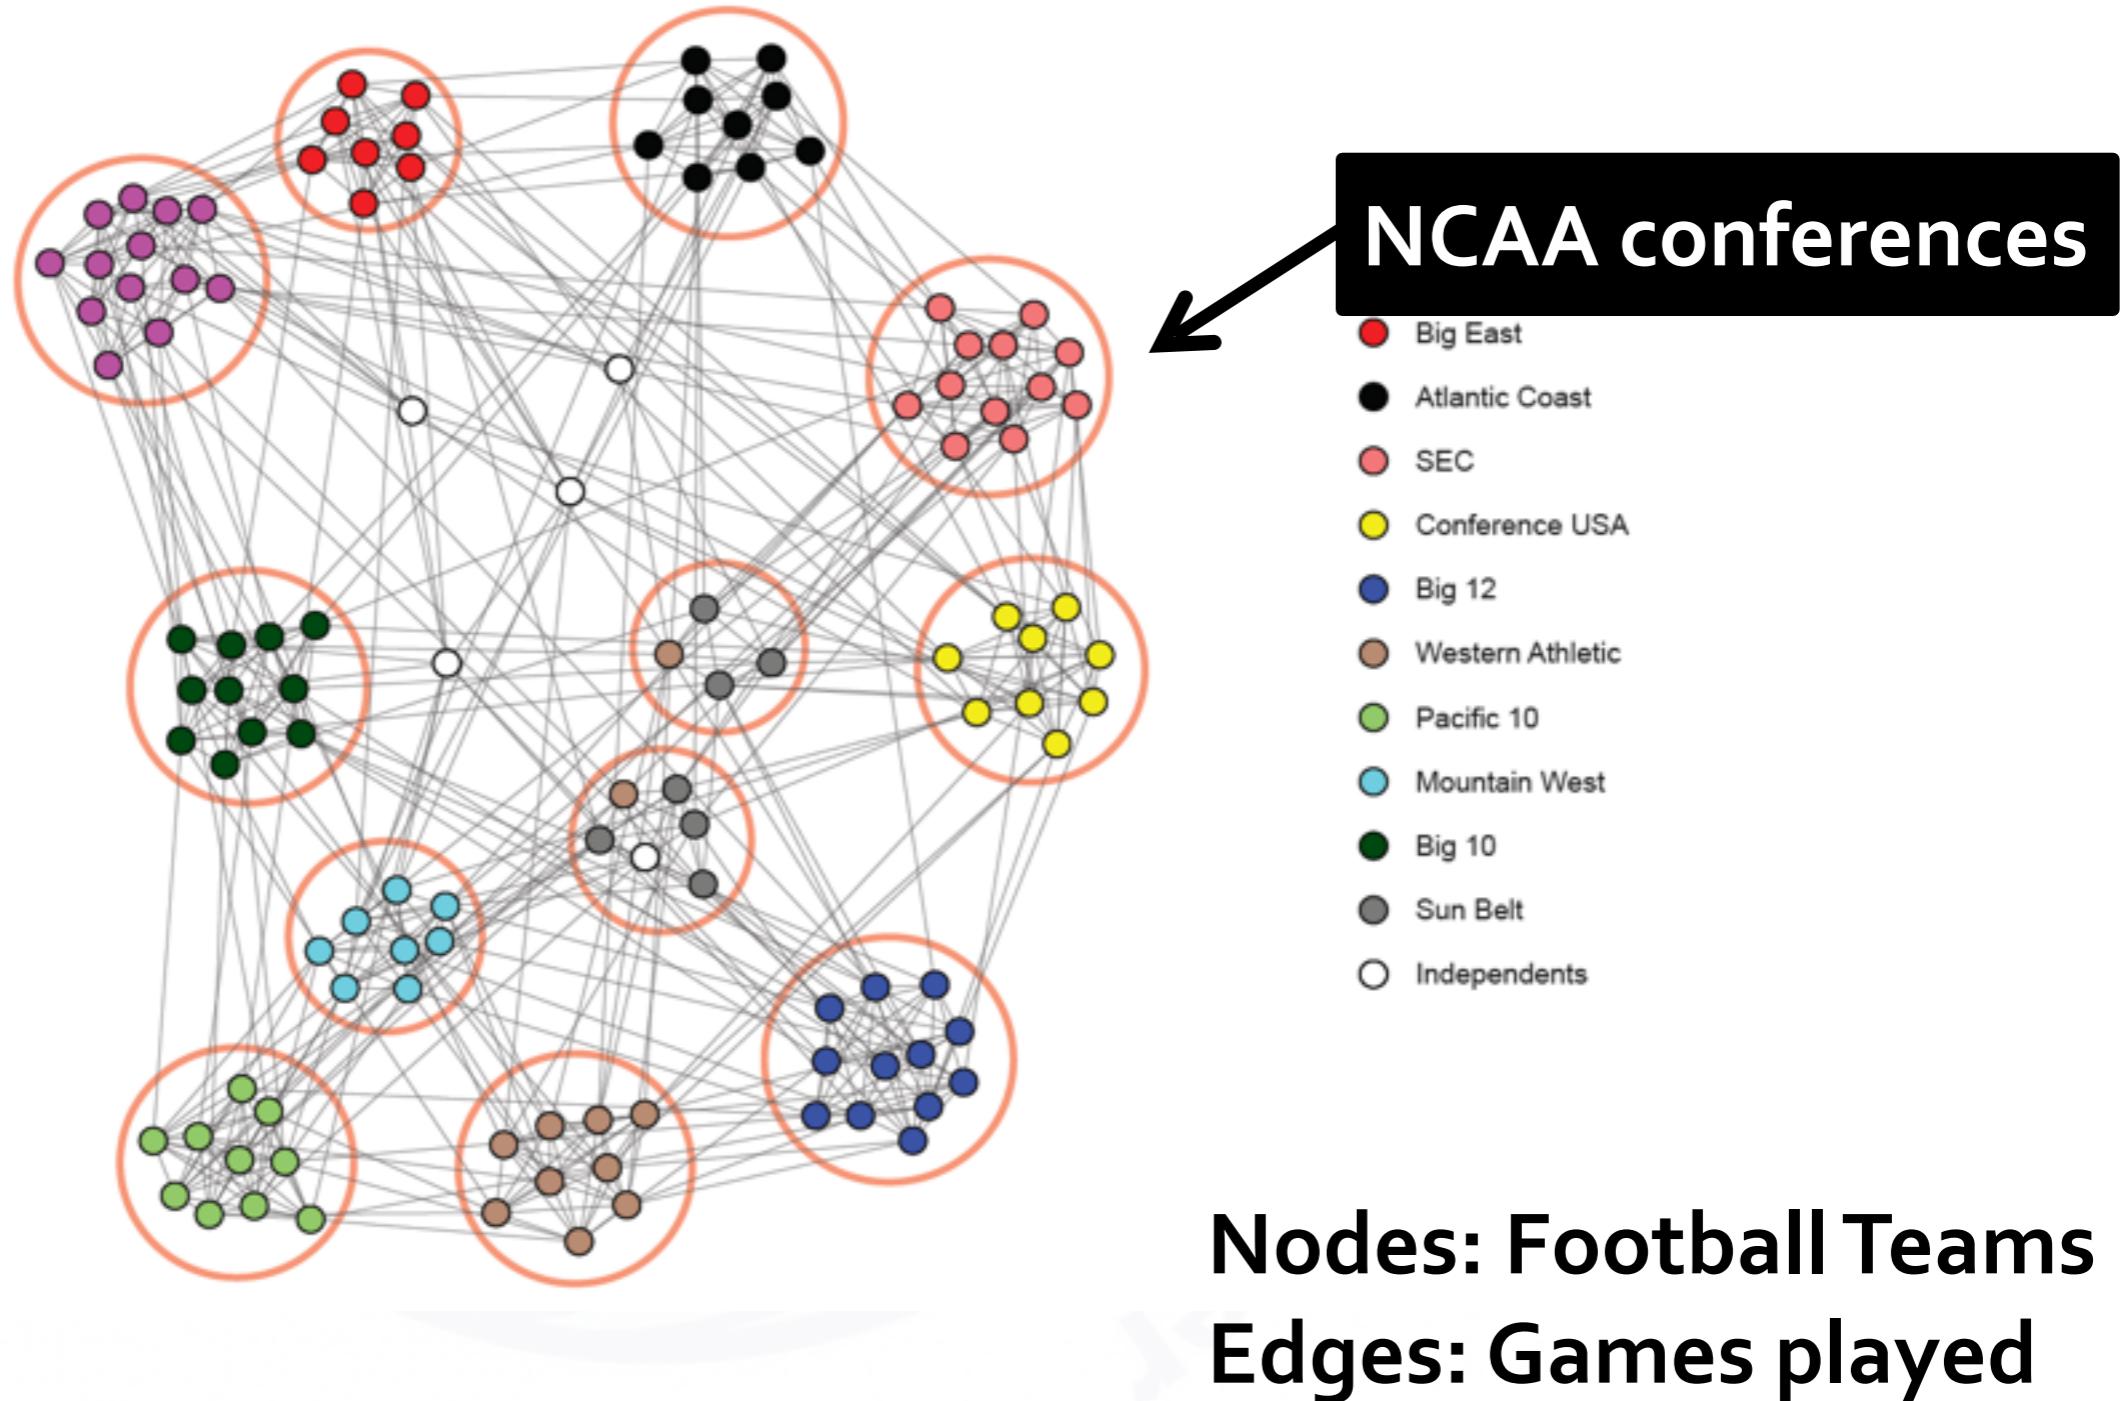
Can we identify node groups?
(communities, modules, clusters)

Nodes: Football Teams
Edges: Games played

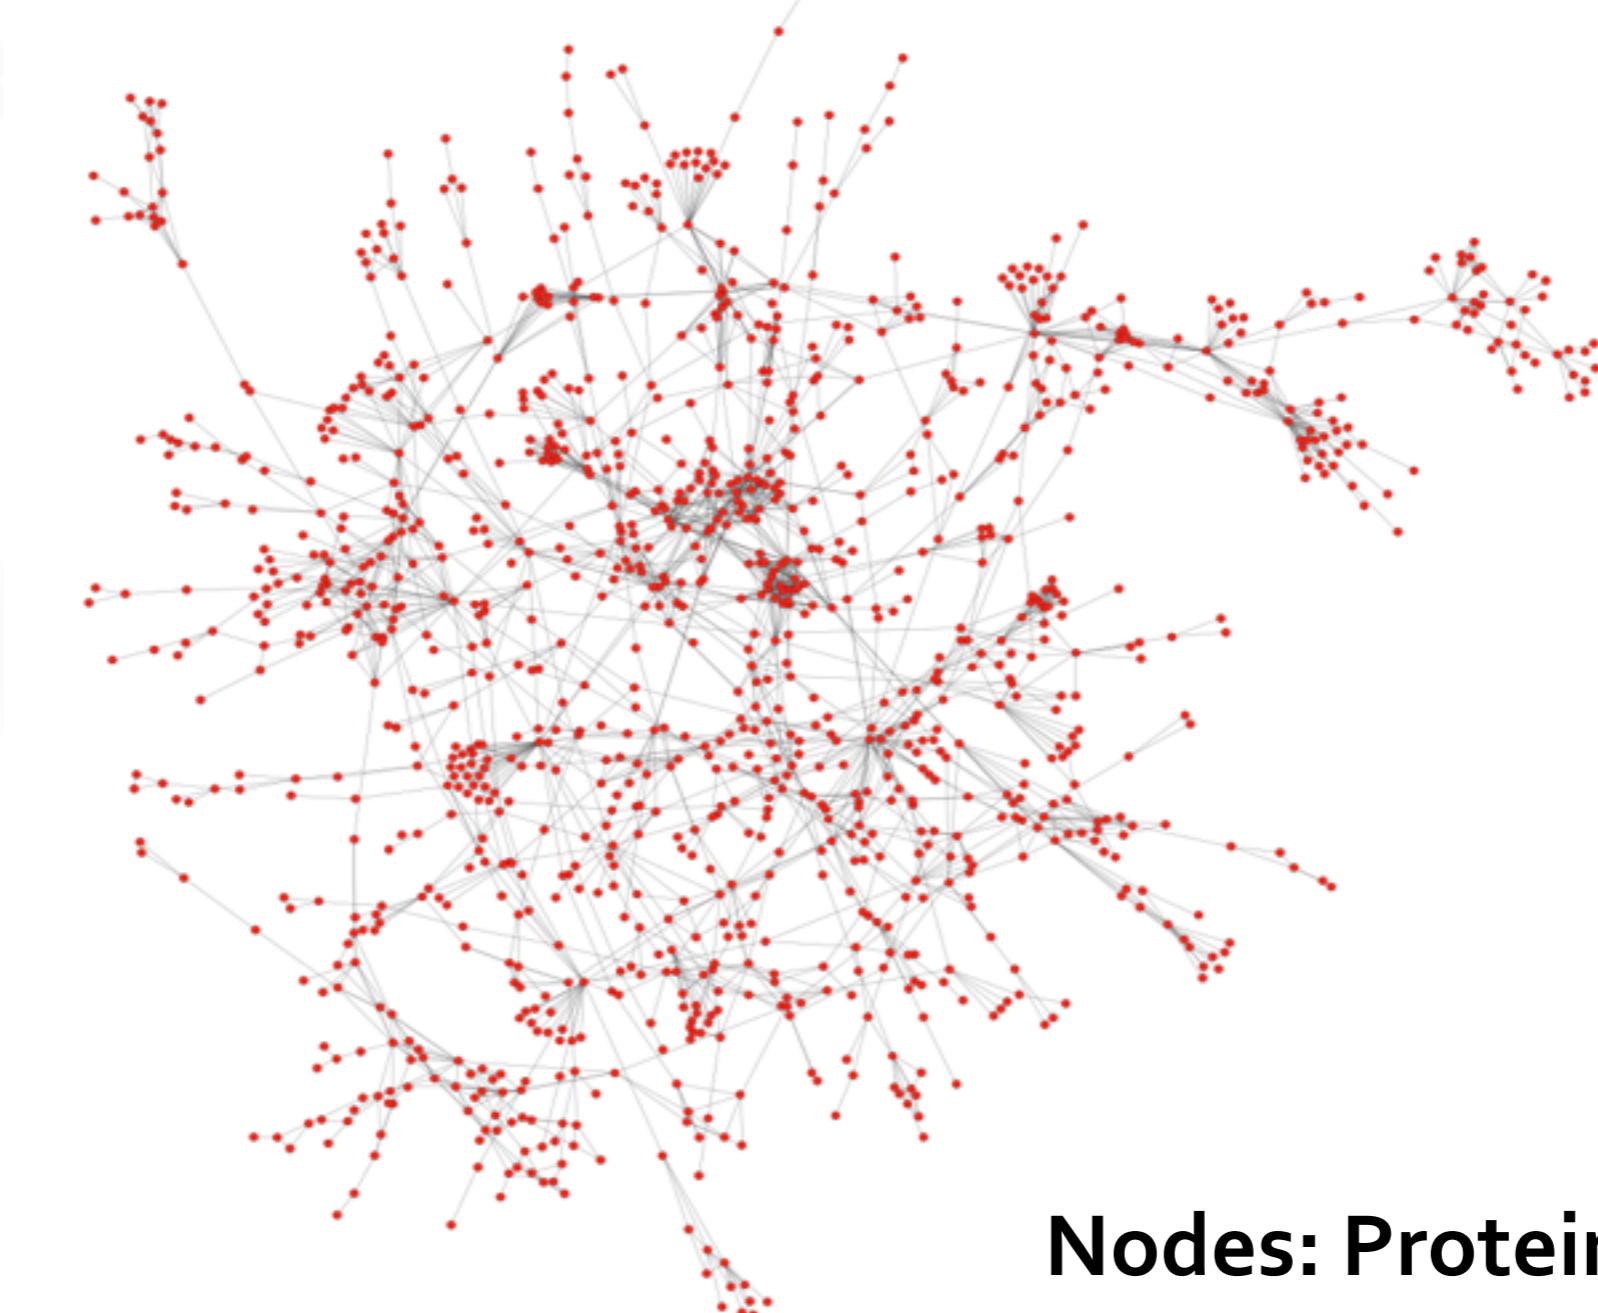
NCAA Football Network



NCAA Football Network

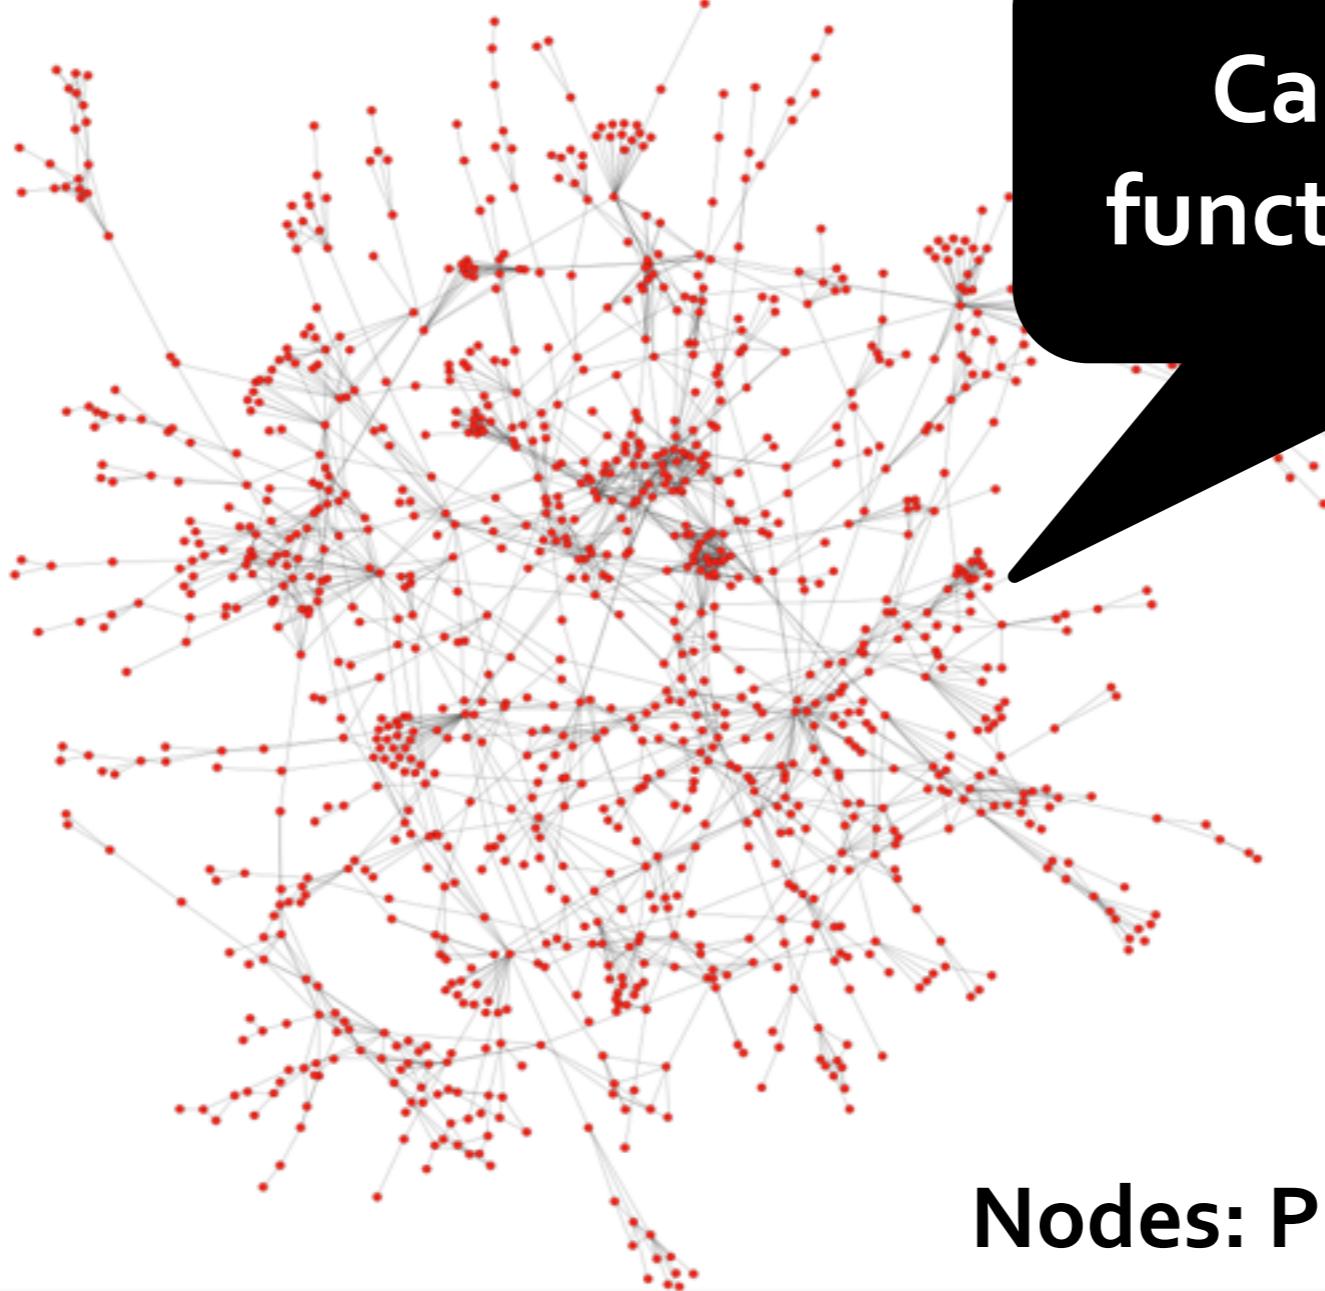


Protein-Protein Interactions



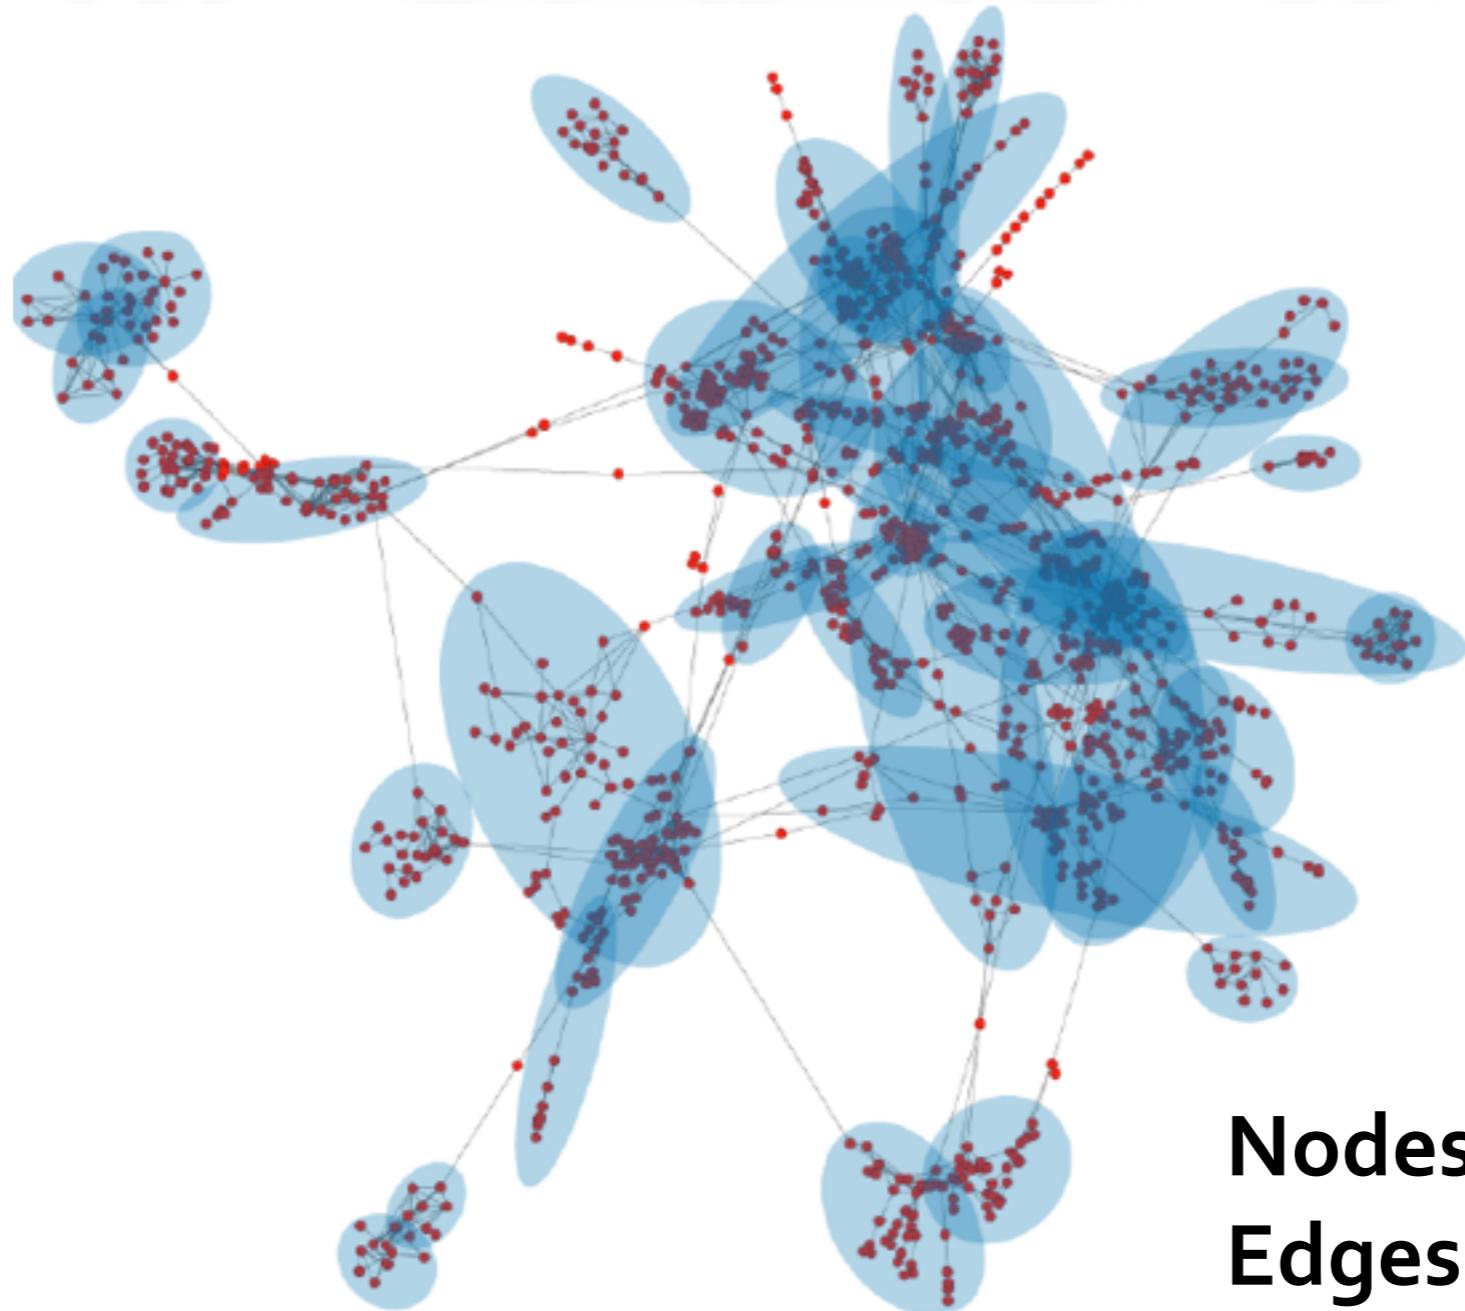
Protein-Protein Interactions

Can we identify
functional modules?



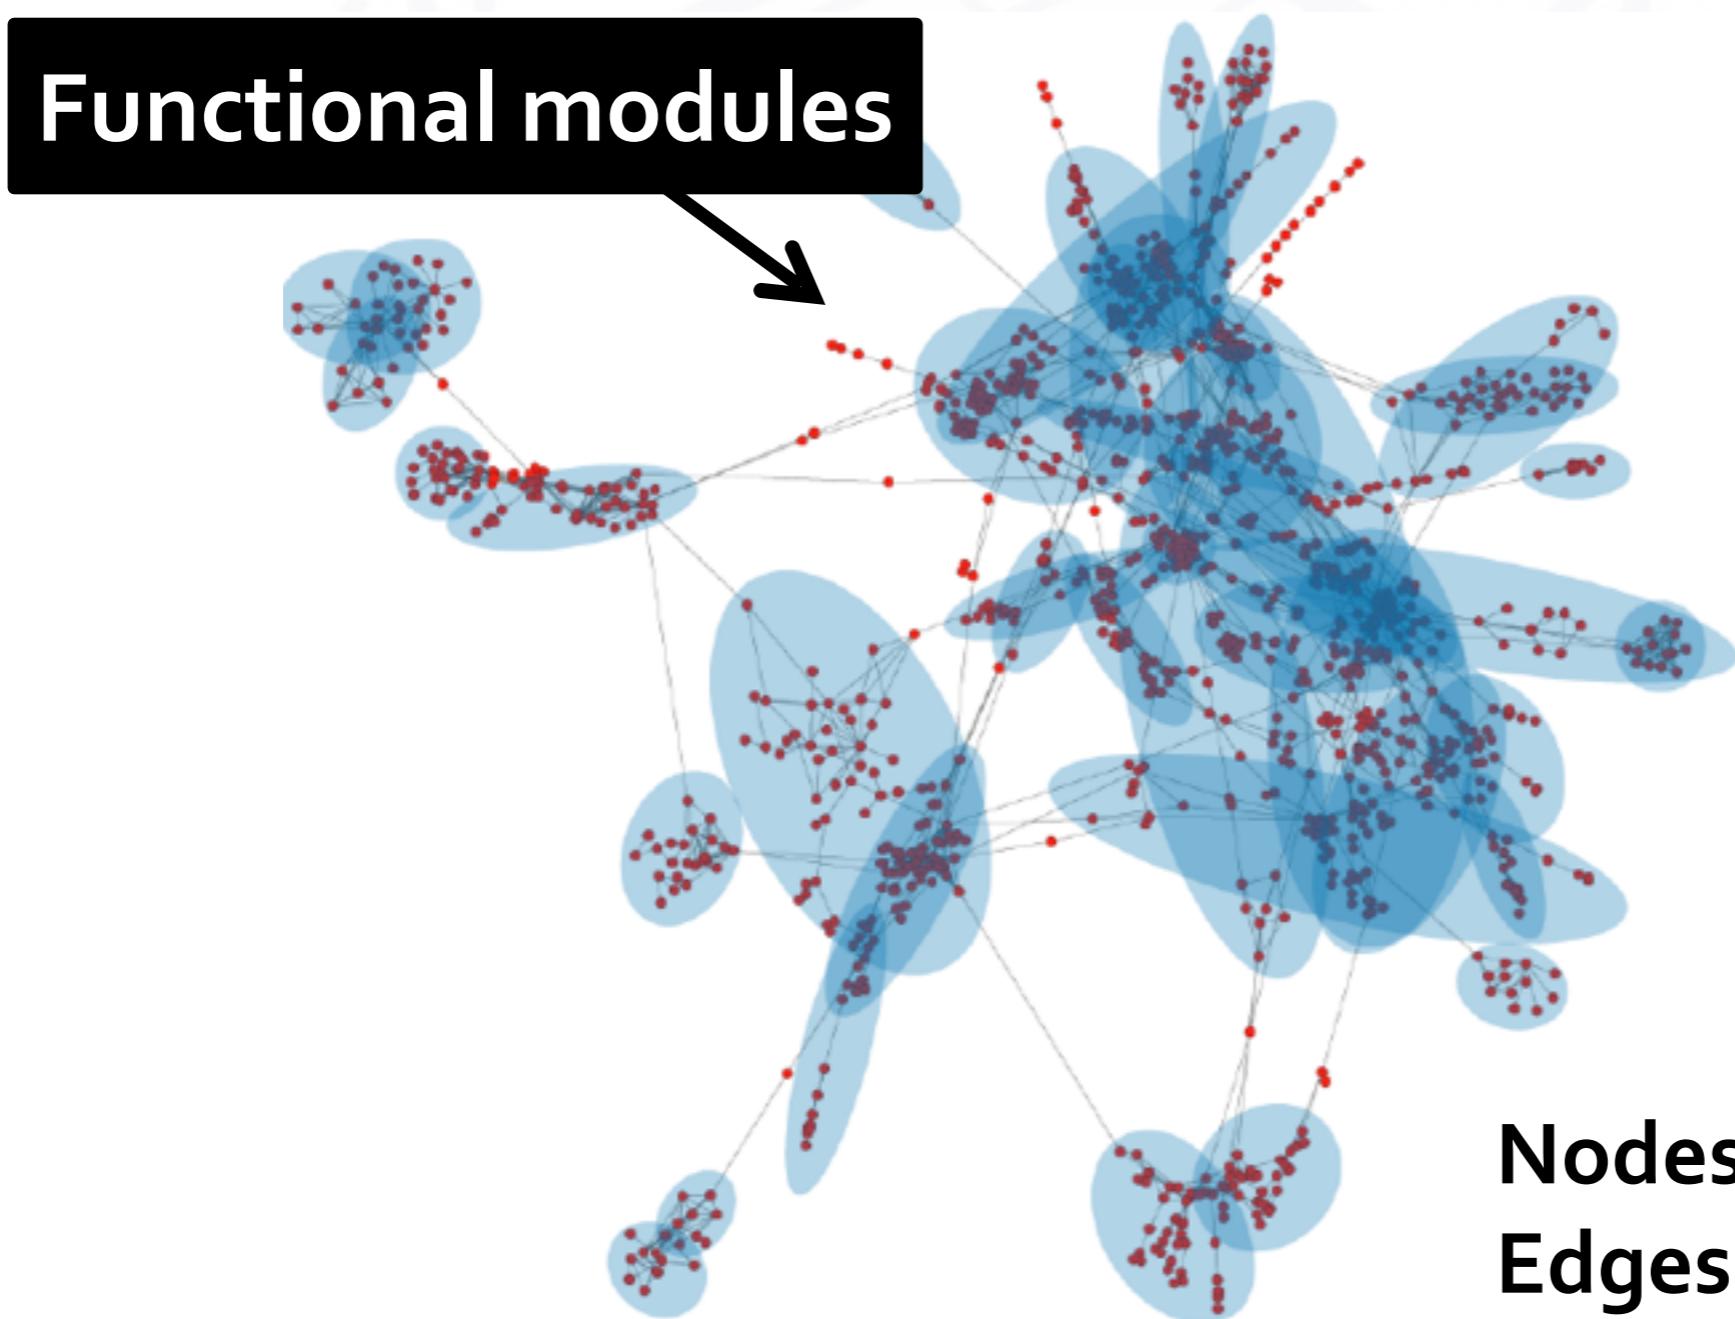
Nodes: Proteins
Edges: Physical interactions

Protein-Protein Interactions



**Nodes: Proteins
Edges: Physical
interactions**

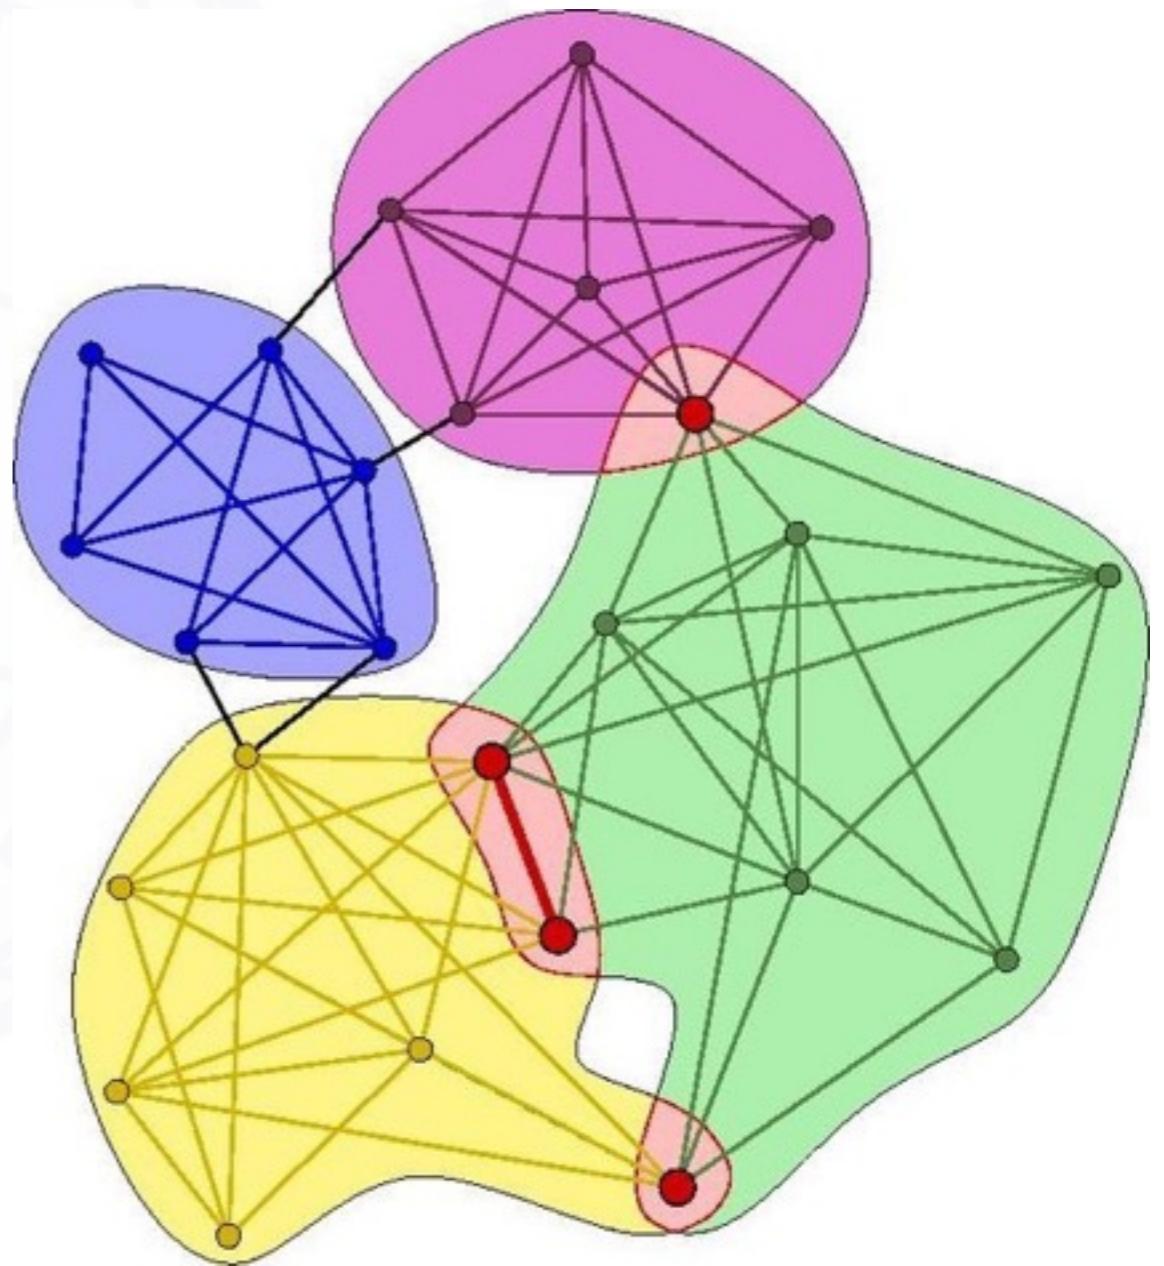
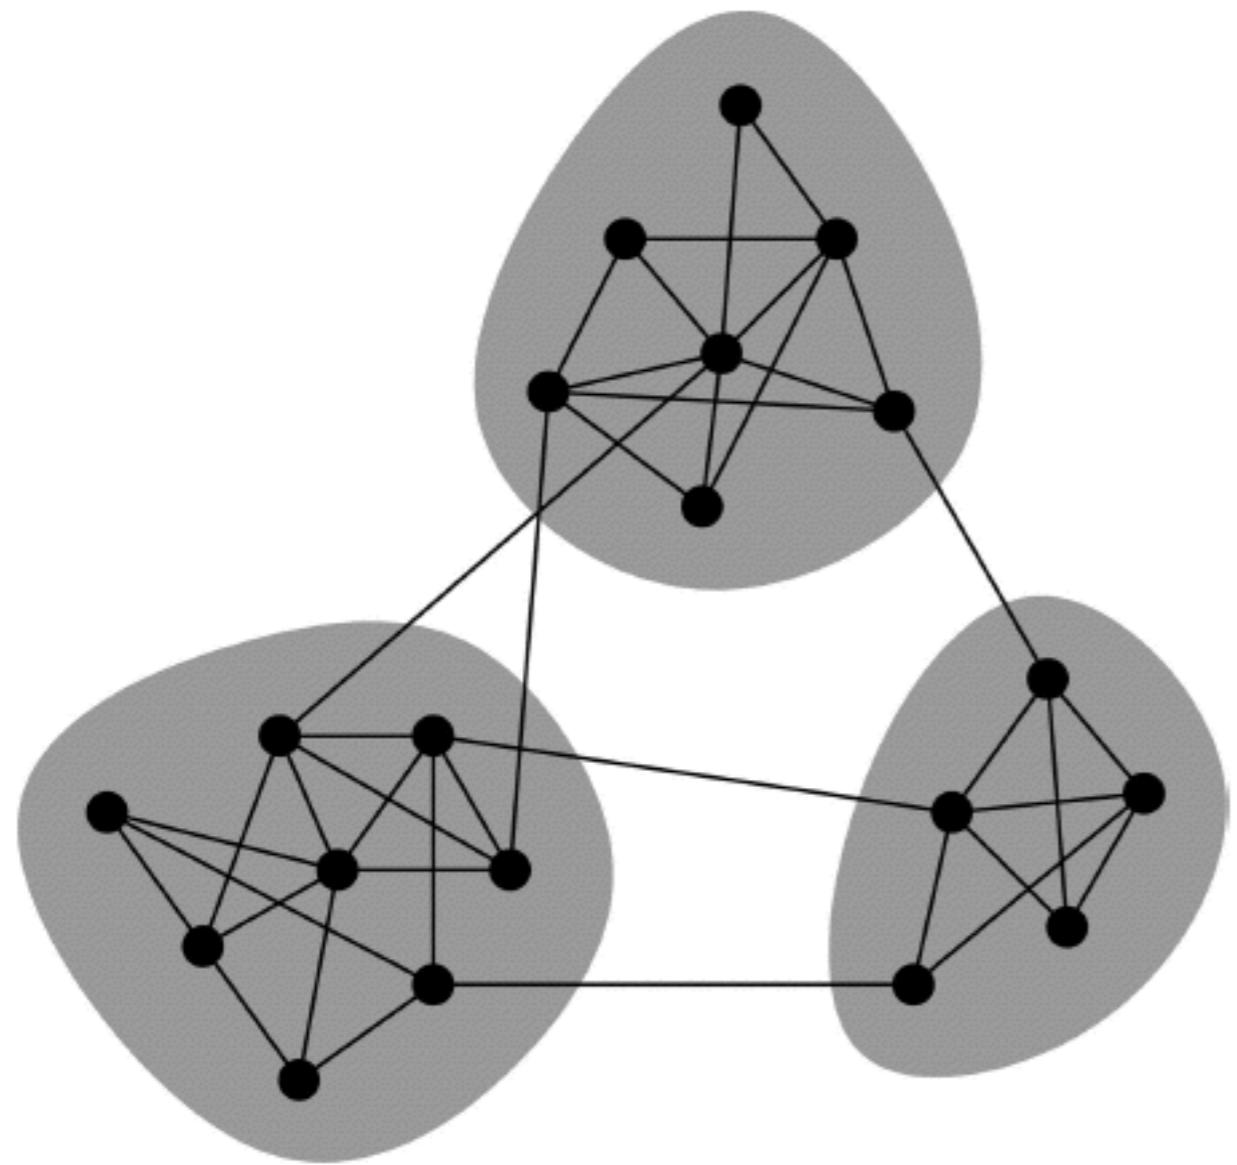
Protein-Protein Interactions



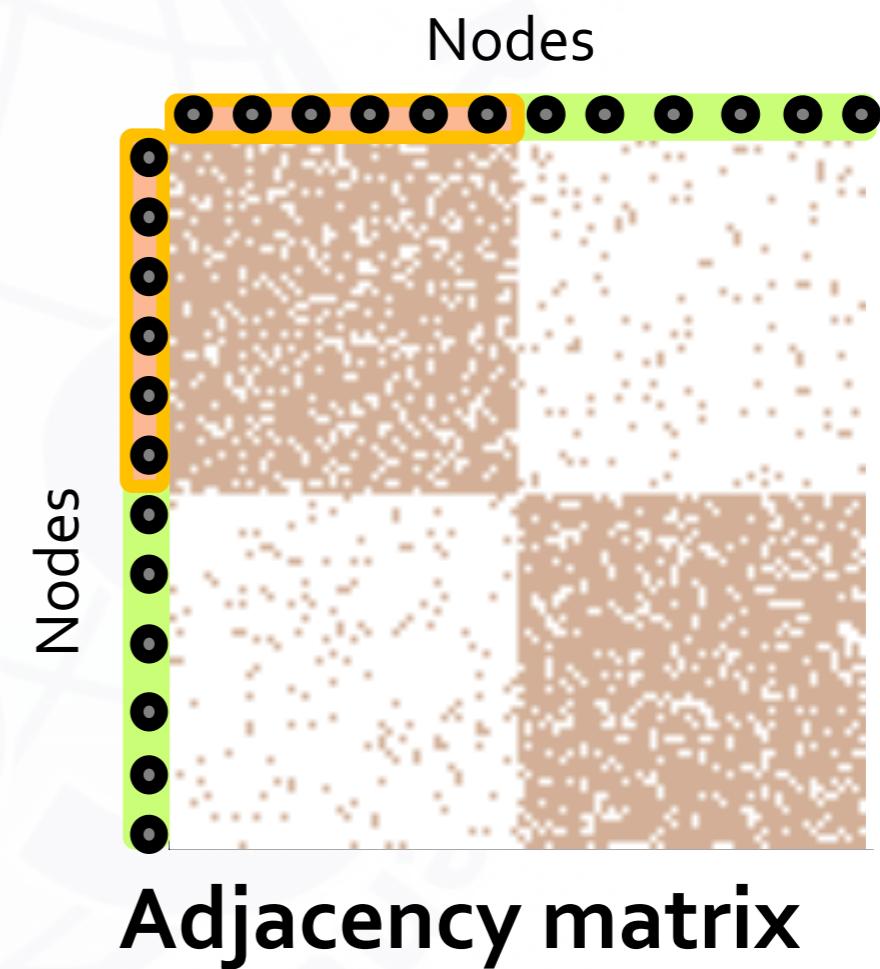
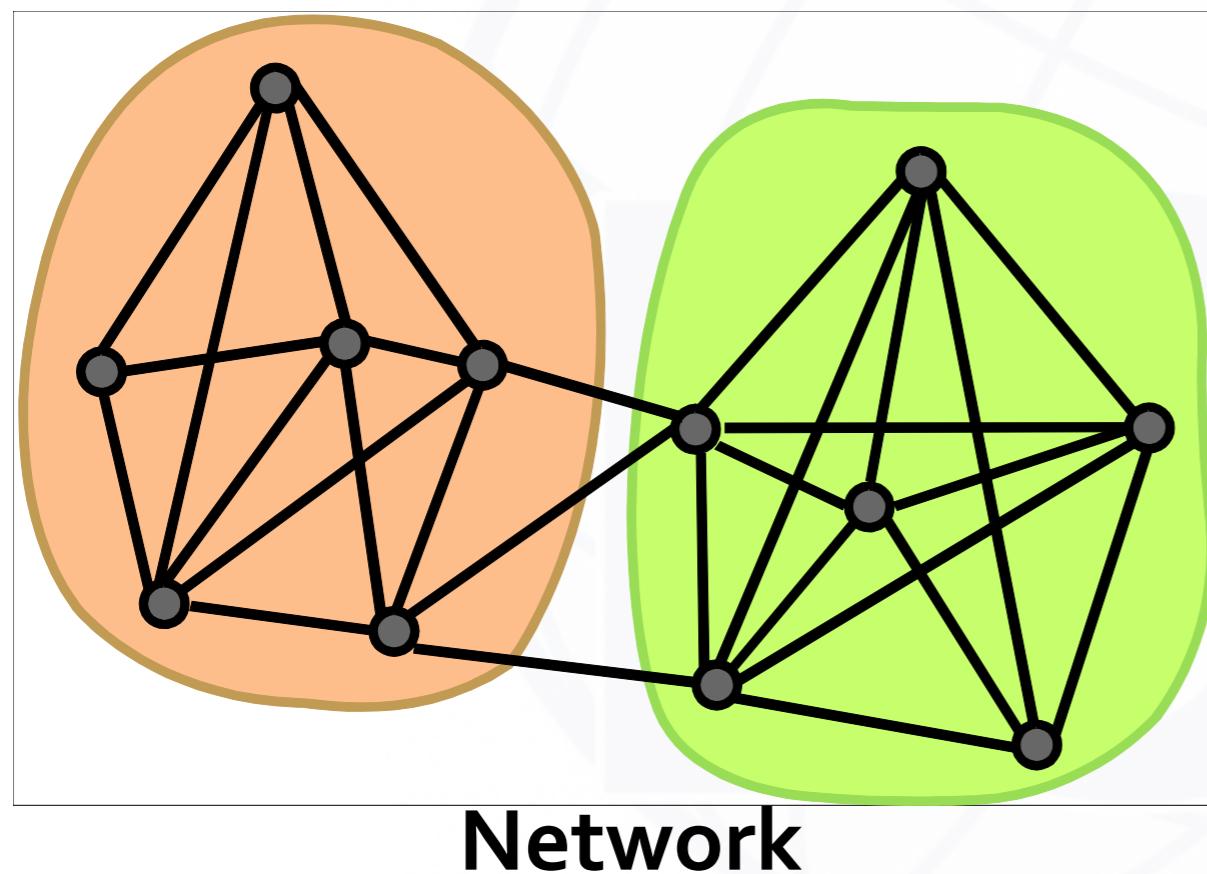
Nodes: Proteins
Edges: Physical
interactions

Overlapping Communities

- Non-overlapping vs. overlapping communities



Non-overlapping Communities

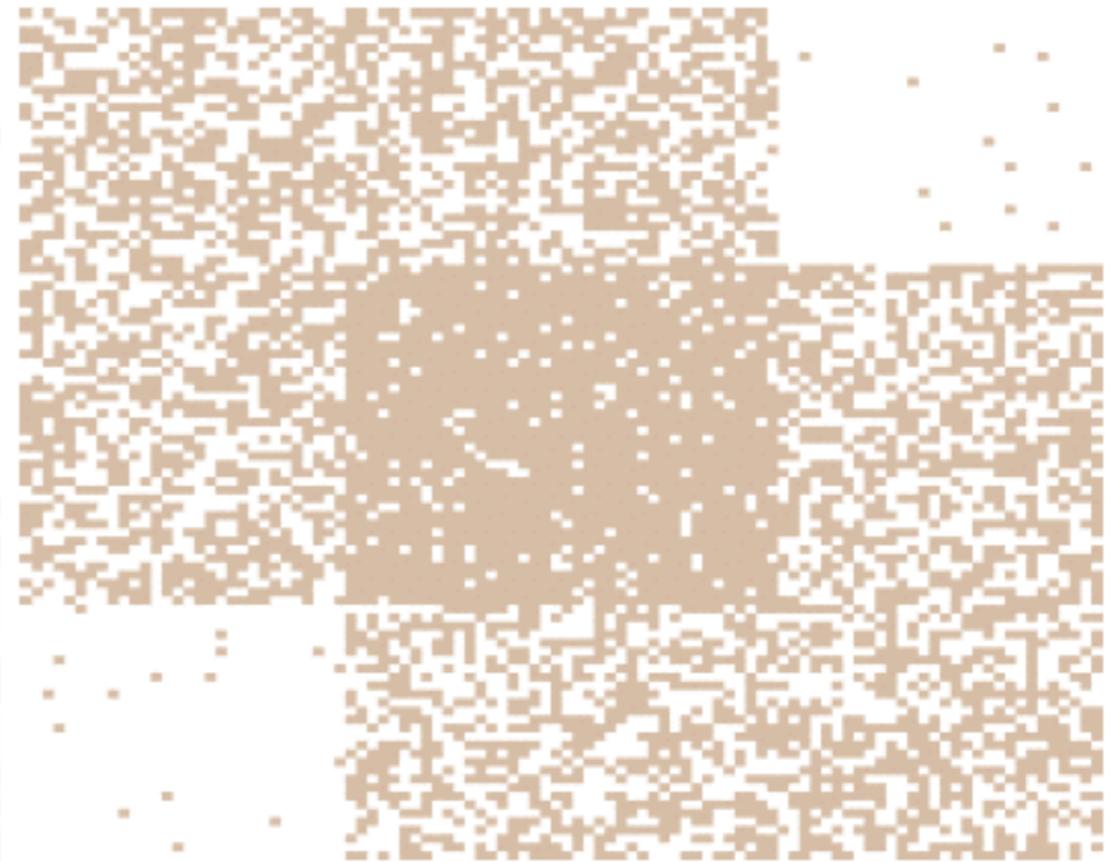
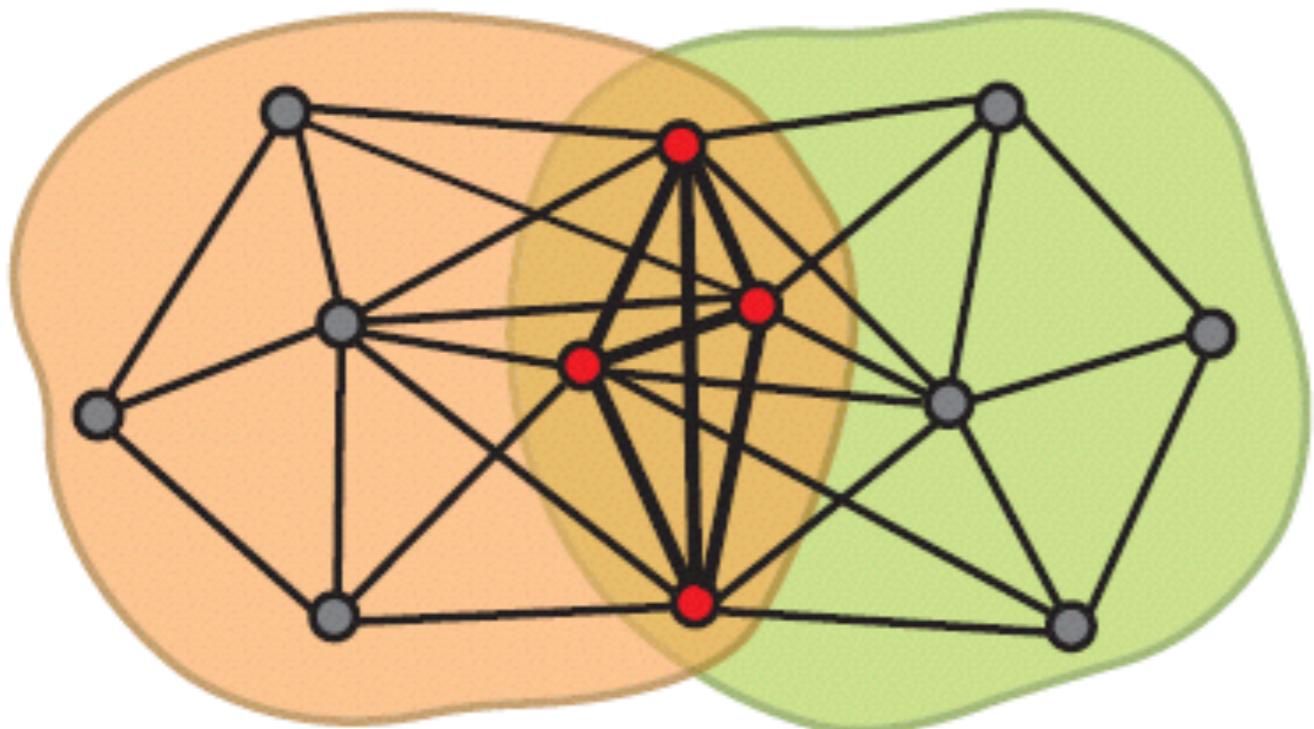


Communities as Tiles!

- **What is the structure of community overlaps:
Edge density in the overlaps is higher!**

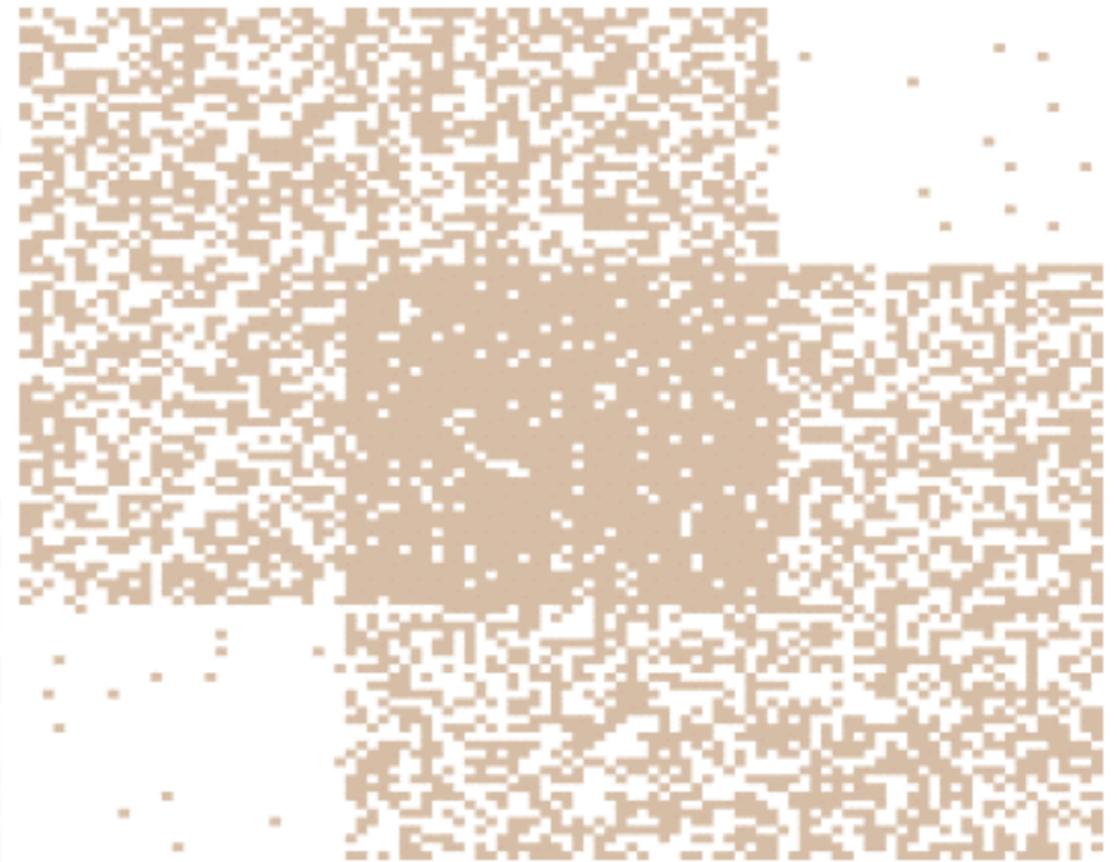
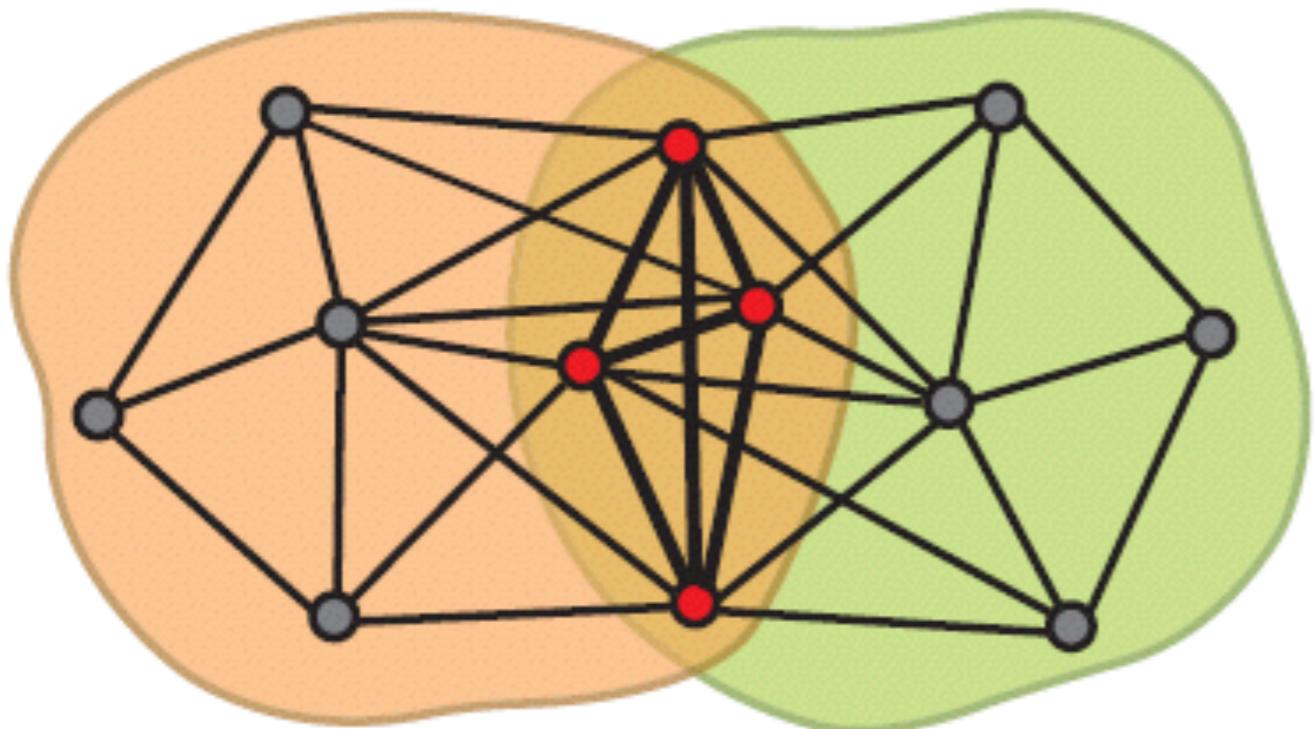
Communities as Tiles!

- What is the structure of community overlaps:
Edge density in the overlaps is higher!



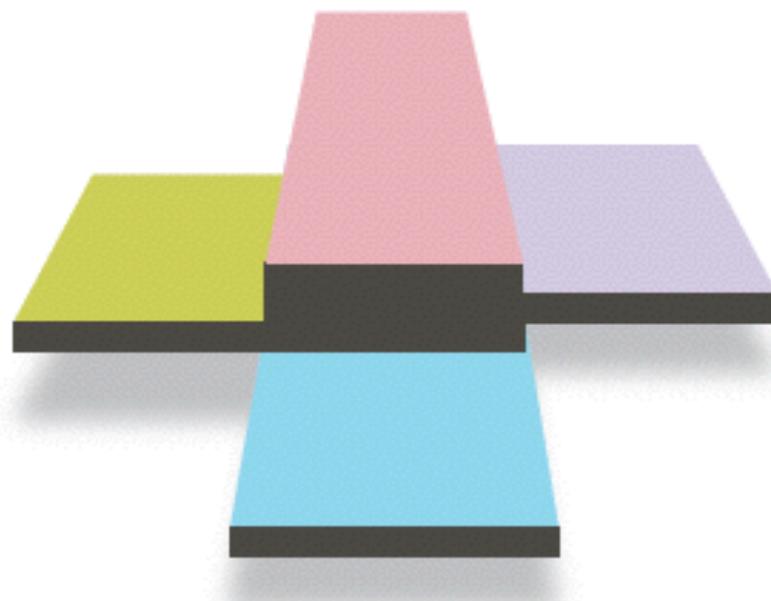
Communities as Tiles!

- What is the structure of community overlaps:
Edge density in the overlaps is higher!



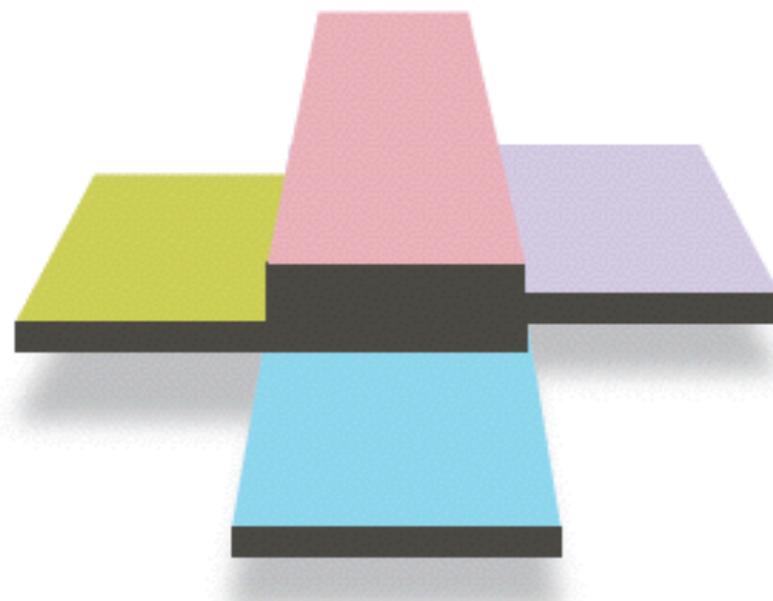
Communities as “tiles”

Recap so far...

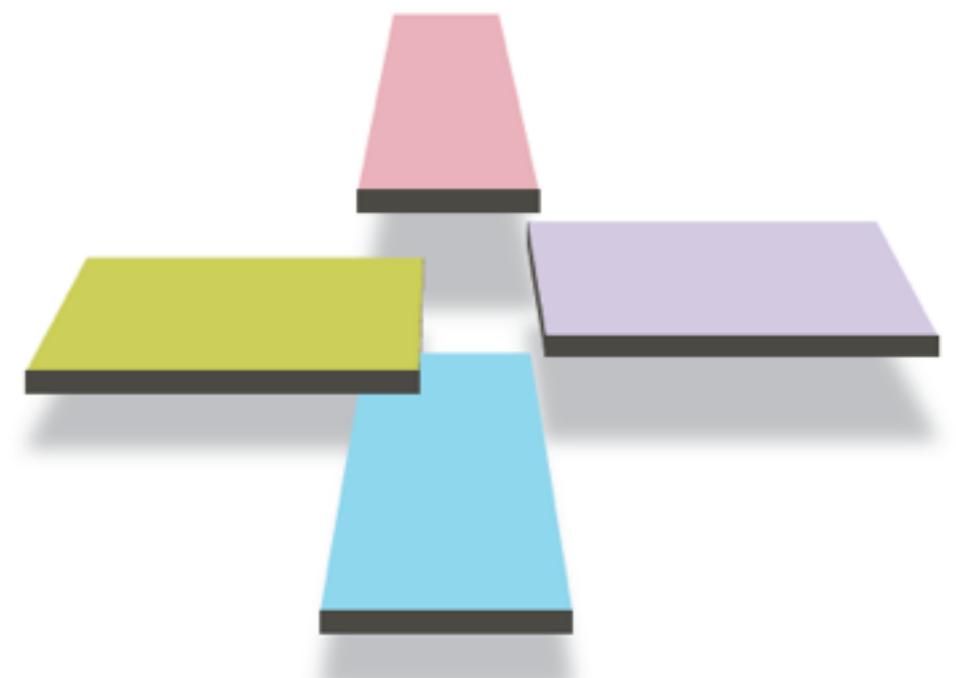


**Communities
in a network**

Recap so far...

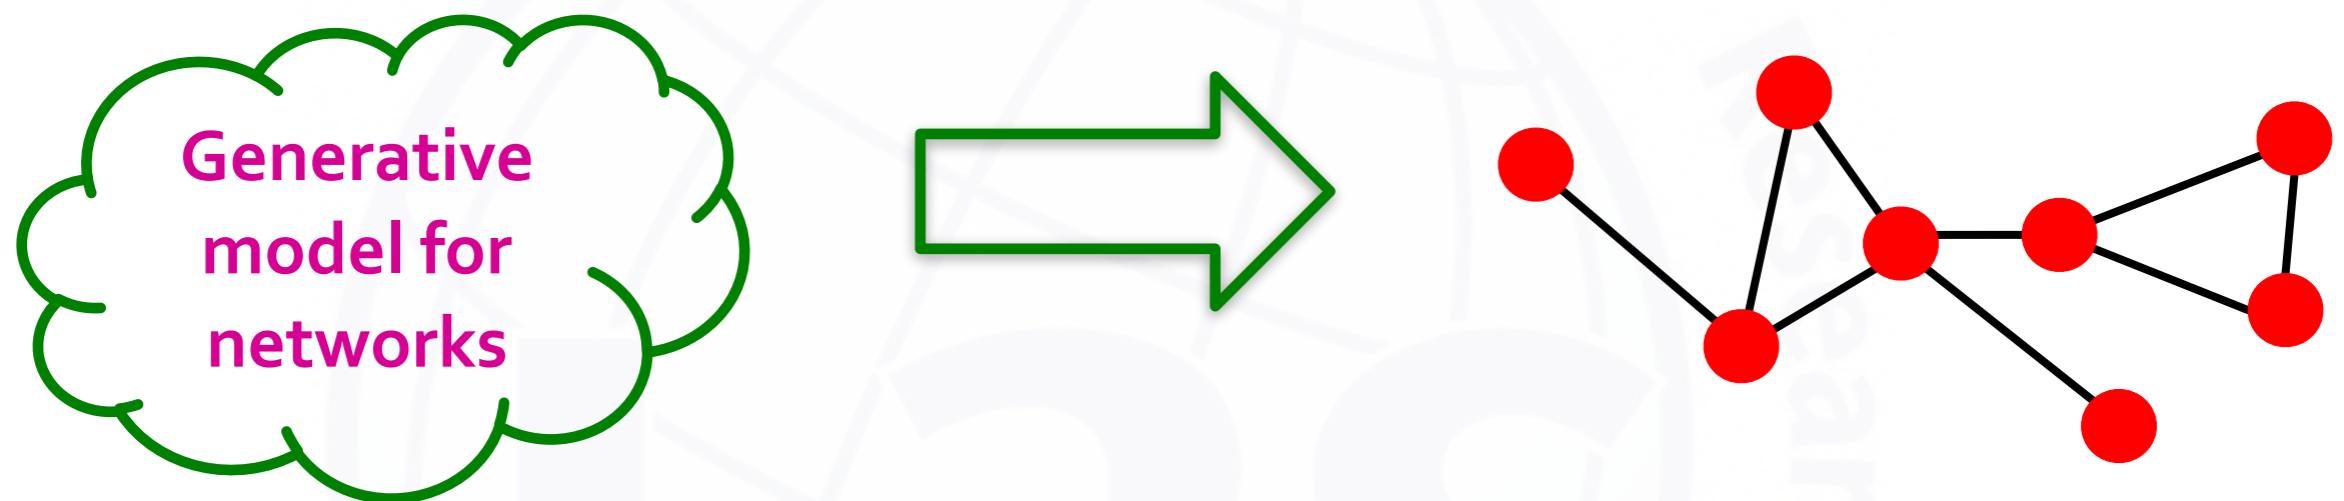


**Communities
in a network**



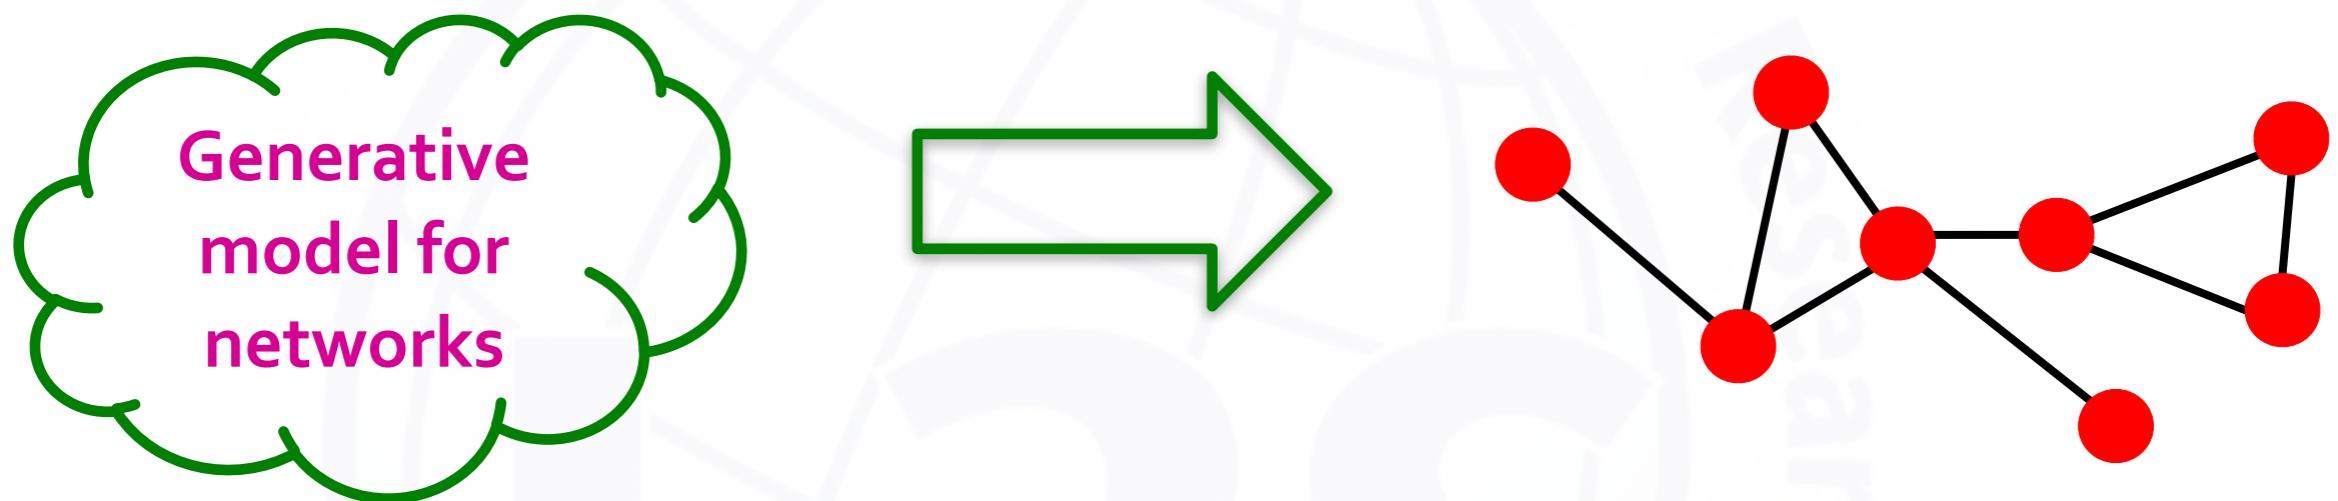
This is what we want!

Plan of attack



Plan of attack

- 1) Given a model, we generate the network:

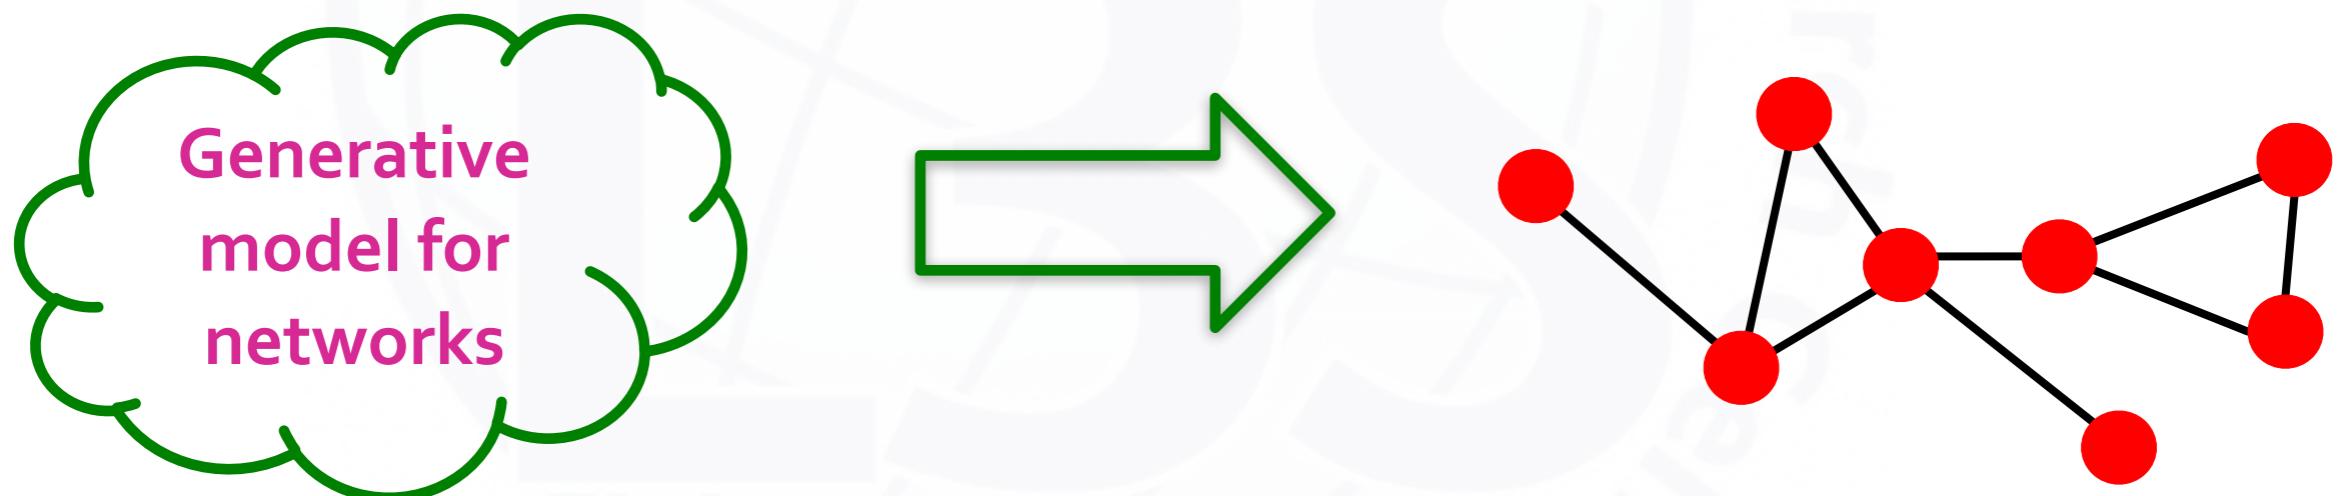


- 2) Given a network, find the “best” model



Model of networks

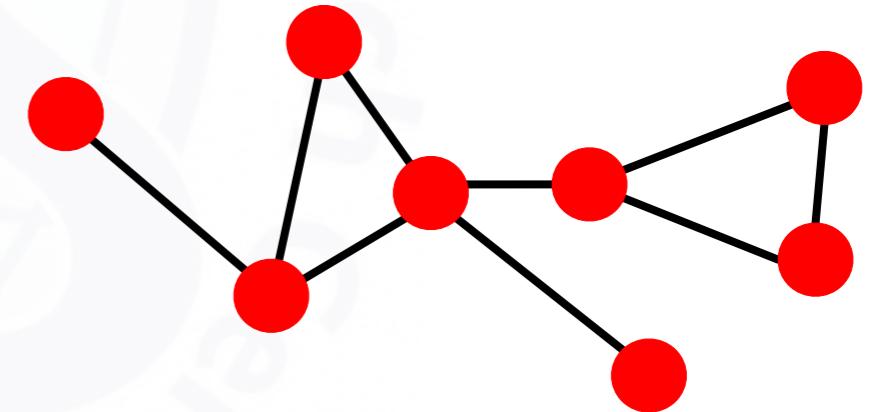
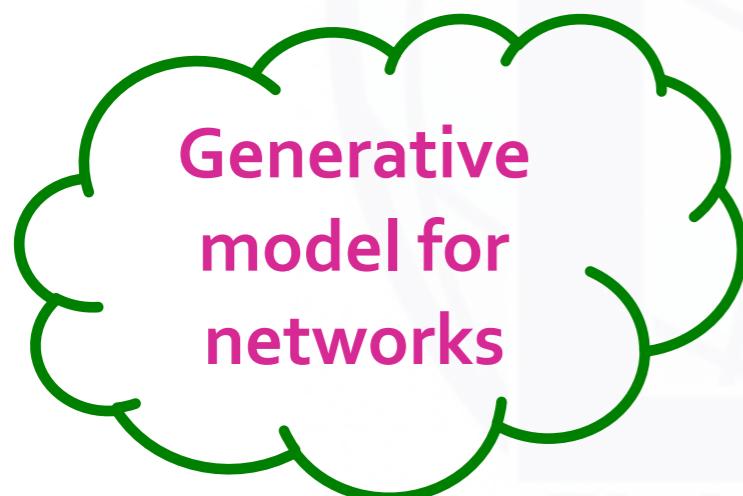
- Goal: Define a model that can generate networks
 - The model will have a set of “parameters” that we will later want to estimate (and detect communities)



- Q: Given a set of nodes, how do communities “generate” edges of the network?

Model of networks

- Goal: Define a model that can generate networks
 - The model will have a set of “parameters” that we will later want to estimate (and detect communities)



- Q: Given a set of nodes, how do communities “generate” edges of the network?

Community-Affiliation Graph

Nodes, V



Model

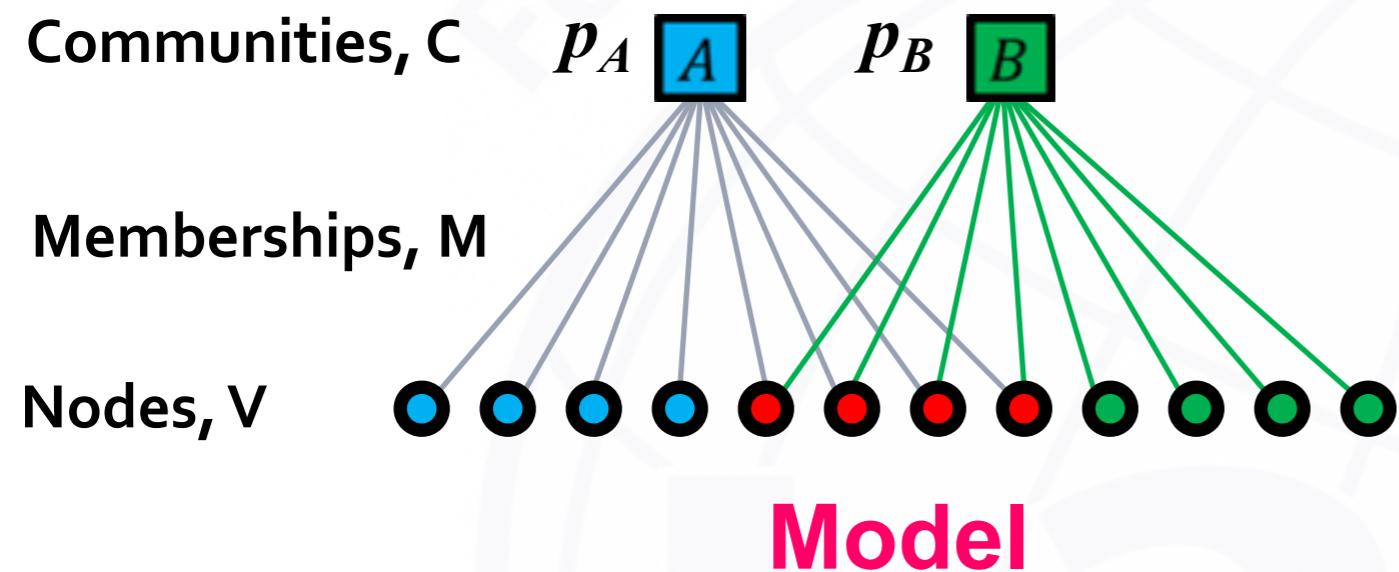
Community-Affiliation Graph

Communities, C p_A  p_B 

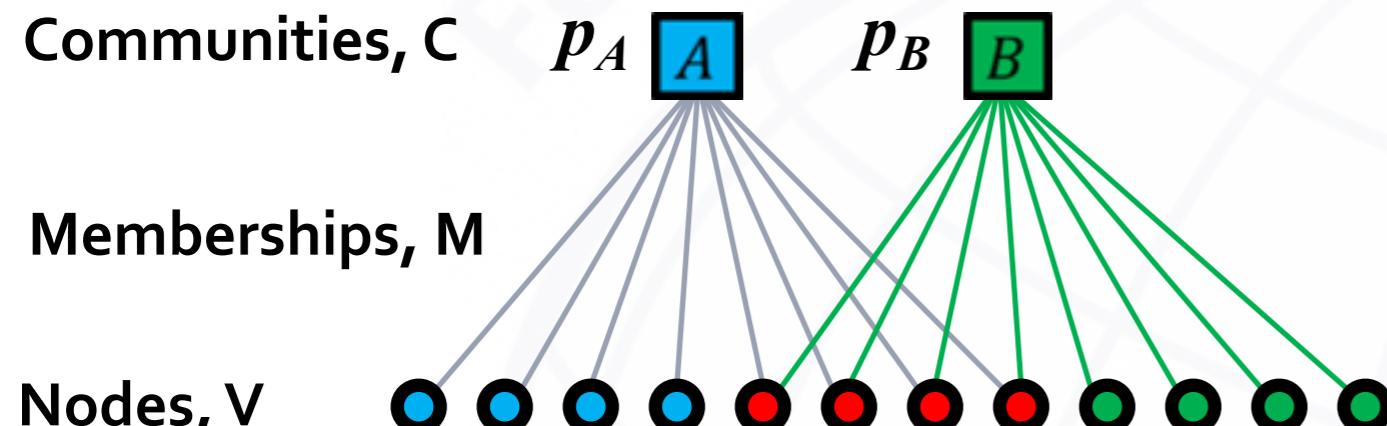
Nodes, V 

Model

Community-Affiliation Graph



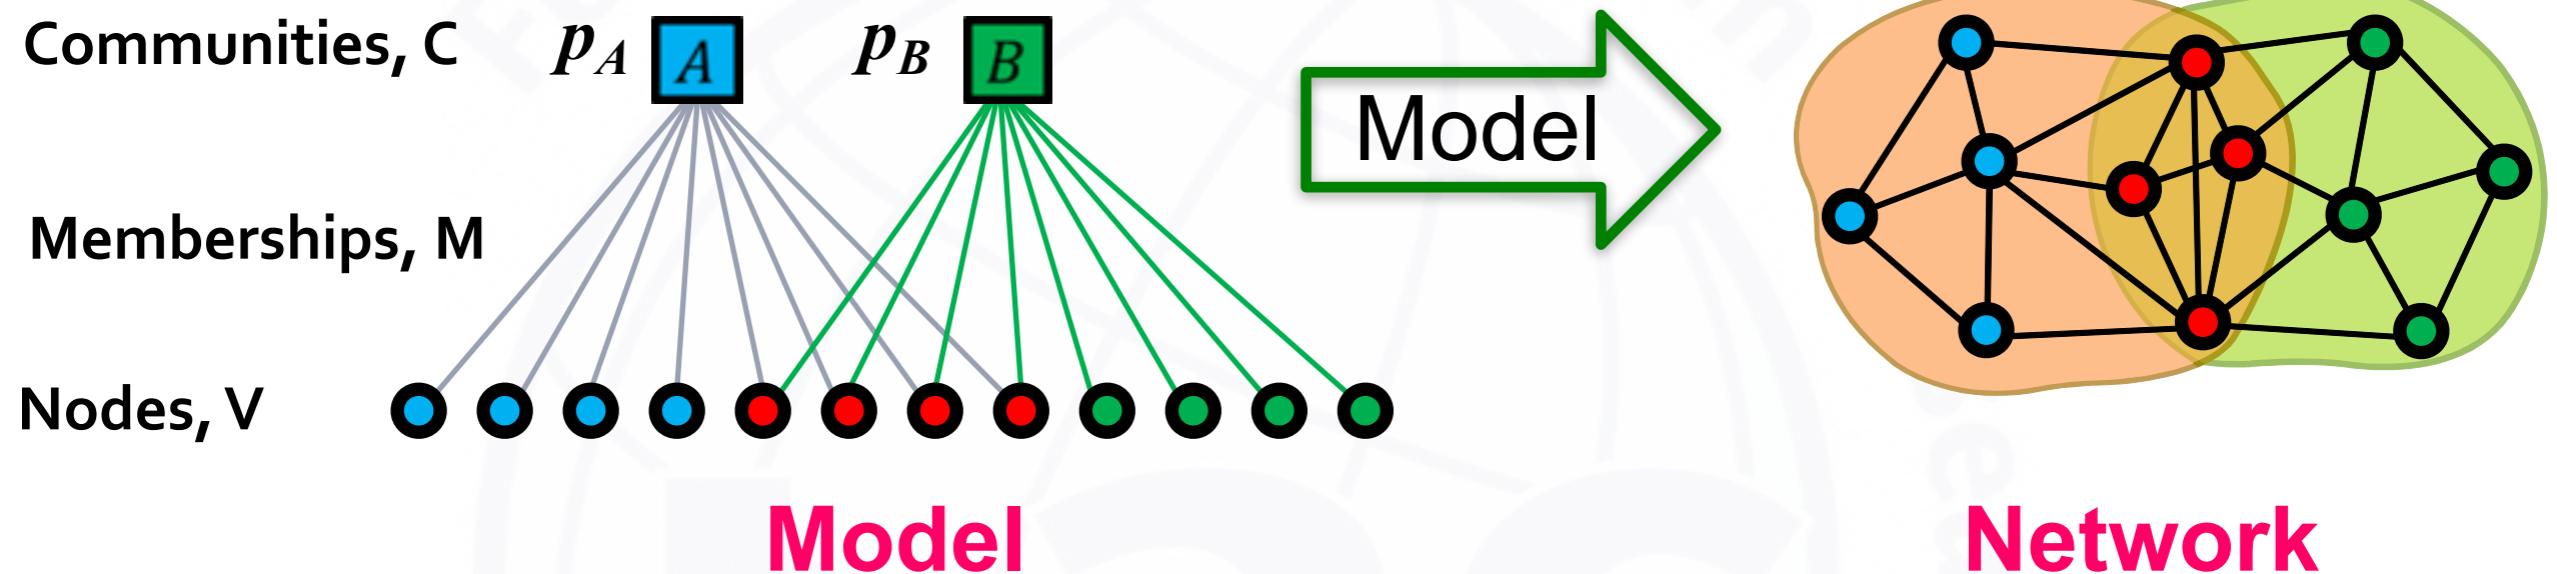
Community-Affiliation Graph



Model

- **Generative model $B(V, C, M, \{p_c\})$ for graphs:**
 - Nodes V , Communities C , Memberships M
 - Each community c has a single probability p_c
- Later we fit the model to networks to detect communities

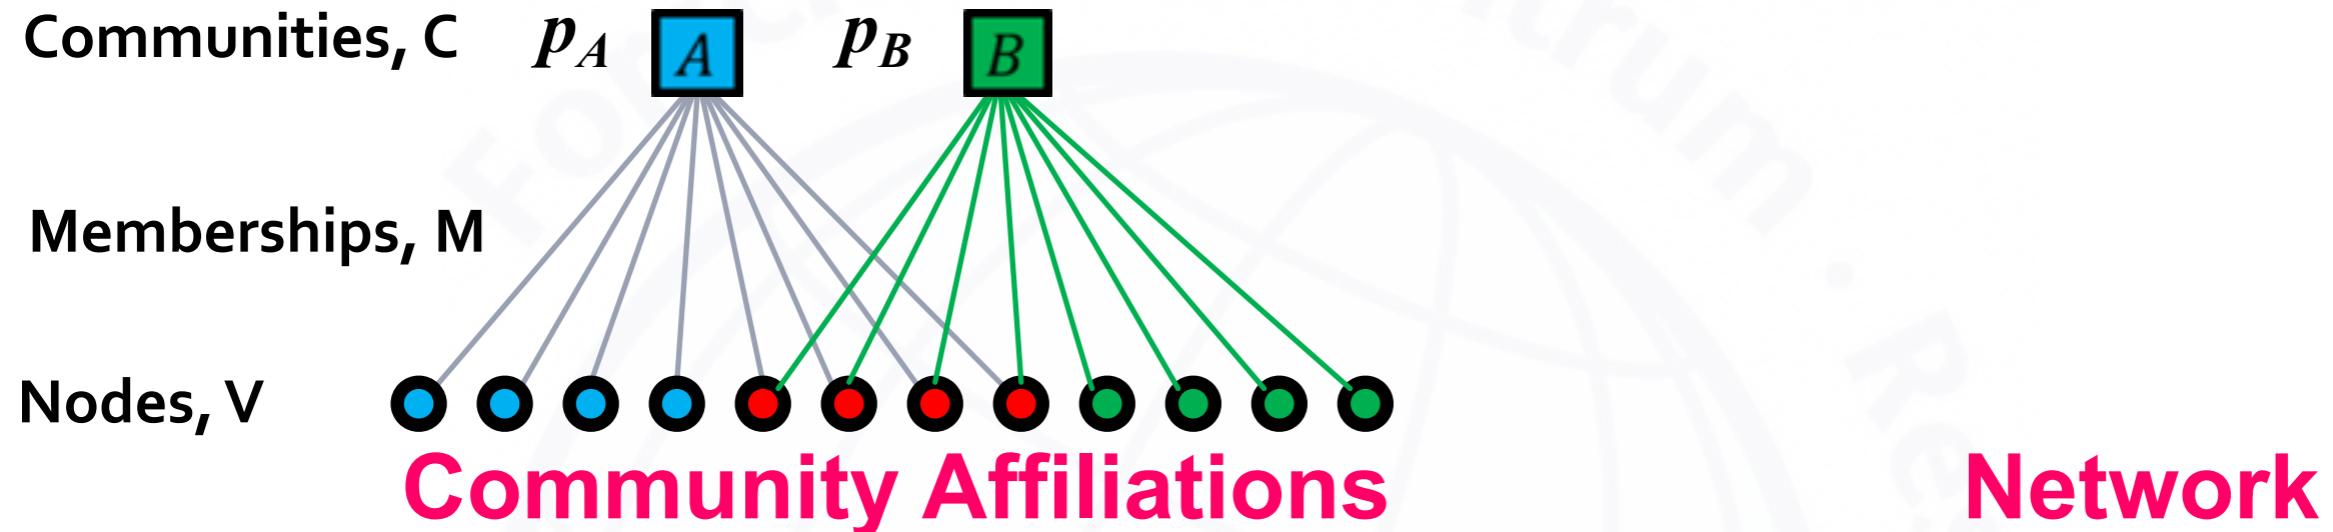
Community-Affiliation Graph



- **Generative model $B(V, C, M, \{p_c\})$ for graphs:**

- Nodes V , Communities C , Memberships M
- Each community c has a single probability p_c
- Later we fit the model to networks to detect communities

AGM: Generative Process

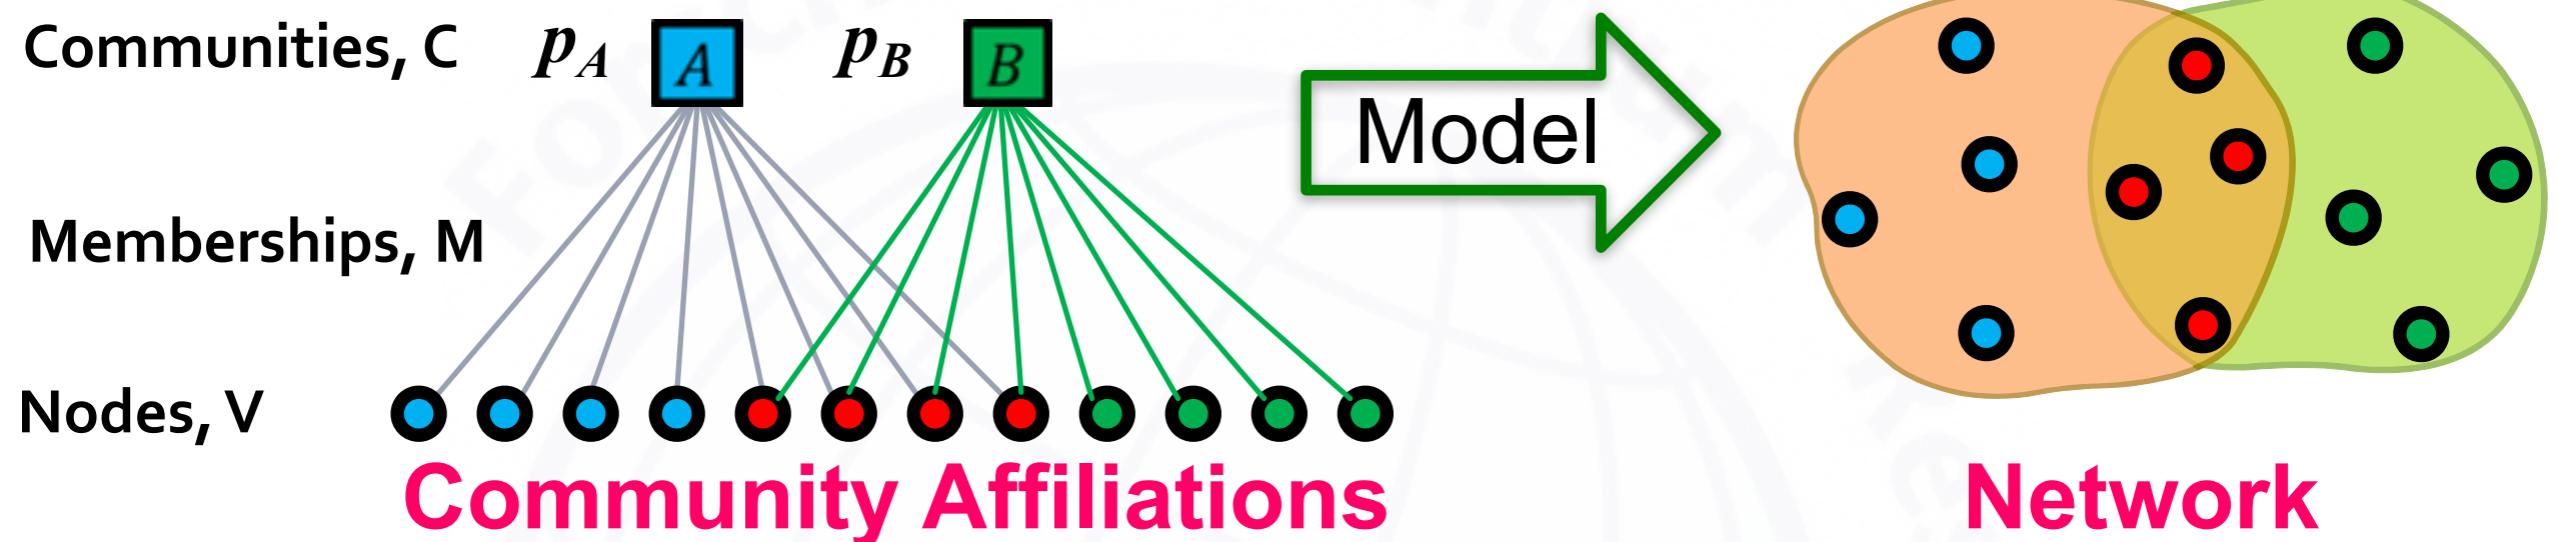


- **AGM generates the links: For each**
 - For each pair of nodes in community A , we connect them with prob. p_A
 - The overall edge probability is:

\mathcal{M}_u ... set of communities
node u belongs to

Think of this as an “OR” function: If at least 1 community says “YES” we create an edge

AGM: Generative Process



- **AGM generates the links: For each**
 - For each pair of nodes in community A , we connect them with prob. p_A
 - The overall edge probability is:

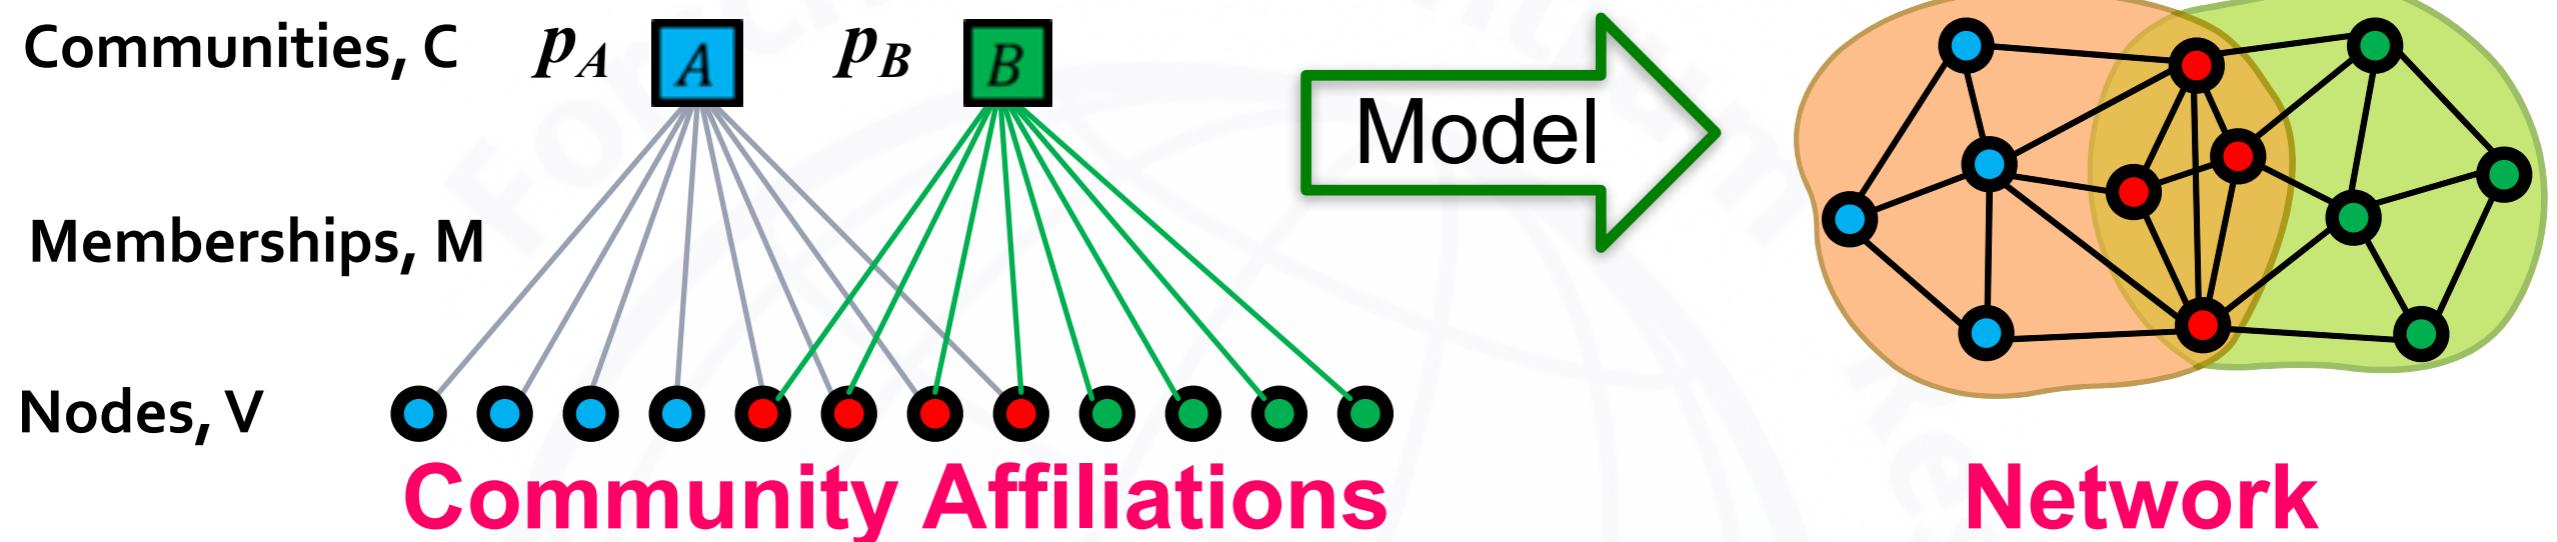
$$P(u, v) = 1 - \prod_{c \in M_u \cap M_v} (1 - p_c)$$

If u, v share no communities: $P(u, v) = \epsilon$

M_u ... set of communities
node u belongs to

Think of this as an “OR” function: If at least 1 community says “YES” we create an edge

AGM: Generative Process



- **AGM generates the links: For each**
 - For each pair of nodes in community A , we connect them with prob. p_A
 - The overall edge probability is:

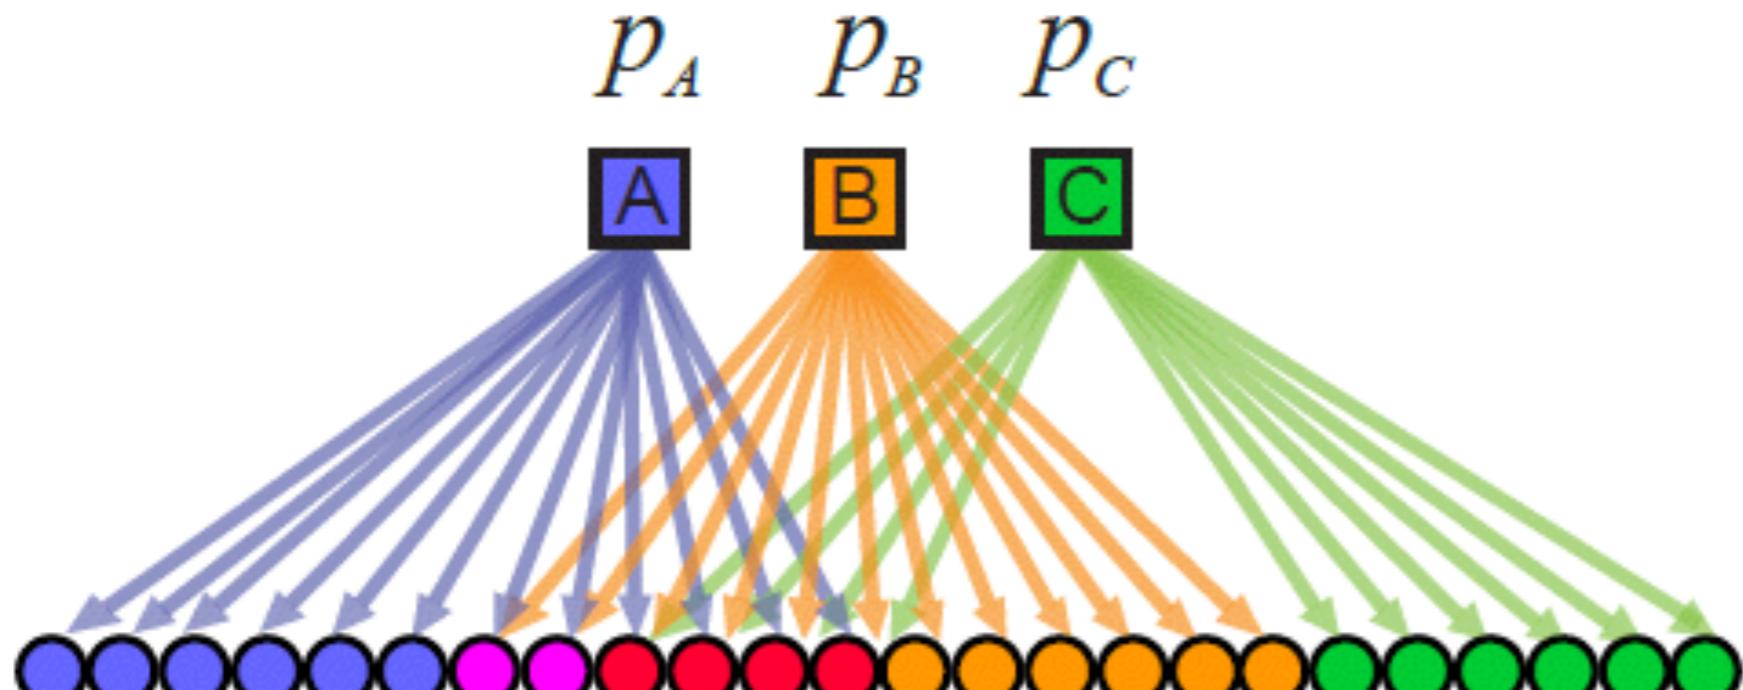
$$P(u, v) = 1 - \prod_{c \in M_u \cap M_v} (1 - p_c)$$

If u, v share no communities: $P(u, v) = \epsilon$

M_u ... set of communities
node u belongs to

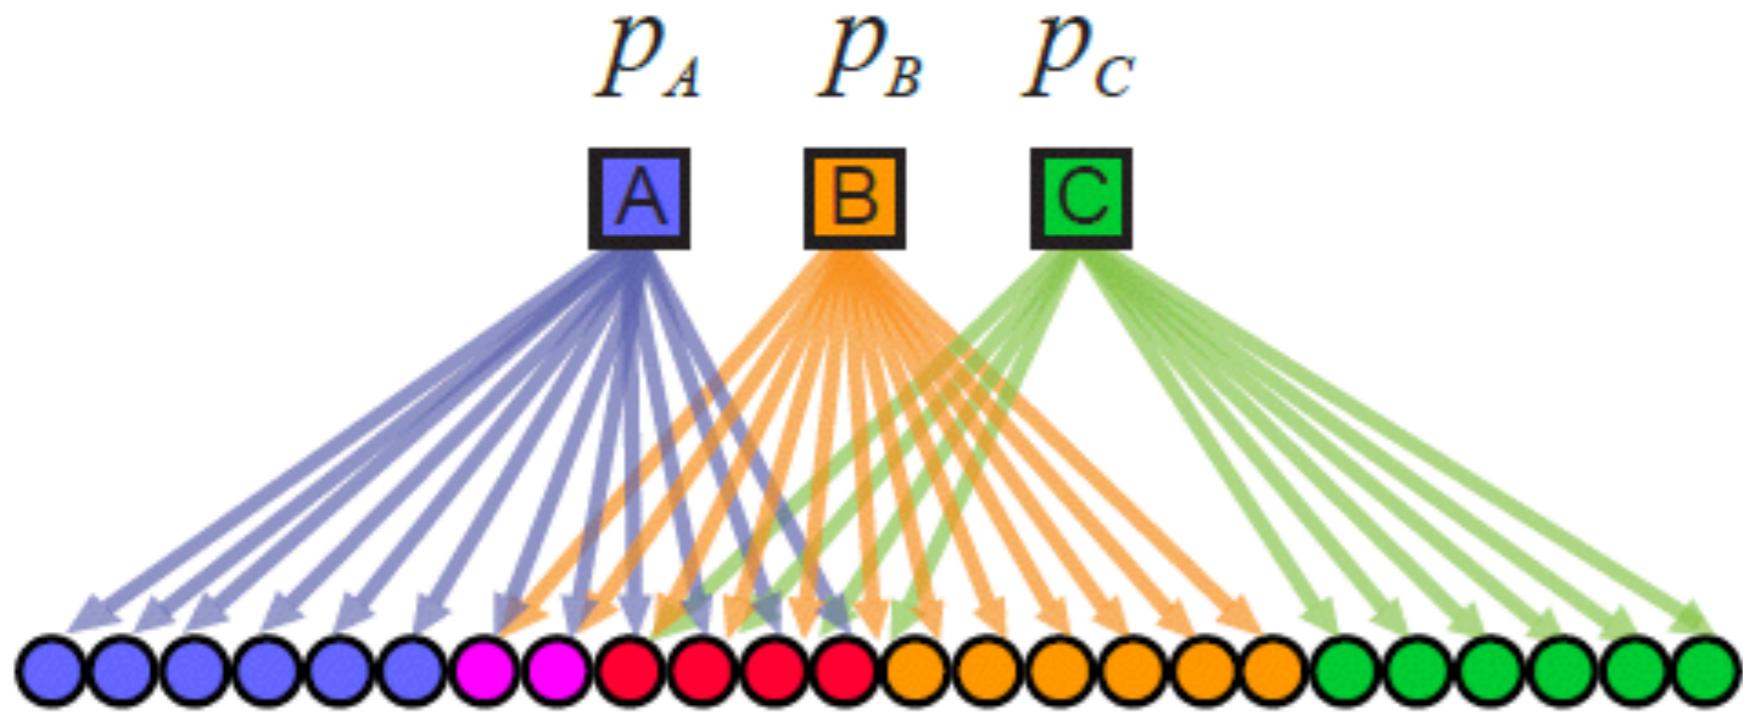
Think of this as an “OR” function: If at least 1 community says “YES” we create an edge

Recap: AGM networks

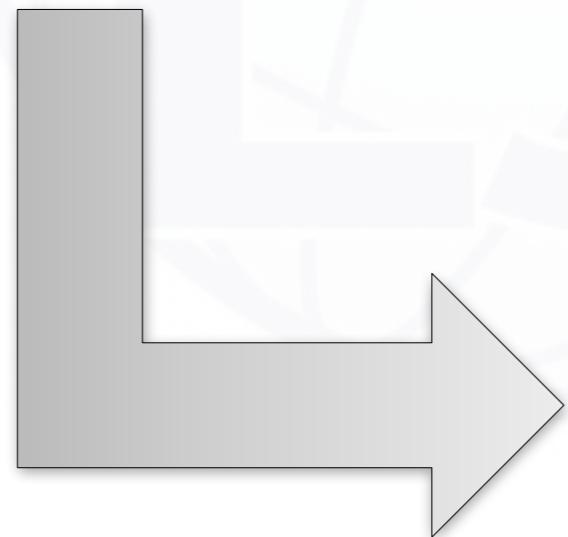


Model

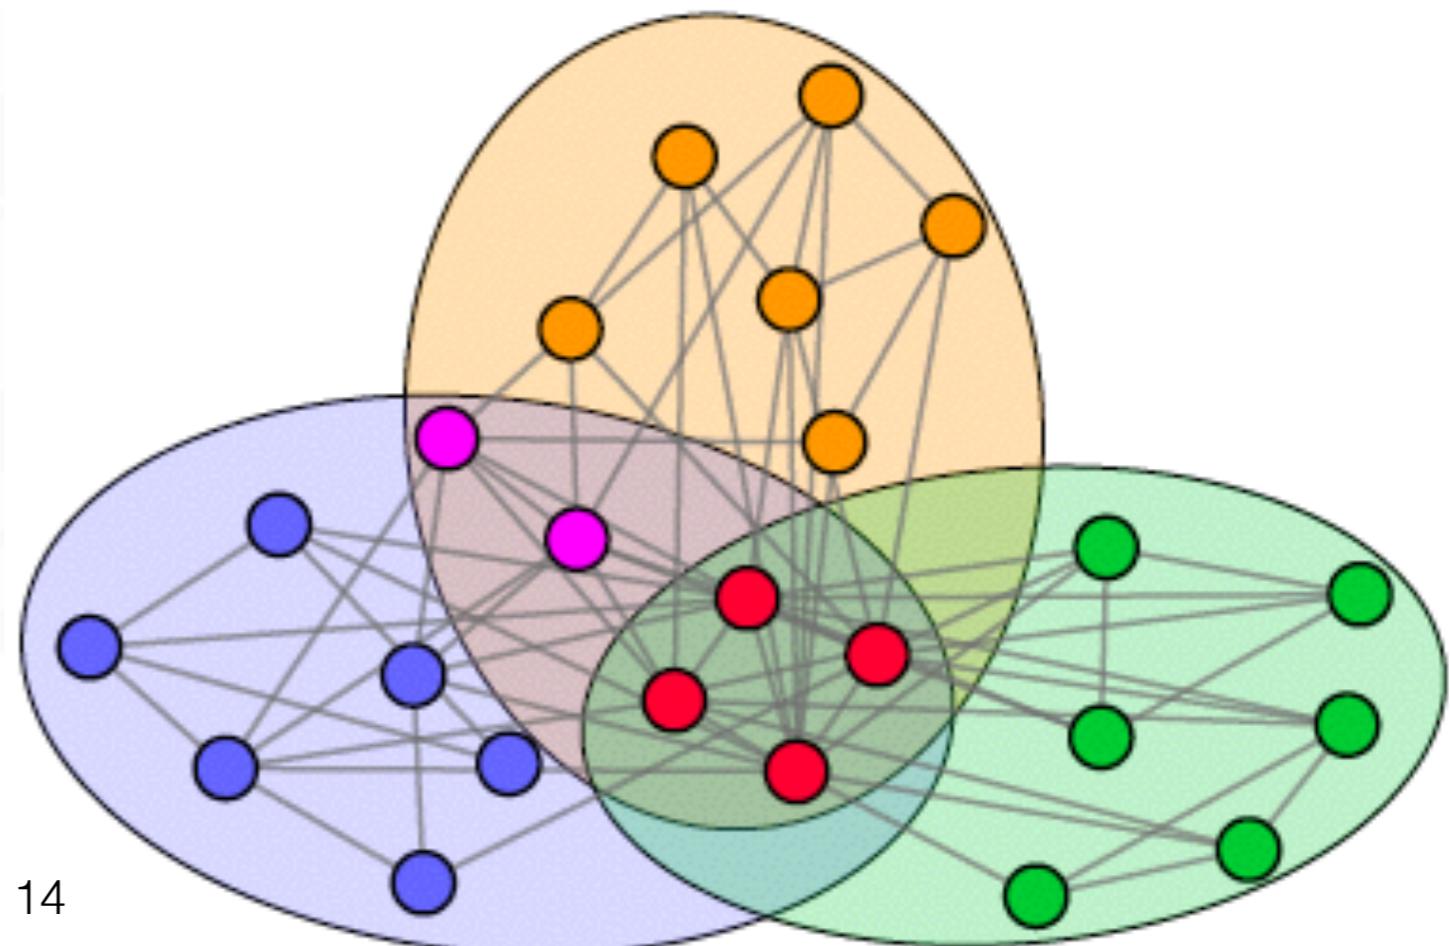
Recap: AGM networks



Model



Network

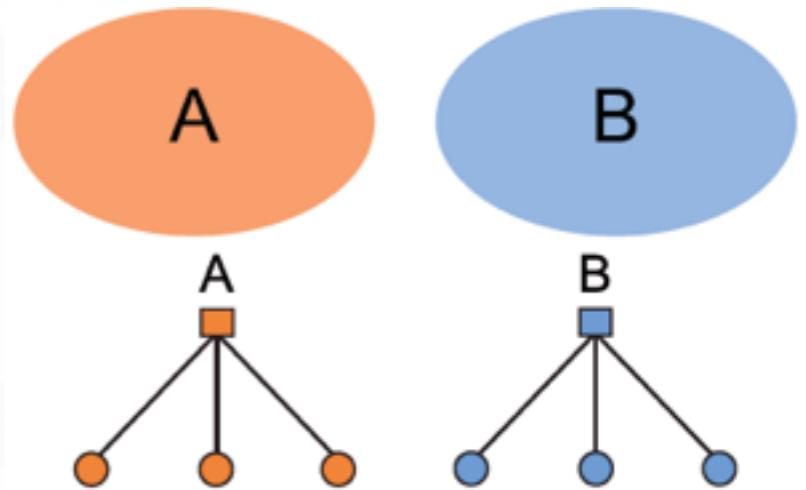


AGM: Flexibility

- AGM can express a variety of community structures:
Non-overlapping, Overlapping,
Nested

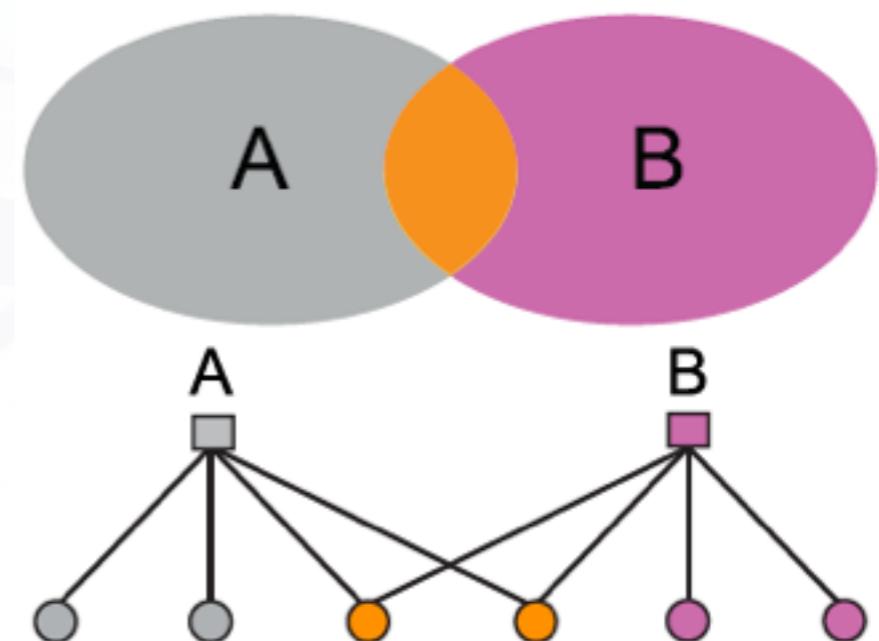
AGM: Flexibility

- AGM can express a variety of community structures:
Non-overlapping, Overlapping,
Nested



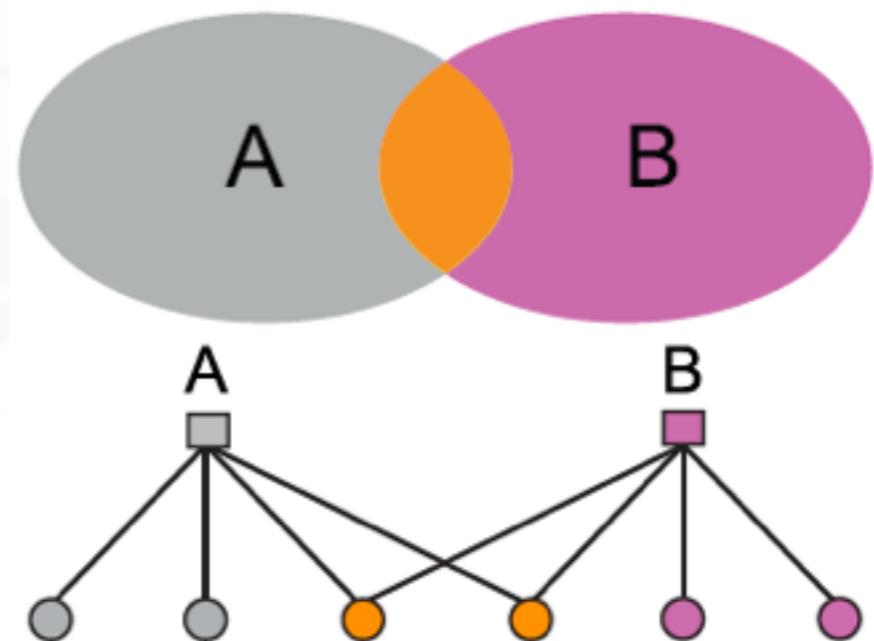
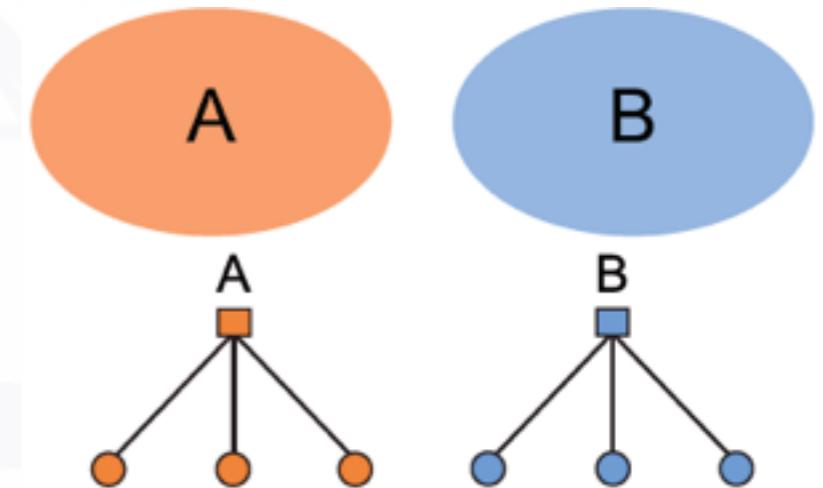
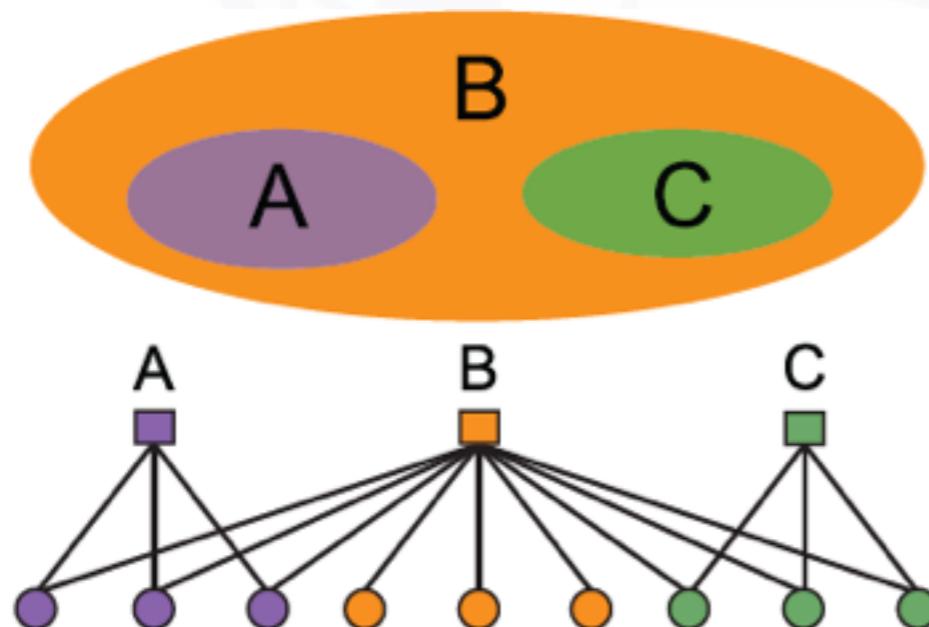
AGM: Flexibility

- AGM can express a variety of community structures:
Non-overlapping, Overlapping,
Nested



AGM: Flexibility

- AGM can express a variety of community structures:
Non-overlapping, Overlapping,
Nested

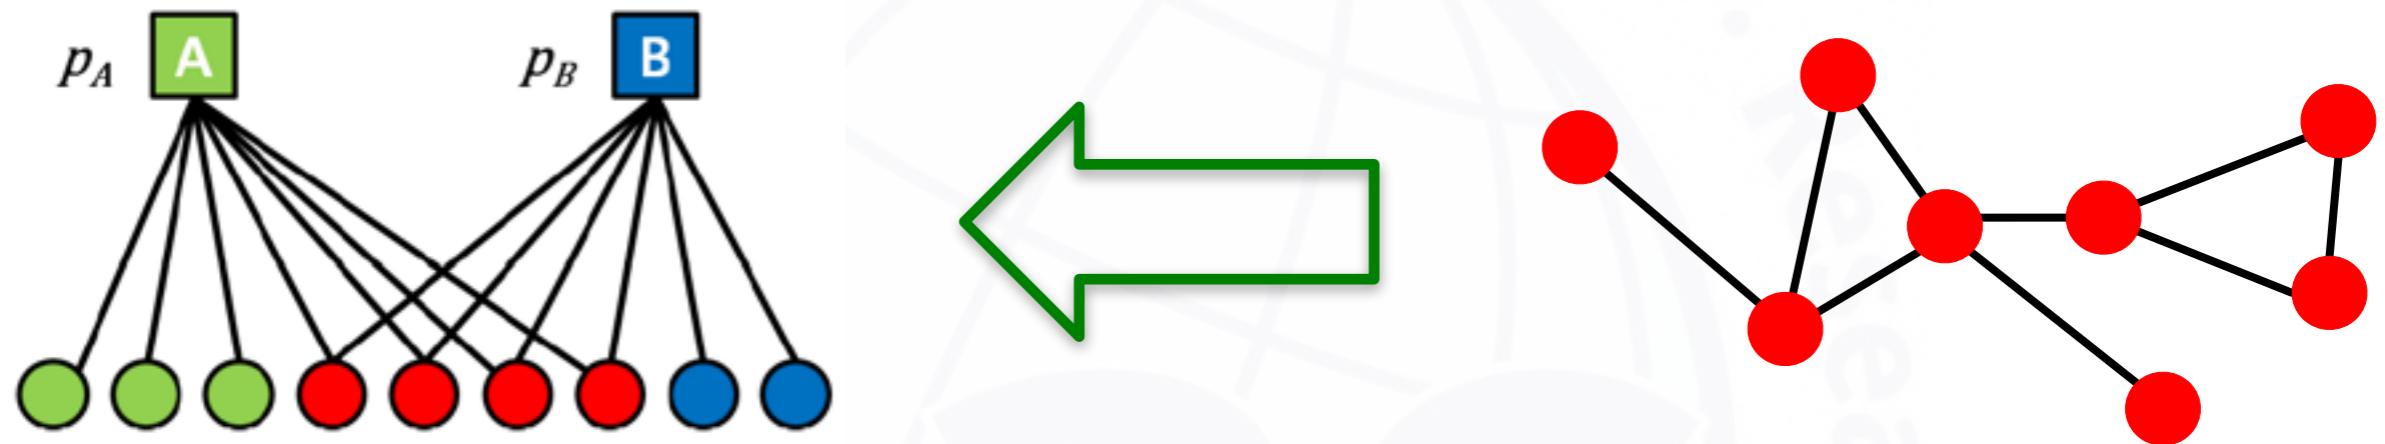


Detecting Communities



- **Detecting communities with AGM:**

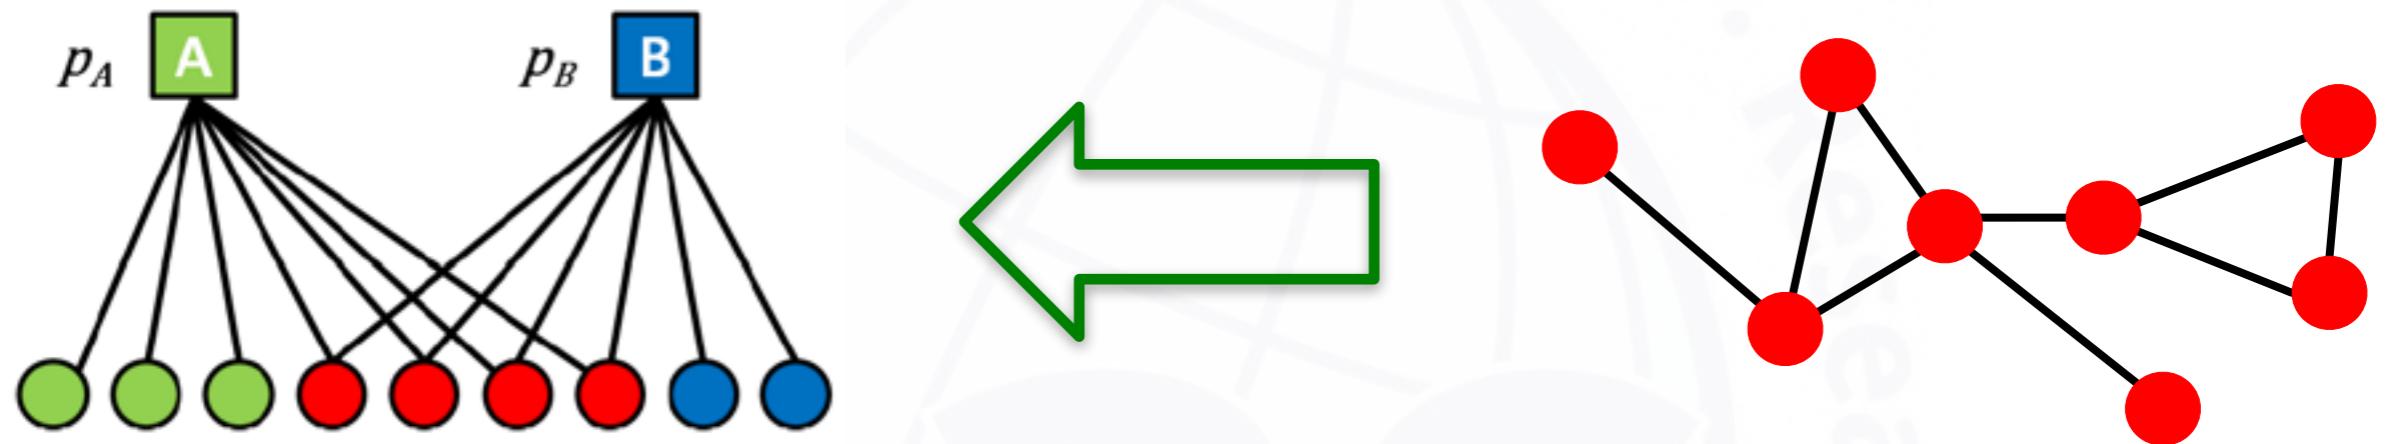
Detecting Communities



- Detecting communities with AGM:

Given a Graph $G(V, E)$, find the Model

Detecting Communities



- Detecting communities with AGM:

Given a Graph $G(V, E)$, find the Model

- 1) Affiliation graph M
- 2) Number of communities C
- 3) Parameters p_c

Likelihood of a Graph

- Given a set of parameters, the likelihood of the graph is computed by

Nodes, V 

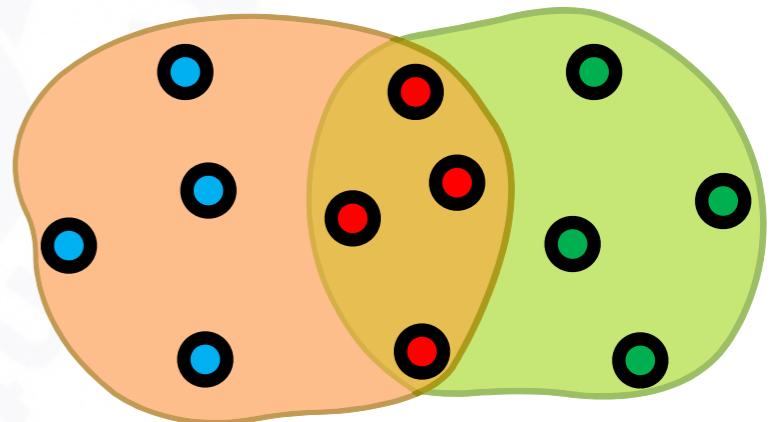
Estimate parameters $\hat{\theta}$ by using Maximum Likelihood estimation

Likelihood of a Graph

- Given a set of parameters, the likelihood of the graph is computed by

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$

Nodes, V 



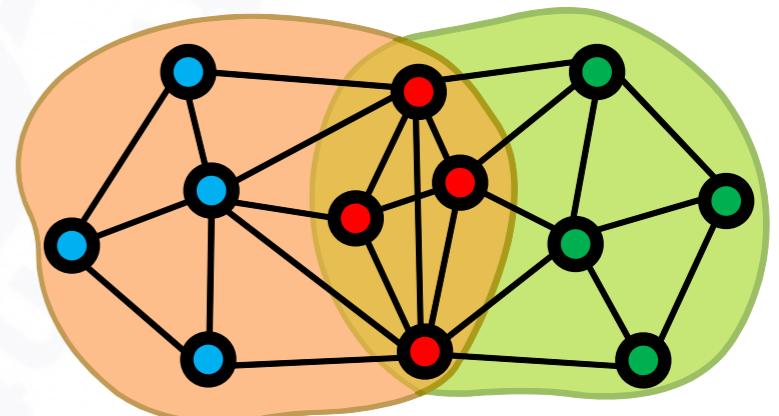
Estimate parameters $\hat{\theta}$ by using Maximum Likelihood estimation

Likelihood of a Graph

- Given a set of parameters, the likelihood of the graph is computed by

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$

Nodes, V ● ● ● ● ● ● ● ● ● ● ● ●



Estimate parameters $\hat{\theta}$ by using Maximum Likelihood estimation

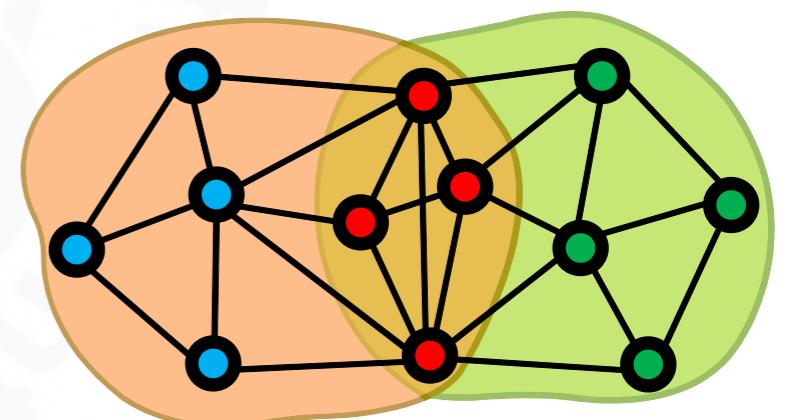
Likelihood of a Graph

- Given a set of parameters, the likelihood of the graph is computed by

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$

Communities, C p_A A p_B B

Nodes, V ● ● ● ● ● ● ● ● ● ●

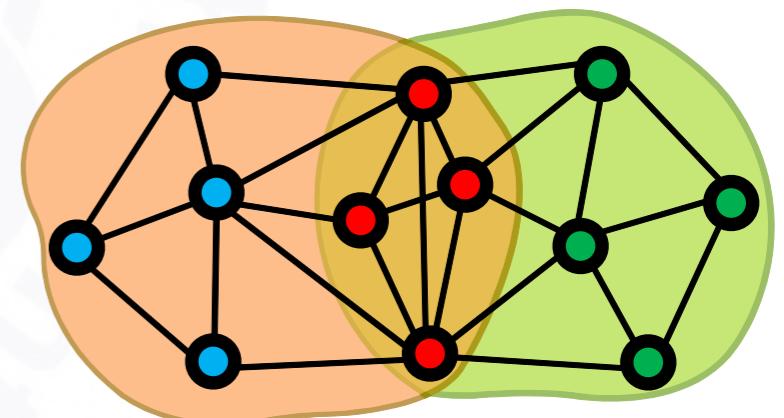
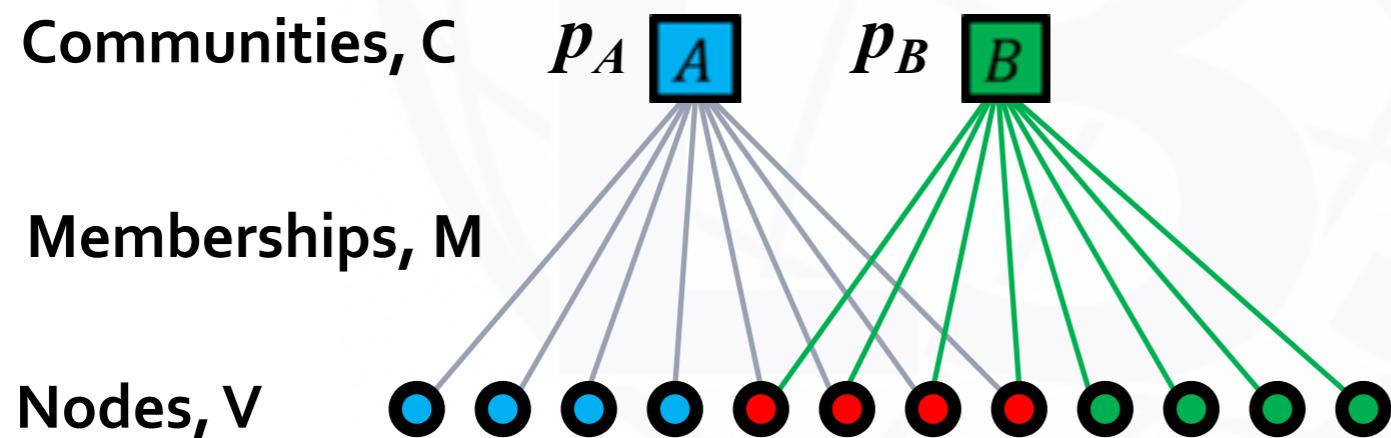


Estimate parameters $\hat{\theta}$ by using Maximum Likelihood estimation

Likelihood of a Graph

- Given a set of parameters, the likelihood of the graph is computed by

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$



Estimate parameters $\hat{\theta}$ by using Maximum Likelihood estimation

Parameter Estimation : MLE

Given observed data D

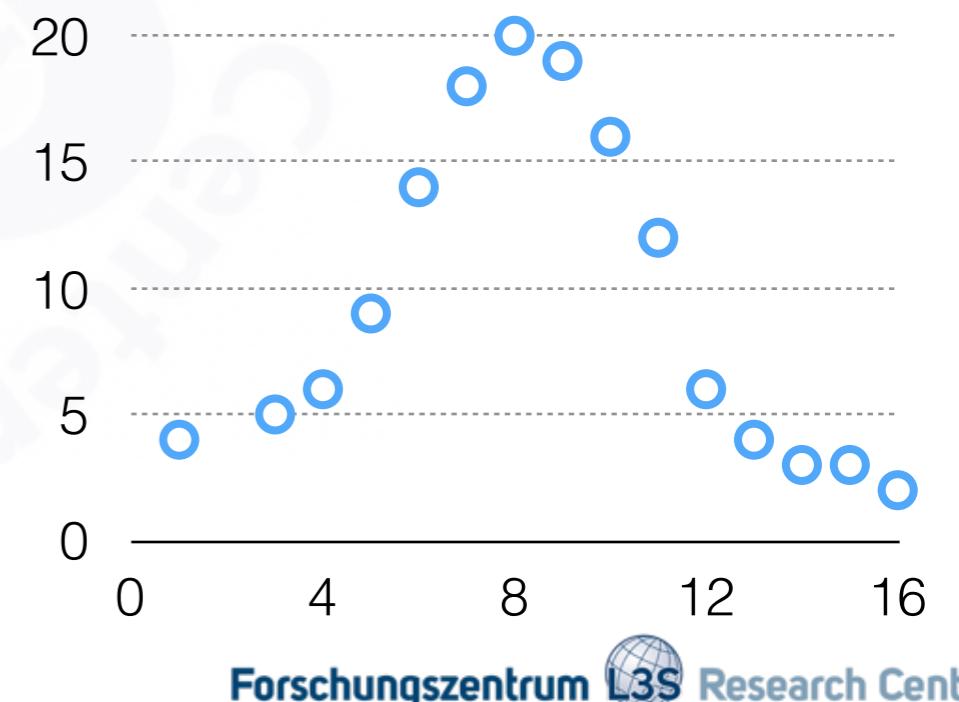
Assume that you have a probability density func. $p(x)$ with parameters θ which generates it

parameter estimation problem

What are the value of the parameters which best explain my data ?

Say you feel that D is generated from a Normal distribution

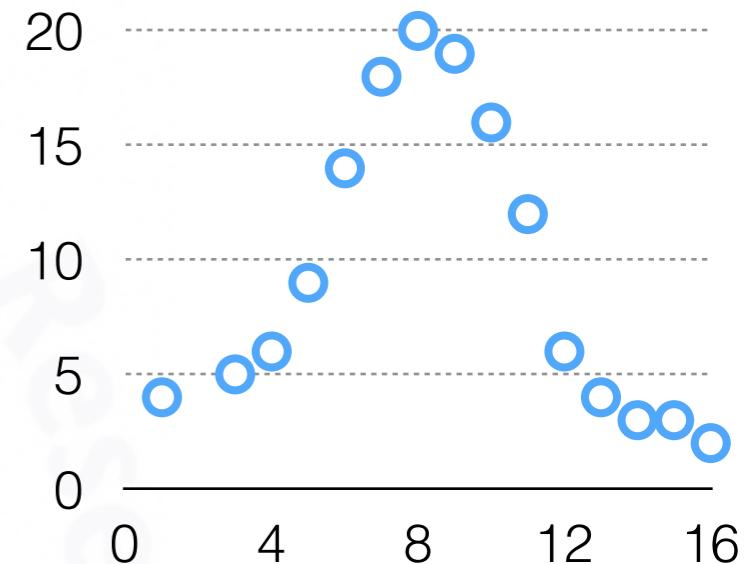
$$\theta = N(\mu, \sigma)$$



Maximum Likelihood Estimator (MLE)

Say you feel that D is generated from a Normal distribution

$$\theta = N(\mu, \sigma)$$



The **likelihood** of observing this **independent and identically distributed sample** is

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = p(x \mid \theta)$$

parameters to be estimated $\hat{\theta}$

observed data x

MLE chooses $\hat{\theta}$ maximizes this likelihood of generating the observed data

Maximum Likelihood Estimator (MLE)

- Desired probability distribution is the one that makes the observed data “most likely,” which means that one must seek the value of the parameter vector $\hat{\theta}$ that maximizes the likelihood function $\mathcal{L}(\hat{\theta}|x)$
- For computational convenience, the MLE estimate is obtained by maximizing the log-likelihood function $\ln(\mathcal{L}(\hat{\theta}|x))$

$$\ln(\mathcal{L}(\hat{\theta}|x)) = \sum_{x_i} \ln p(x_i | \theta)$$

- If log-likelihood is differentiable find $\hat{\theta}$ partial derivatives

Maximum Likelihood Estimator (MLE)

- Desired probability distribution is the one that makes the observed data “most likely,” which means that one must seek the value of the parameter vector $\hat{\theta}$ that maximizes the likelihood function $\mathcal{L}(\hat{\theta}|x)$

$$p(x = (x_1, x_2, \dots, x_n) | \theta) = p(x_1|\theta) \cdot p(x_2|\theta) \dots p(x_n|\theta)$$


- For computational convenience, the MLE estimate is obtained by maximizing the log-likelihood function $\ln(\mathcal{L}(\hat{\theta}|x))$

$$\ln(\mathcal{L}(\hat{\theta}|x)) = \sum_{x_i} \ln p(x_i | \theta)$$

- If log-likelihood is differentiable find $\hat{\theta}$ partial derivatives

Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \longrightarrow \quad \ln(\mathcal{L}(\hat{\theta}|x)) = \sum_{x_i} \ln p(x_i \mid \theta)$$

Likelihood → *Log Likelihood*

- Find maxima by using partial derivatives i.e.

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

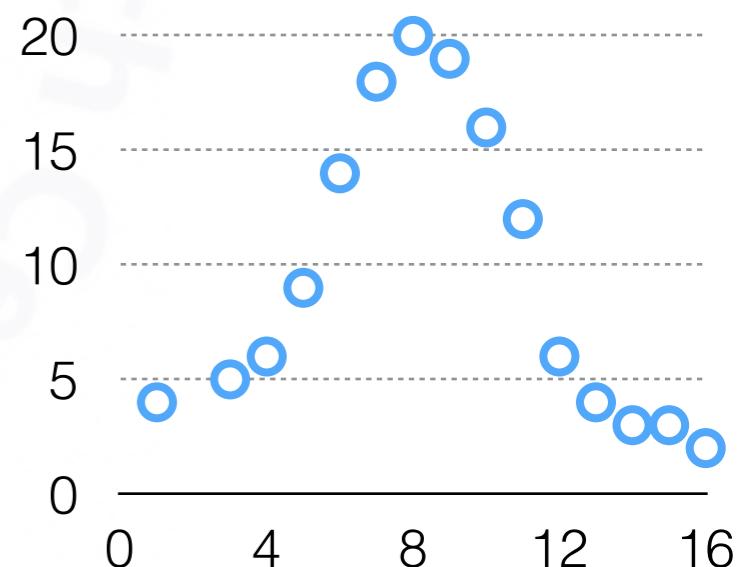
Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \xrightarrow{\text{Likelihood}} \quad \ln(\mathcal{L}(\hat{\theta} \mid x)) = \sum_{x_i} \ln p(x_i \mid \theta) \quad \xrightarrow{\text{Log Likelihood}}$$

- Find maxima by using partial derivatives i.e.

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

- Example :



Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \longrightarrow \quad \ln(\mathcal{L}(\hat{\theta} \mid x)) = \sum_{x_i} \ln p(x_i \mid \theta)$$

Likelihood → **log Likelihood**

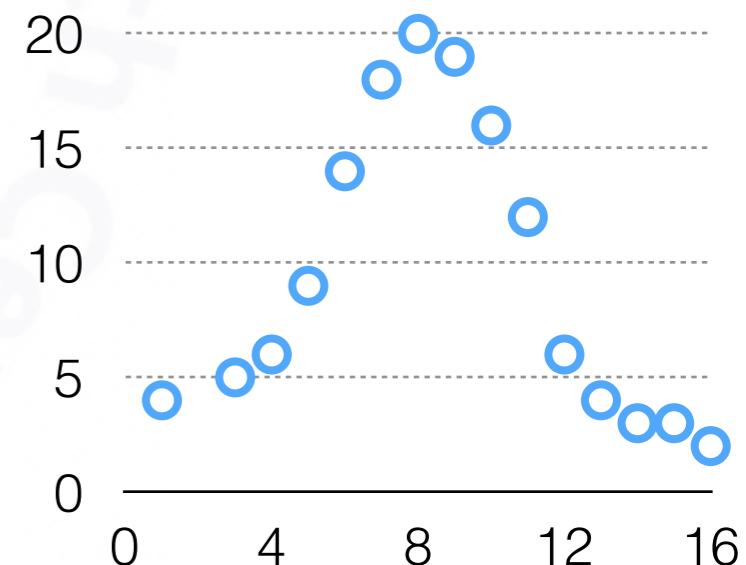
- Find maxima by using partial derivatives i.e.

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

- Example :

$$p(x_i \mid (\mu, \sigma)) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

generative model gaussian



Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \longrightarrow \quad \ln(\mathcal{L}(\hat{\theta} \mid x)) = \sum_{x_i} \ln p(x_i \mid \theta)$$

Likelihood → **log Likelihood**

- Find maxima by using partial derivatives i.e.

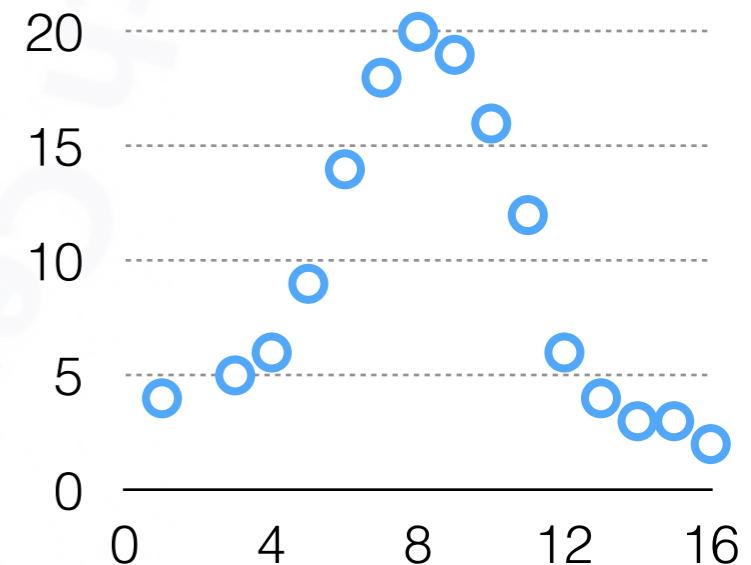
$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

- Example :

$$p(x_i \mid (\mu, \sigma)) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

generative model ↓ **gaussian**

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \mu} = 0$$



Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \xrightarrow{\text{likelihood}} \quad \ln(\mathcal{L}(\hat{\theta} \mid x)) = \sum_{x_i} \ln p(x_i \mid \theta) \quad \xrightarrow{\text{log likelihood}}$$

- Find maxima by using partial derivatives i.e.

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

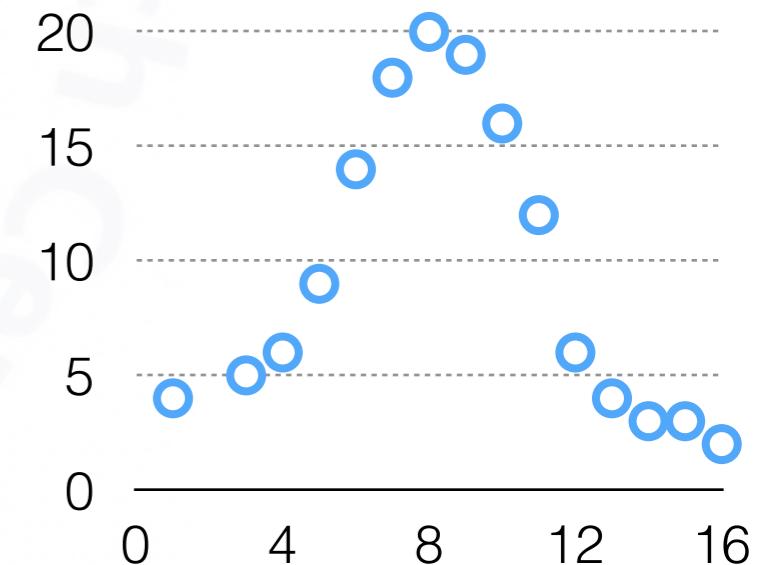
- Example :

generative model

$$p(x_i \mid (\mu, \sigma)) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

gaussian

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \mu} = 0 \quad \frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \sigma} = 0$$



Maximum Likelihood Estimator (MLE)

$$\mathcal{L}(\theta \mid x = (x_1, x_2, \dots, x_n)) = \prod_{x_i} p(x_i \mid \theta) \quad \xrightarrow{\text{likelihood}} \quad \ln(\mathcal{L}(\hat{\theta} \mid x)) = \sum_{x_i} \ln p(x_i \mid \theta) \quad \xrightarrow{\text{log likelihood}}$$

- Find maxima by using partial derivatives i.e.

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \theta_i} = 0$$

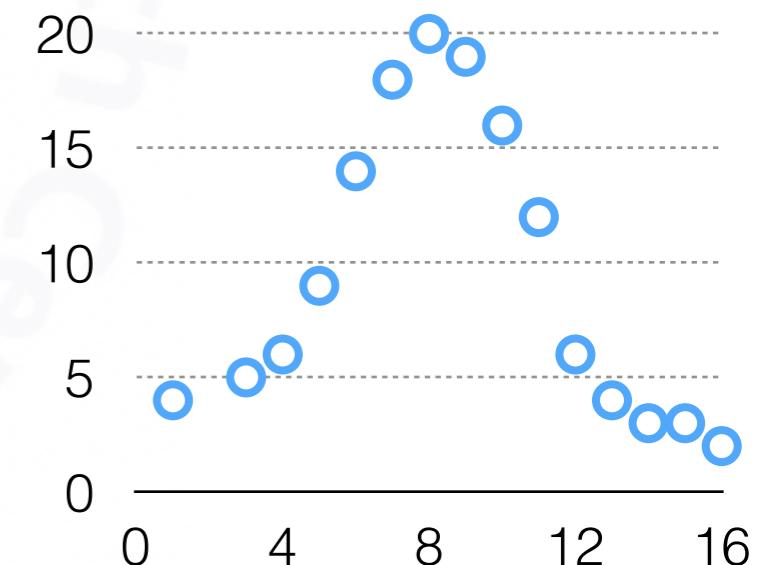
- Example :

generative model

$$p(x_i \mid (\mu, \sigma)) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

gaussian

$$\frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \mu} = 0 \quad \frac{\partial \ln \mathcal{L}(\theta \mid x)}{\partial \sigma} = 0$$



- What are the estimated parameters of the gaussian ?

MLE for Graphs

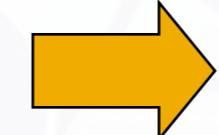
- How do we do MLE for graphs?
 - Model generates a **probabilistic adjacency matrix**
 - We then flip all the entries of the probabilistic matrix to obtain the **binary adjacency matrix A**
-
- The likelihood of AGM generating graph G:

MLE for Graphs

- How do we do MLE for graphs?
 - Model generates a **probabilistic adjacency matrix**
 - We then flip all the entries of the probabilistic matrix to obtain the **binary adjacency matrix A**

For every pair of nodes u, v AGM gives the prob. p_{uv} of them being linked

0	0.10	0.10	0.04
0.10	0	0.02	0.06
0.10	0.02	0	0.06
0.04	0.06	0.06	0

Flip biased coins 

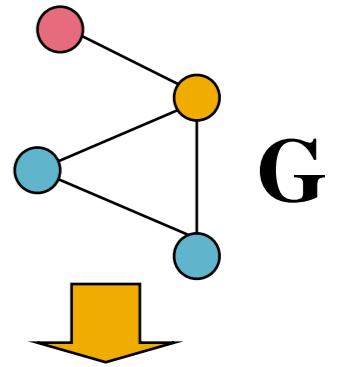
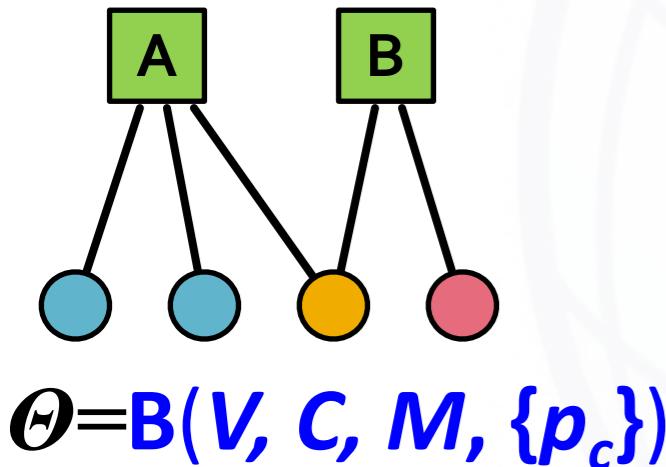
A	0	1	0	0
0	1	0	1	1
1	0	1	0	1
0	1	0	1	0
0	1	1	1	0

- The likelihood of AGM generating graph G:

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u,v) \prod_{(u,v) \notin E} (1 - P(u,v))$$

Graphs: Likelihood $P(G|\Theta)$

- Given graph $G(V,E)$ and Θ , we calculate likelihood that Θ generated G : $P(G|\Theta)$

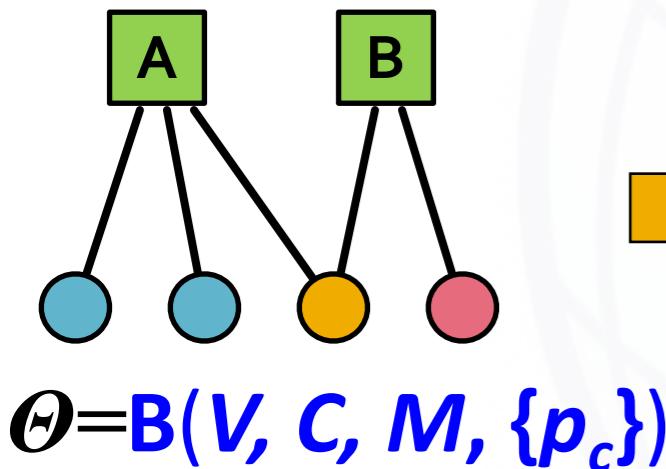


0	1	1	0
1	0	1	0
1	1	0	1
0	0	1	0

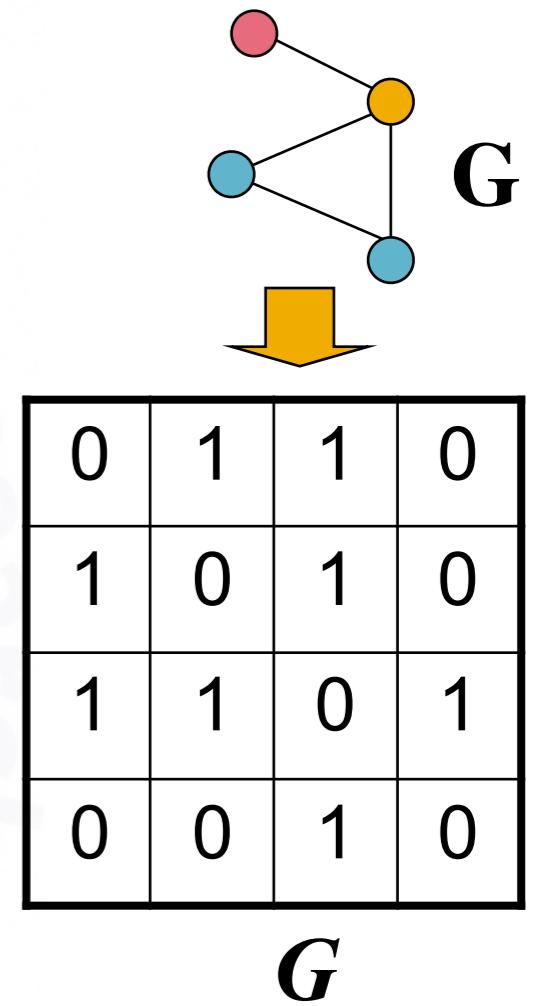
G

Graphs: Likelihood $P(G|\Theta)$

- Given graph $G(V,E)$ and Θ , we calculate likelihood that Θ generated G : $P(G|\Theta)$

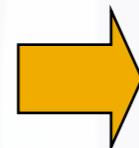
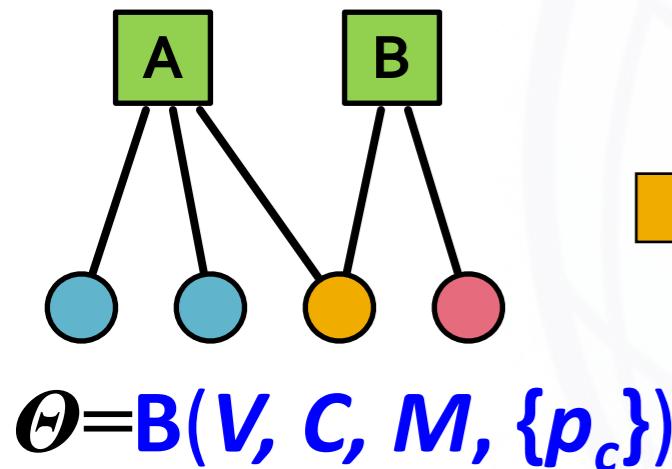


0	0.9	0.9	0
0.9	0	0.9	0
0.9	0.9	0	0.9
0	0	0.9	0

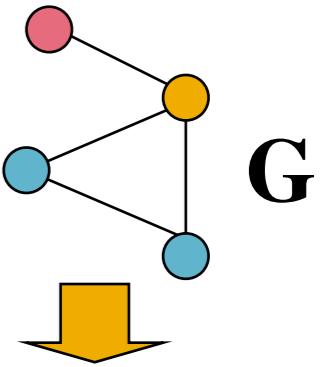


Graphs: Likelihood $P(G|\Theta)$

- Given graph $G(V,E)$ and Θ , we calculate likelihood that Θ generated G : $P(G|\Theta)$



0	0.9	0.9	0
0.9	0	0.9	0
0.9	0.9	0	0.9
0	0	0.9	0



G

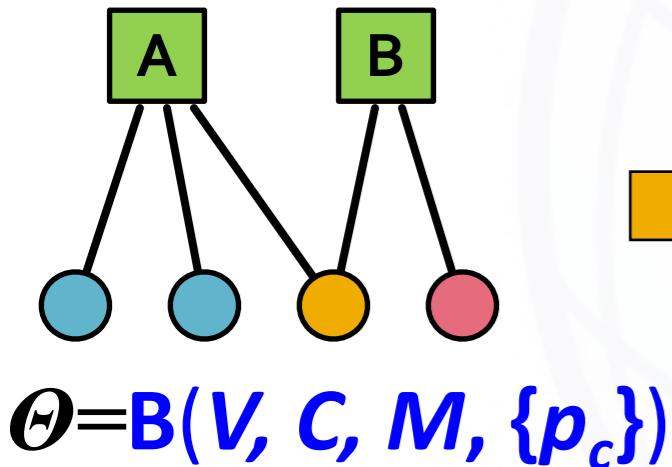
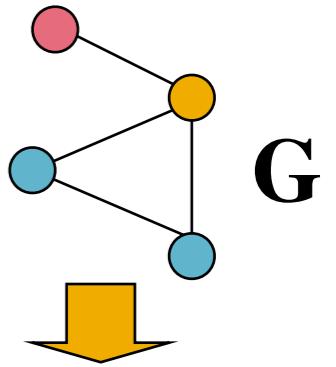
0	1	1	0
1	0	1	0
1	1	0	1
0	0	1	0

G

$P(G|\Theta)$

Graphs: Likelihood $P(G|\Theta)$

- Given graph $G(V,E)$ and Θ , we calculate likelihood that Θ generated G : $P(G|\Theta)$



0	0.9	0.9	0
0.9	0	0.9	0
0.9	0.9	0	0.9
0	0	0.9	0

0	1	1	0
1	0	1	0
1	1	0	1
0	0	1	0

G

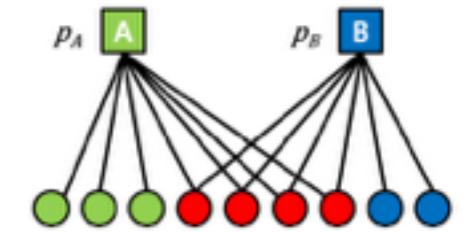
$P(G|\Theta)$

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u, v) \prod_{(u,v) \notin E} (1 - P(u, v))$$

MLE for Graphs

- **Our goal:** Find $\Theta = B(V, C, M, \{p_C\})$ such that:

G

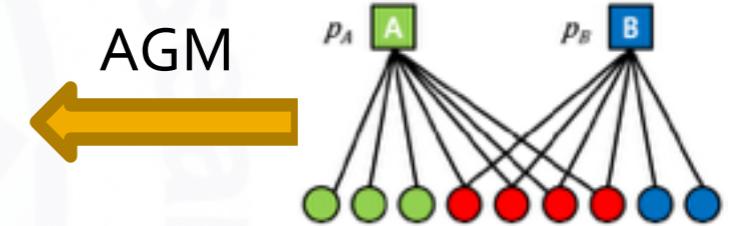


-
- **How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?**

MLE for Graphs

- Our goal: Find $\Theta = B(V, C, M, \{p_C\})$ such that:

G

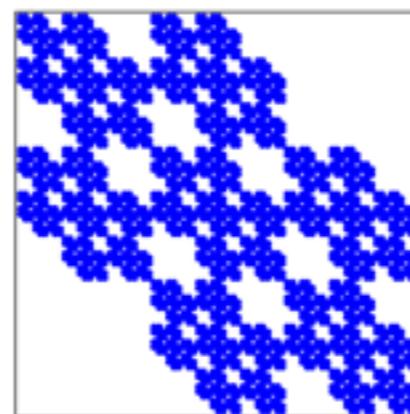


-
- How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?

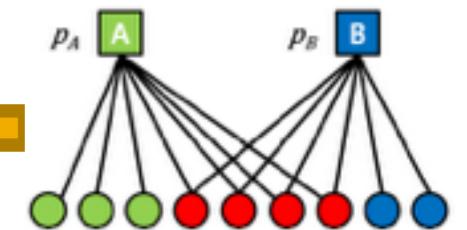
MLE for Graphs

- Our goal: Find $\Theta = B(V, C, M, \{p_C\})$ such that:

G



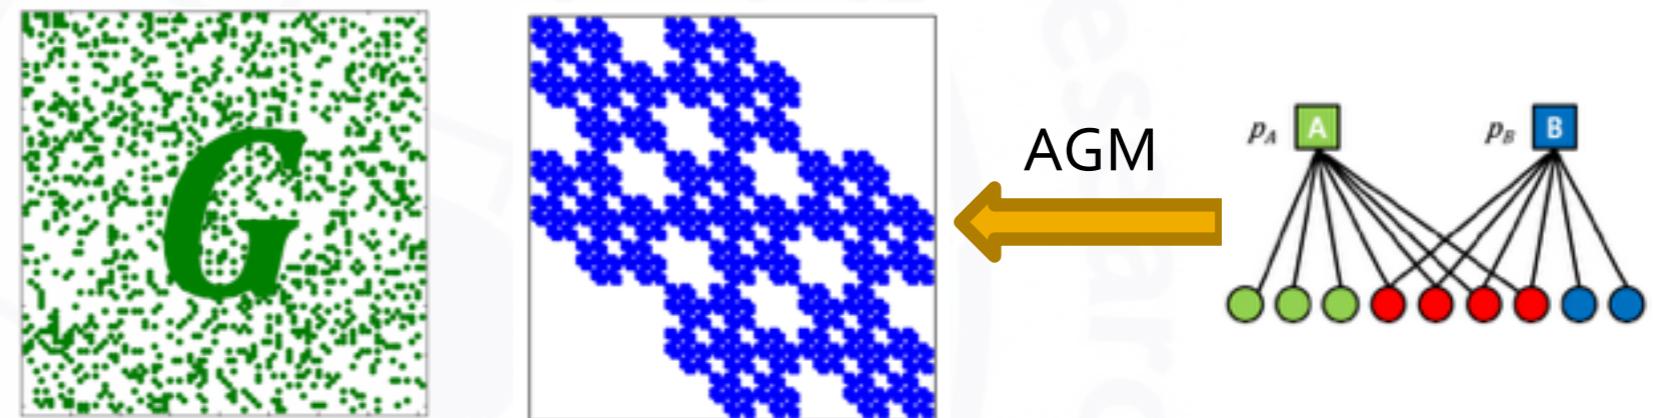
AGM



-
- How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?

MLE for Graphs

- Our goal: Find $\Theta = B(V, C, M, \{p_C\})$ such that:



-
- How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?

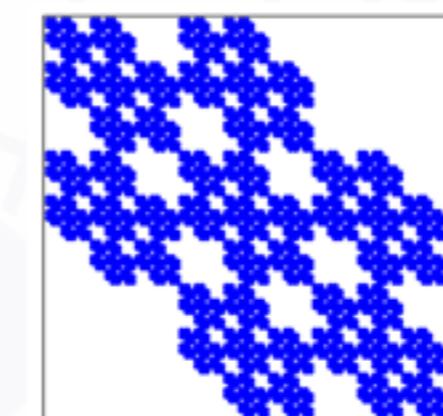
MLE for Graphs

- Our goal: Find $\Theta = B(V, C, M, \{p_C\})$ such that:

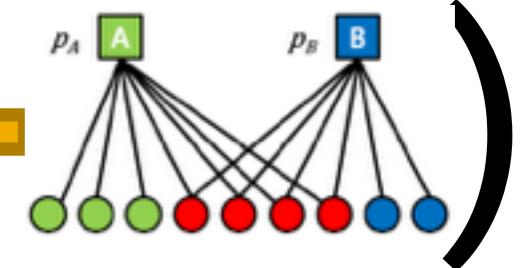
P(



|



AGM



- How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?

MLE for Graphs

- **Our goal:** Find $\Theta = B(V, C, M, \{p_C\})$ such that:

$$\arg \max_{\Theta} P(G | \Theta)$$

-
- **How do we find $B(V, C, M, \{p_C\})$ that maximizes the likelihood?**

$$P(G | \Theta) = \prod_{(u,v) \in E} P(u, v) \prod_{(u,v) \notin E} (1 - P(u, v))$$

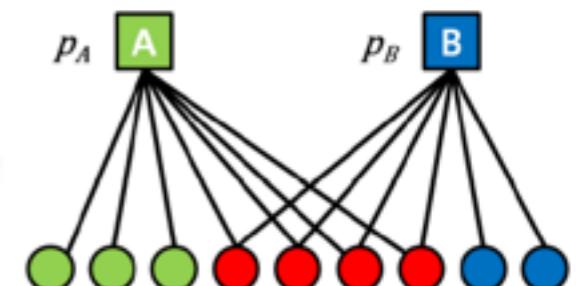
MLE for AGM

- Our goal is to find $B(V, C, M, \{p_C\})$ such that:

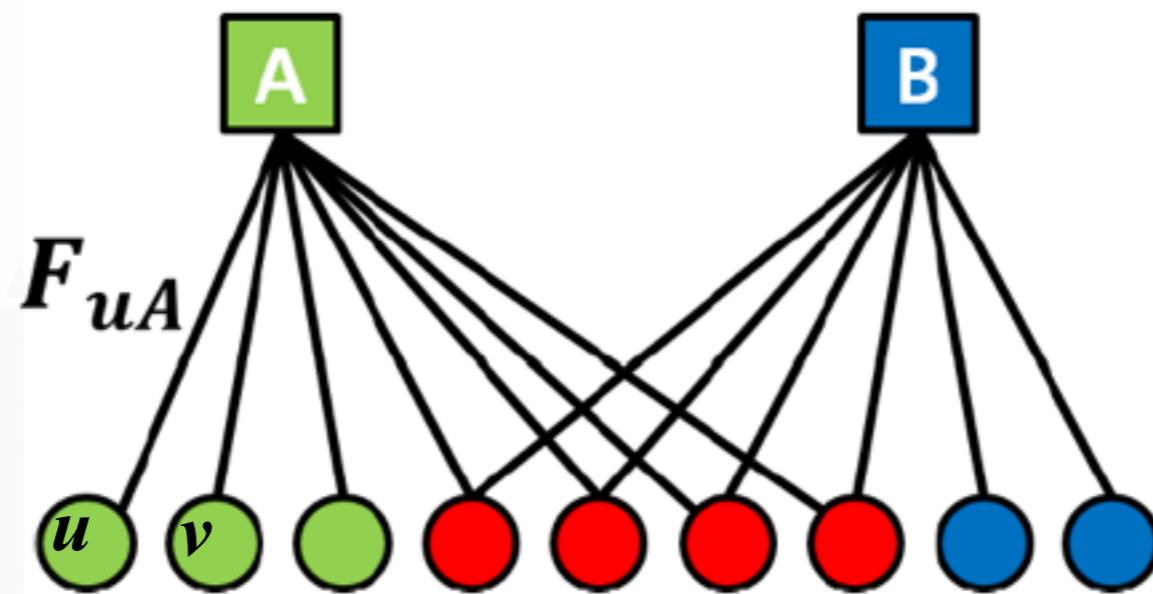
$$\arg \max_{B(V, C, M, \{p_C\})} \prod_{u, v \in E} P(u, v) \prod_{uv \notin E} (1 - P(u, v))$$

- Problem: Finding B means finding the bipartite affiliation network.

- There is no nice way to do this.
- Fitting $B(V, C, M, \{p_C\})$ is too hard, let's change the model (so it is easier to fit)!

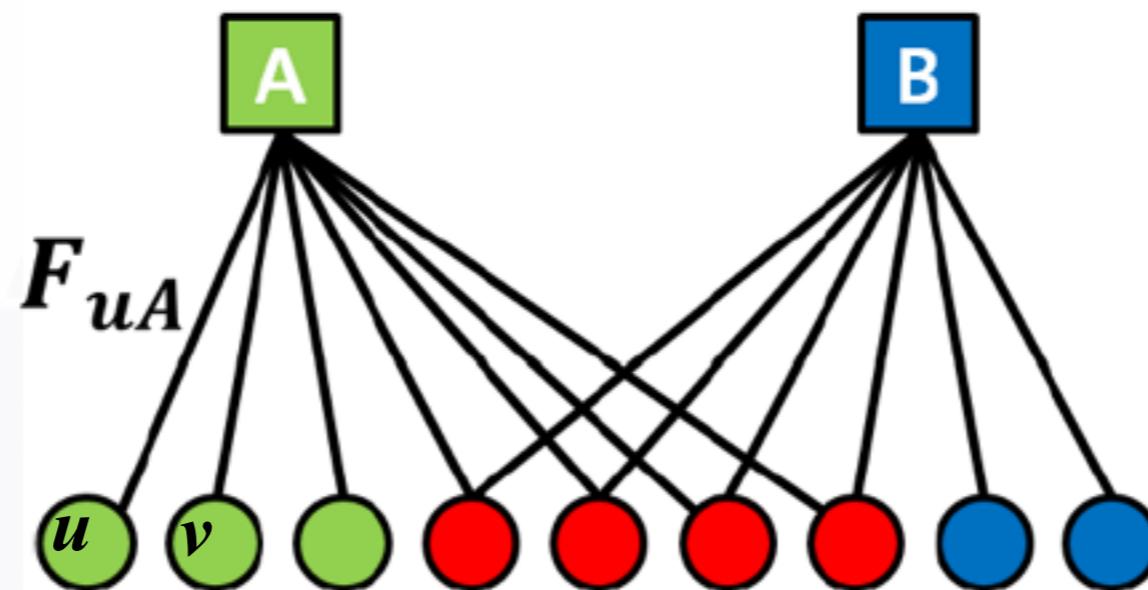


From AGM to BigCLAM



From AGM to BigCLAM

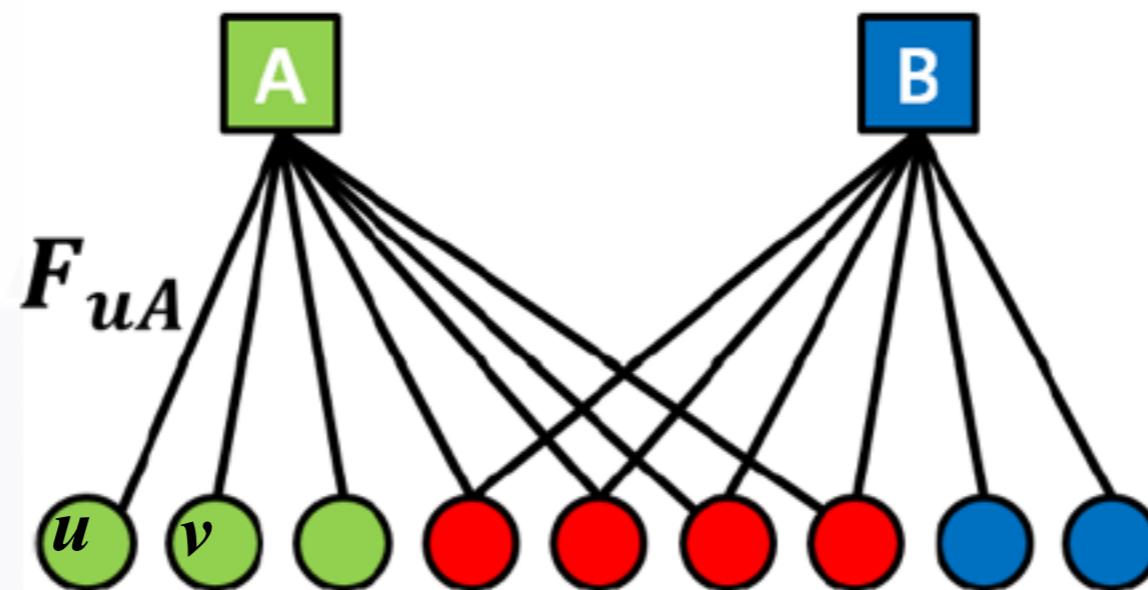
■ Relaxation: Memberships have strengths



- F_{uA} : The membership strength of node u to community A ($F_{uA} = 0$: no membership)
- **Each community A links nodes independently:**
$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

From AGM to BigCLAM

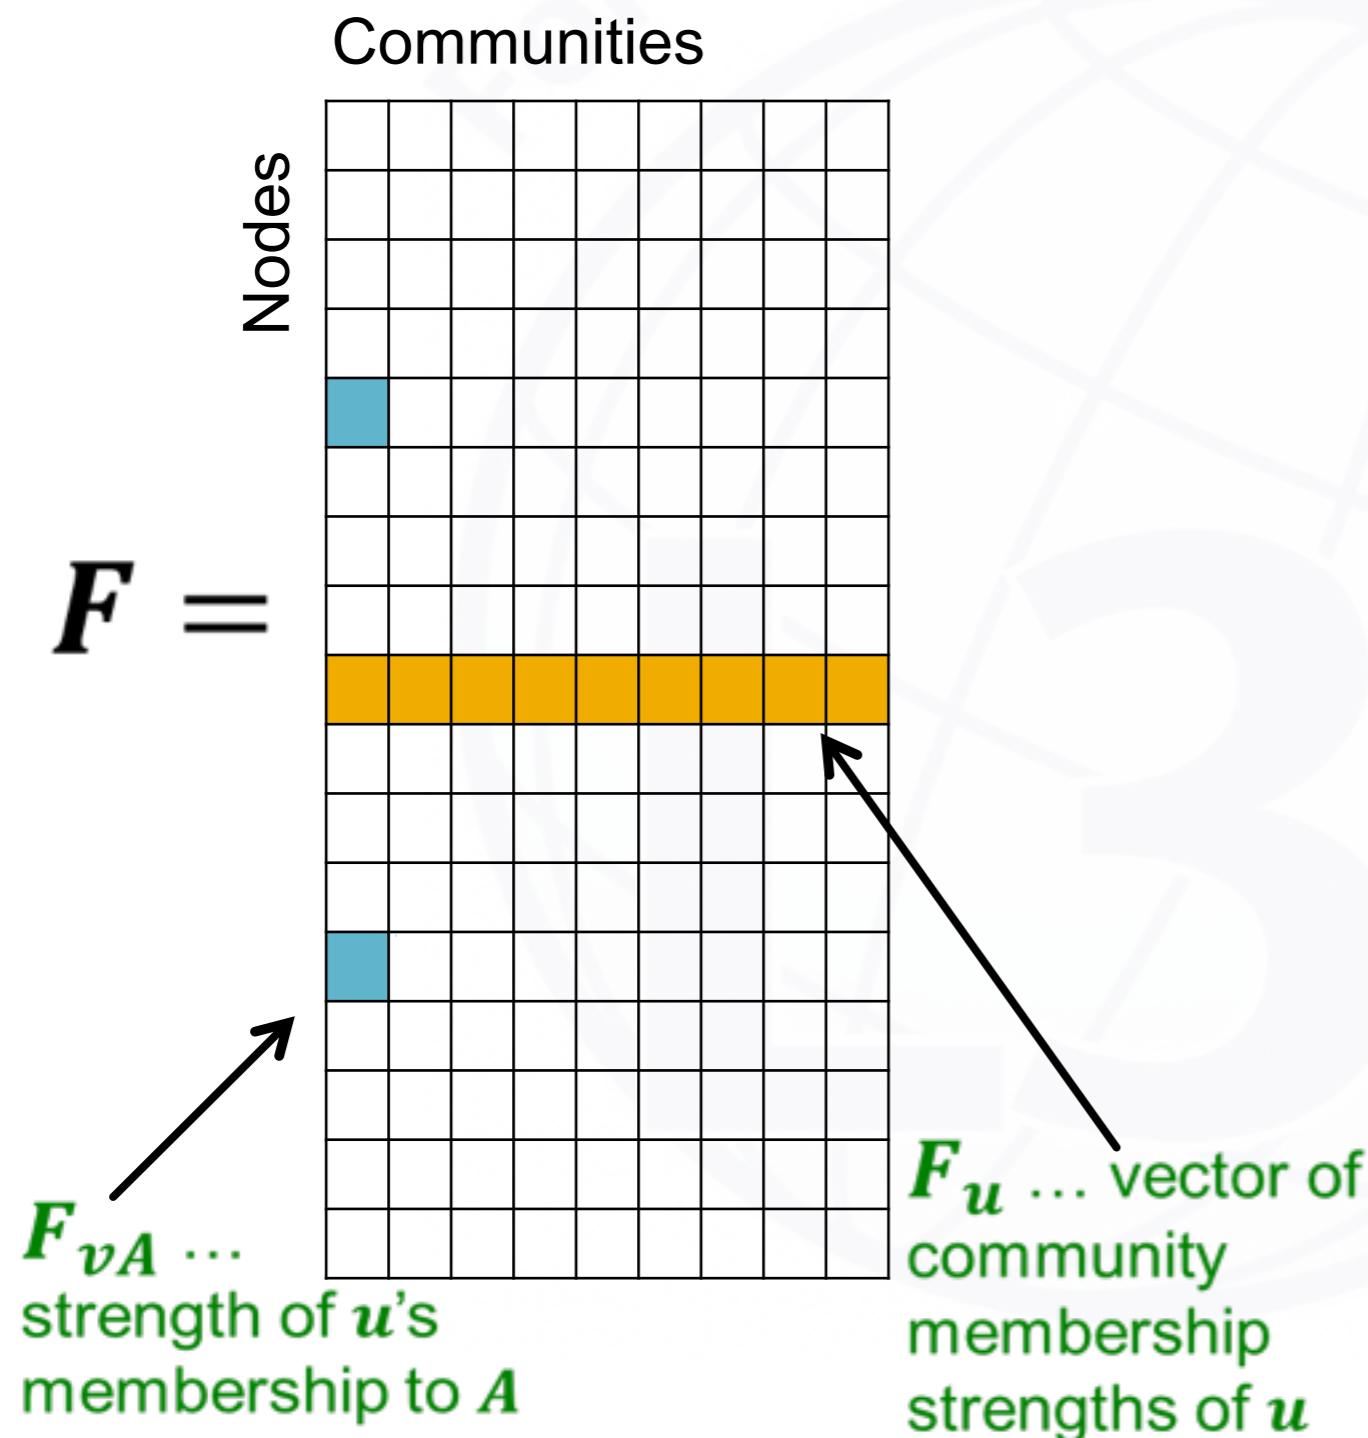
■ Relaxation: Memberships have strengths



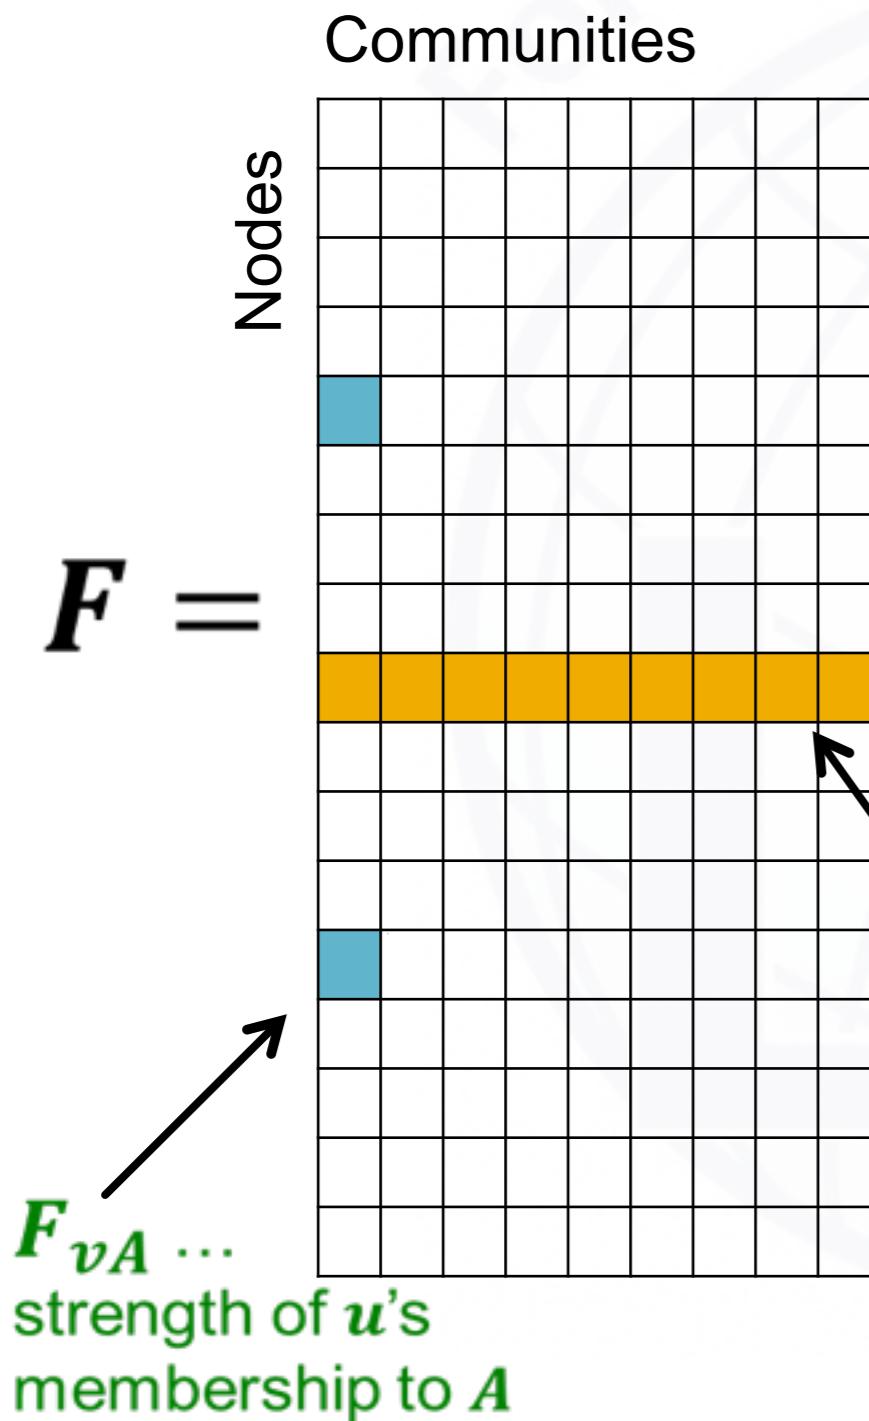
- F_{uA} : The membership strength of node u to community A ($F_{uA} = 0$: no membership)
- **Each community A links nodes independently:**
$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

Assume if a pair of nodes are strongly related to a community then they have a high probability of a connection

■ Community membership strength matrix F



■ Community membership strength matrix F



- $P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$
 - Probability of connection is proportional to the product of strengths
 - Notice: If one node doesn't belong to the community ($F_{uC} = 0$) then $P(u, v) = 0$
- Prob. that at least one common community C links the nodes:
 - $P(u, v) = 1 - \prod_C (1 - P_C(u, v))$

F_u ... vector of community membership strengths of u

From AGM to BigCLAM

From AGM to BigCLAM

- Community \mathbf{A} links nodes u, v independently:

$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

- Then prob. at least one common C links them:

$$\begin{aligned} P(u, v) &= 1 - \prod_C (1 - P_C(u, v)) \\ &= 1 - \exp(-\sum_C F_{uC} \cdot F_{vC}) \\ &= 1 - \exp(-F_u \cdot F_v^T) \end{aligned}$$

- Example F matrix:

From AGM to BigCLAM

- Community \mathbf{A} links nodes u, v independently:

$$P_A(u, v) = 1 - \exp(-\mathbf{F}_{uA} \cdot \mathbf{F}_{vA})$$

- Then prob. at least one common C links them:

$$\begin{aligned} P(u, v) &= 1 - \prod_C (1 - P_C(u, v)) \\ &= 1 - \exp(-\sum_C \mathbf{F}_{uC} \cdot \mathbf{F}_{vC}) \\ &= 1 - \exp(-\mathbf{F}_u \cdot \mathbf{F}_v^T) \end{aligned}$$

- Example F matrix:

$\mathbf{F}_u:$	0	1.2	0	0.2
-----------------	---	-----	---	-----

$\mathbf{F}_v:$	0.5	0	0	0.8
-----------------	-----	---	---	-----

$\mathbf{F}_w:$	0	1.8	1	0
-----------------	---	-----	---	---

Node community
membership strengths

From AGM to BigCLAM

- Community \mathbf{A} links nodes u, v independently:

$$P_A(u, v) = 1 - \exp(-F_{uA} \cdot F_{vA})$$

- Then prob. at least one common C links them:

$$\begin{aligned} P(u, v) &= 1 - \prod_C (1 - P_C(u, v)) \\ &= 1 - \exp(-\sum_C F_{uC} \cdot F_{vC}) \\ &= 1 - \exp(-F_u \cdot F_v^T) \end{aligned}$$

- Example F matrix:

$F_u:$	0	1.2	0	0.2
$F_v:$	0.5	0	0	0.8
$F_w:$	0	1.8	1	0

Node community
membership strengths

Then: $F_u \cdot F_v^T = 0.16$

And: $P(u, v) = 1 - \exp(-0.16) = 0.14$

But: $P(u, w) = 0.88$

$P(v, w) = 0$

BigCLAM: How to find F

- **Task:** Given a network $G(V, E)$, estimate F

- Find F that maximizes the likelihood:

$$\arg \max_F \prod_{(u,v) \in E} P(u, v) \prod_{(u,v) \notin E} (1 - P(u, v))$$

- where: $P(u, v) = 1 - \exp(-F_u \cdot F_v^T)$
 - Many times we take the logarithm of the likelihood, and call it log-likelihood: $l(F) = \log P(G|F)$
- **Goal:** Find F that maximizes $l(F)$:

BigCLAM: How to find F

- **Task:** Given a network $G(V, E)$, estimate F
 - Find F that maximizes the likelihood:

$$\arg \max_F \prod_{(u,v) \in E} P(u, v) \prod_{(u,v) \notin E} (1 - P(u, v))$$

- - where: $P(u, v) = 1 - \exp(-F_u \cdot F_v^T)$
 - Many times we take the logarithm of the likelihood, and call it log-likelihood: $l(F) = \log P(G|F)$
- **Goal:** Find F that maximizes $l(F)$:

$$l(F) = \sum_{(u,v) \in E} \log(1 - \exp(-F_u F_v^T)) - \sum_{(u,v) \notin E} F_u F_v^T$$

BigCLAM: V1.0

$$l(F_u) = \sum_{v \in \mathcal{N}(u)} \log(1 - \exp(-F_u F_v^T)) - \sum_{v \notin \mathcal{N}(u)} F_u F_v^T$$

$\mathcal{N}(u)$.. Set out
outgoing neighbors

BigCLAM: V1.0

$$l(F_u) = \sum_{v \in \mathcal{N}(u)} \log(1 - \exp(-F_u F_v^T)) - \sum_{v \notin \mathcal{N}(u)} F_u F_v^T$$

- Compute gradient of a single row F_u , of F :

$$\nabla l(F_u) = \sum_{v \in \mathcal{N}(u)} F_v \frac{\exp(-F_u F_v^T)}{1 - \exp(-F_u F_v^T)} - \sum_{v \notin \mathcal{N}(u)} F_v$$

- Coordinate gradient ascent:

- Iterate over the rows of F :

- Compute gradient $\nabla l(F_u)$ of row u (while keeping others fixed)
 - Update the row F_u : $F_u \leftarrow F_u + \eta \nabla l(F_u)$
 - Project F_u back to a non-negative vector: If $F_{uC} < 0$: $F_{uC} = 0$

$\mathcal{N}(u)$.. Set out
outgoing neighbors

- This is slow! Computing $\nabla l(F_u)$ takes linear time!

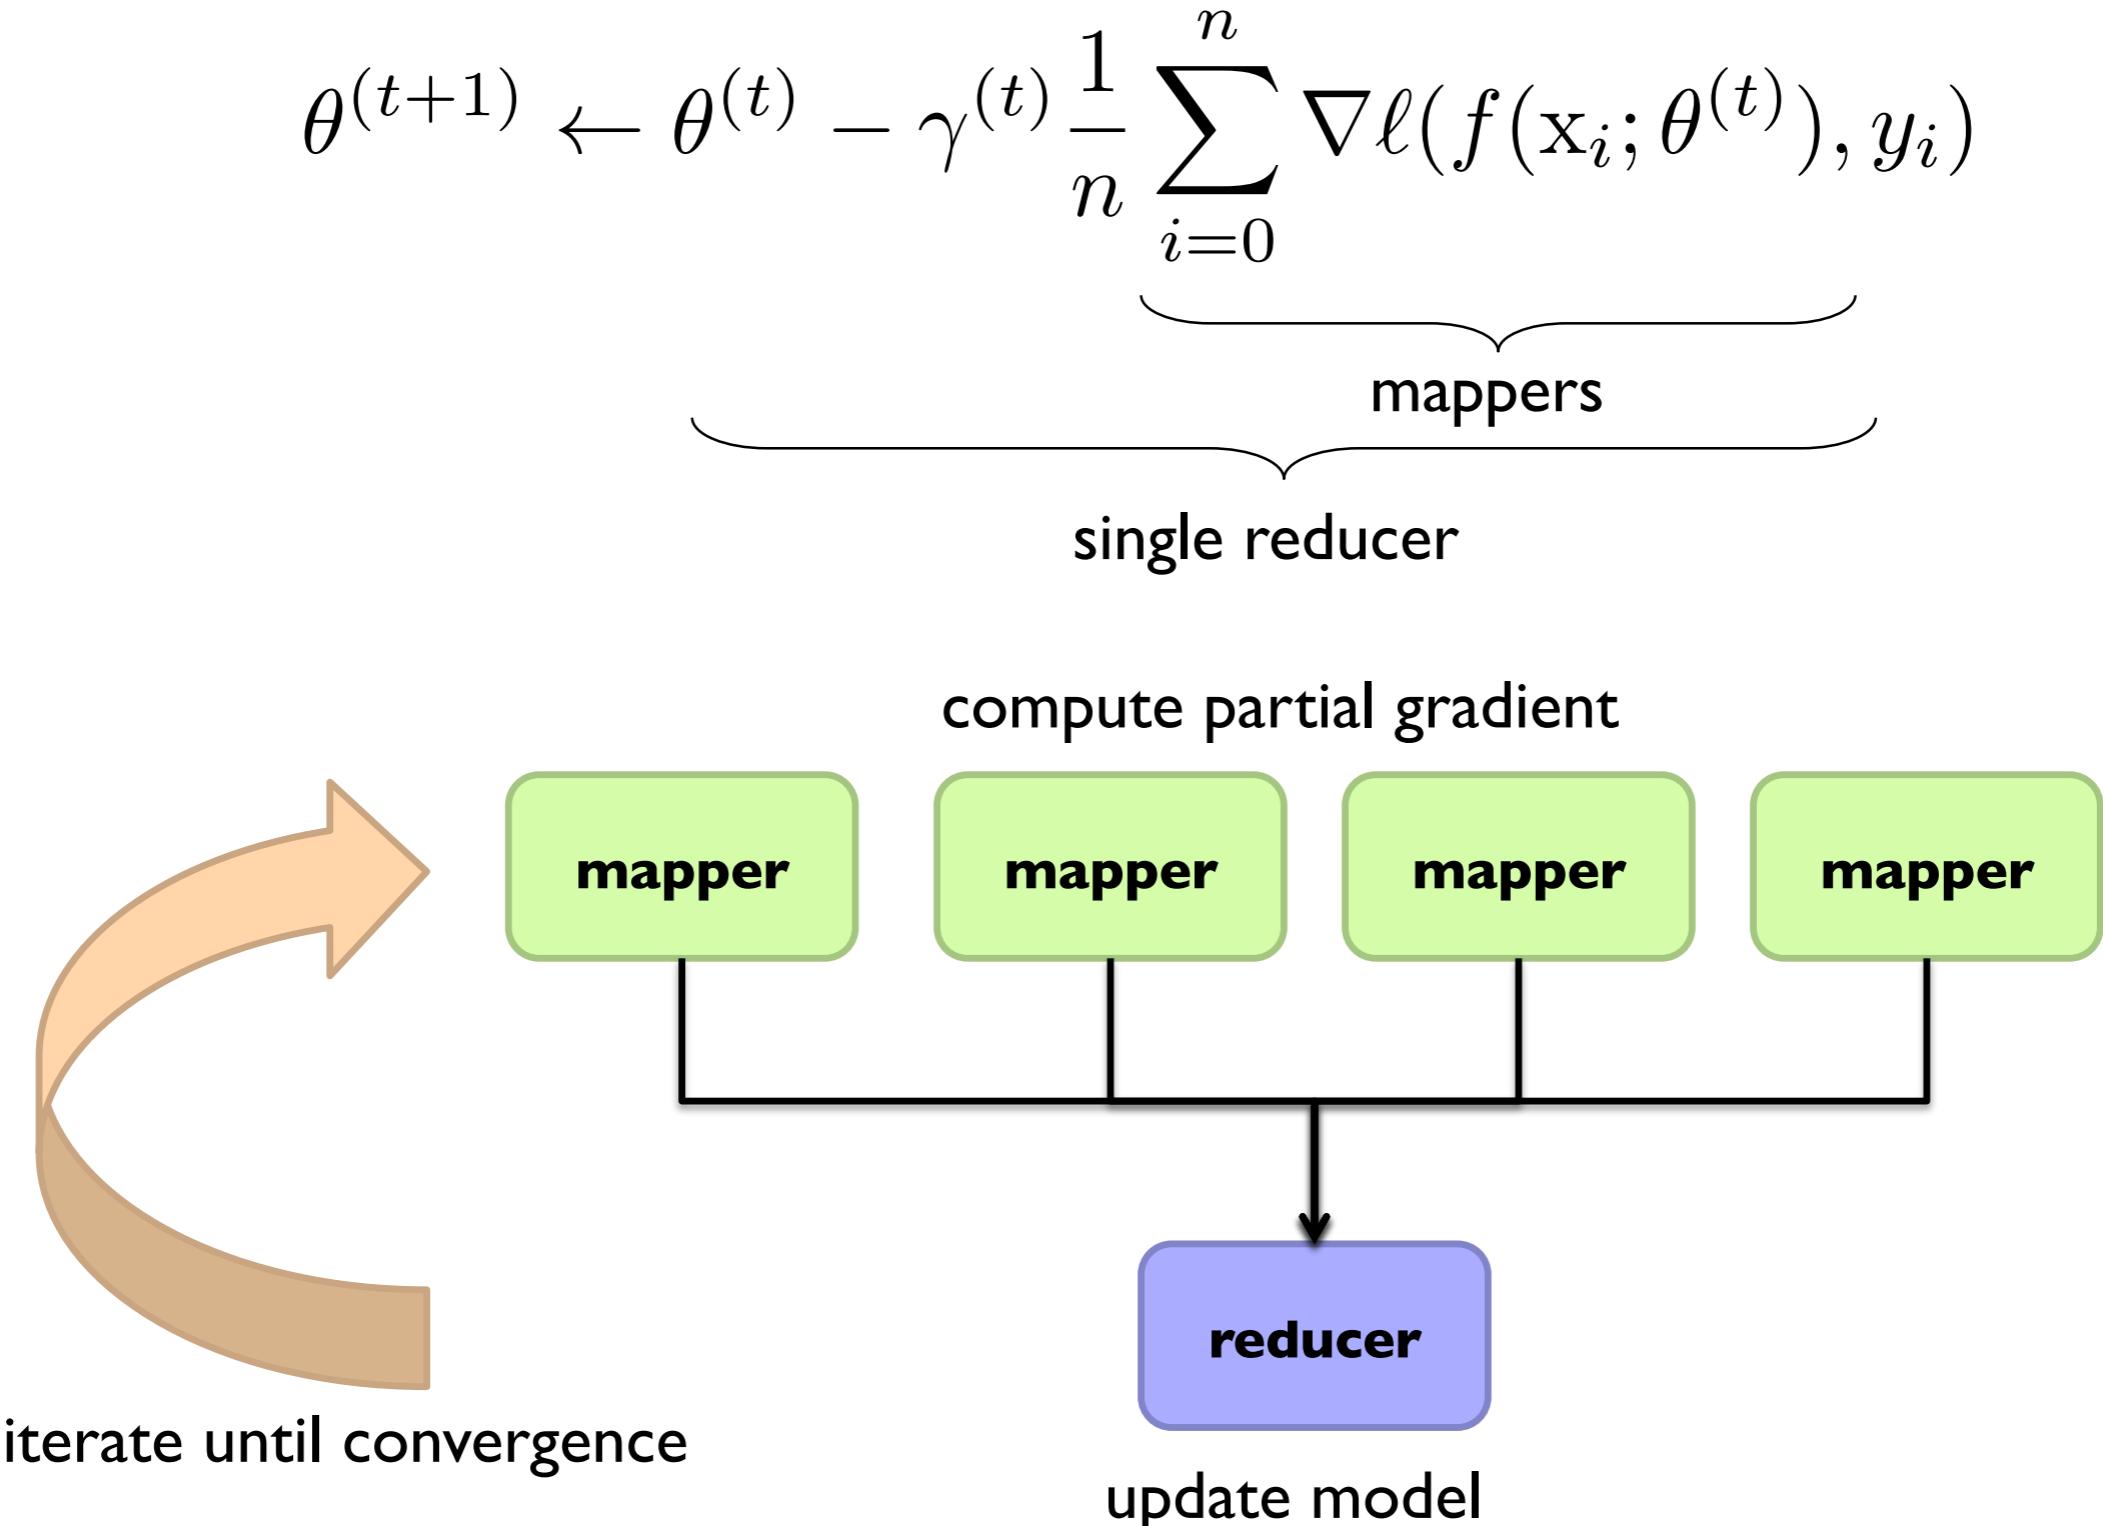
BigCLAM: V2.0

- **However, we notice:**

$$\sum_{v \notin \mathcal{N}(u)} F_v = \left(\sum_v F_v - F_u - \sum_{v \in \mathcal{N}(u)} F_v \right)$$

- - We cache $\sum_v F_v$
 - So, computing $\sum_{v \notin \mathcal{N}(u)} F_v$ now takes **linear time** in the degree $|\mathcal{N}(u)|$ of u
 - In networks degree of a node is much smaller to the total number of nodes in the network, so this is a significant speedup!

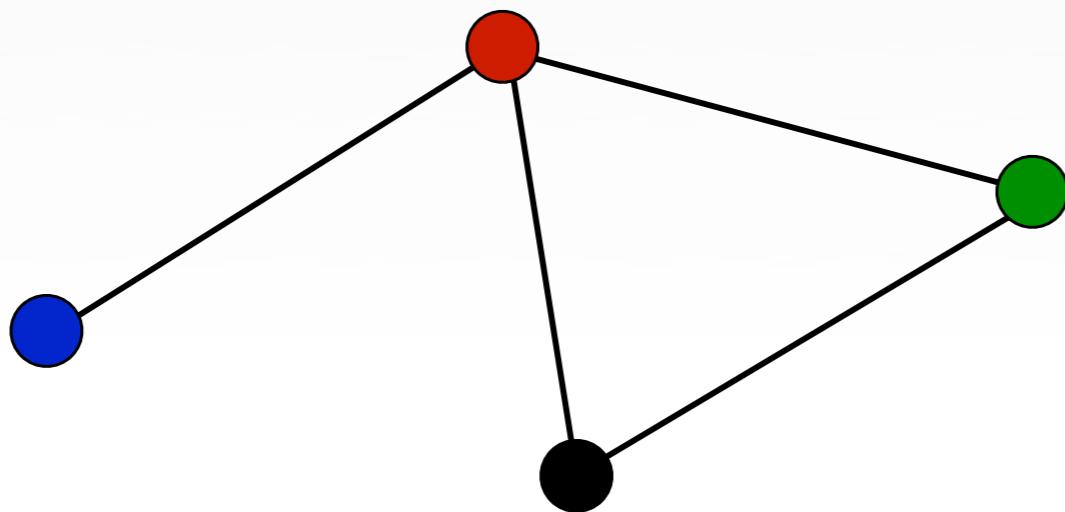
Batch Gradient Descent in Mapreduce



Counting Triangles

- Counting triangles in a graph is an important operation for determining the **degree of connectedness** in a neighbourhood
- Clustering Coefficient : fraction of a node's neighbors who are neighbors themselves
- $\text{cc}(v) = \frac{|\{(u, w) \in E | u \in \Gamma(v) \wedge w \in \Gamma(v)\}|}{\binom{d_v}{2}} = \frac{\#\Delta s \text{ incident on } v}{\binom{d_v}{2}}$

Clustering Co-efficient



$$cc(\text{Blue}) = \text{N/A}$$

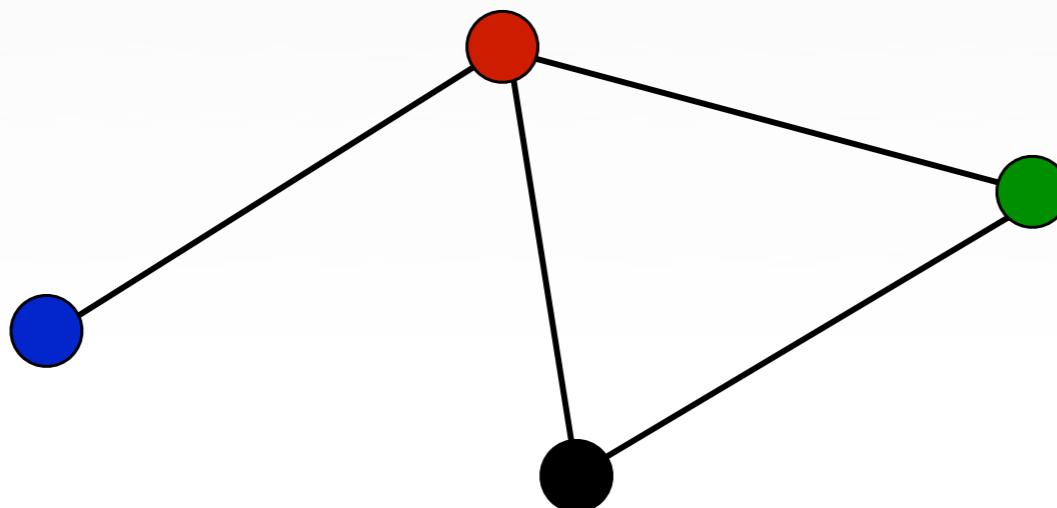
$$cc(\text{Red}) = 1/3$$

$$cc(\text{Green}) = 1$$

$$cc(\text{Black}) = 1$$

$$\frac{\#\Delta s \text{ incident on } v}{\binom{d_v}{2}}$$

Clustering Co-efficient



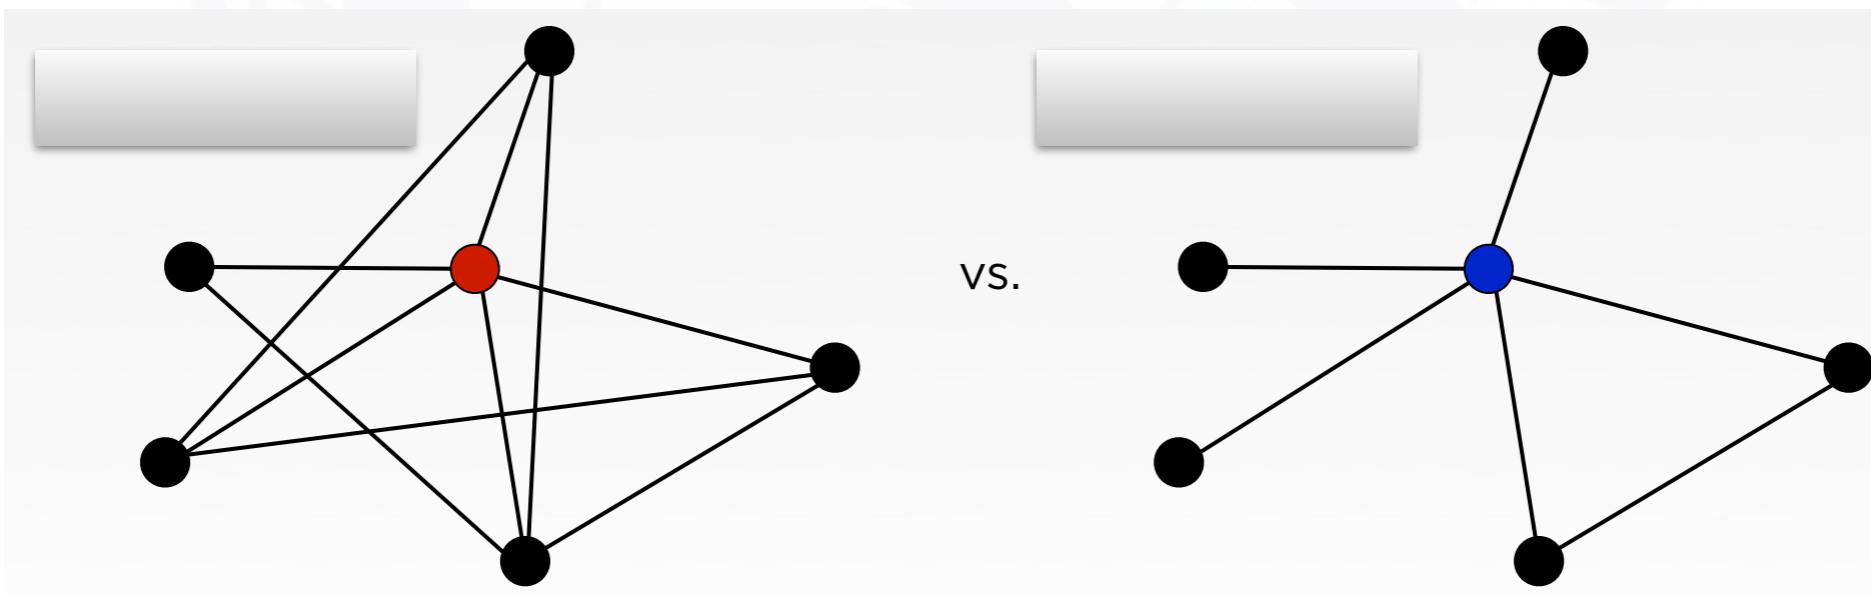
$$cc(\text{blue}) = \text{N/A}$$

$$cc(\text{red}) = 1/3$$

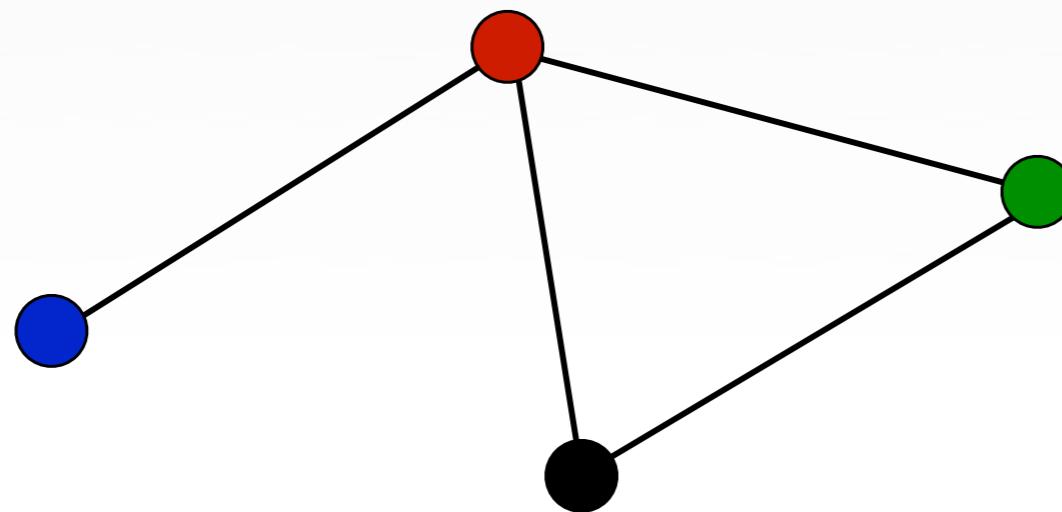
$$cc(\text{green}) = 1$$

$$cc(\text{black}) = 1$$

$$\frac{\#\Delta s \text{ incident on } v}{\binom{d_v}{2}}$$



Clustering Co-efficient



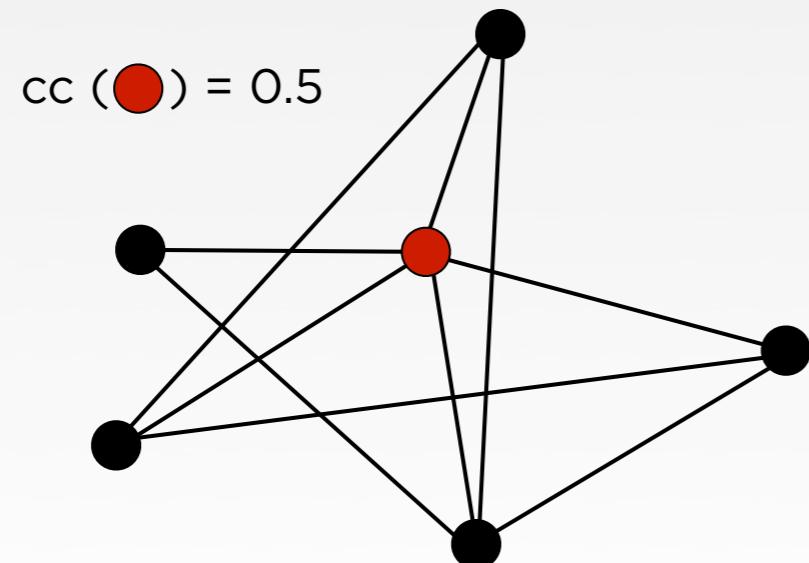
$$cc(\text{blue}) = \text{N/A}$$

$$cc(\text{red}) = 1/3$$

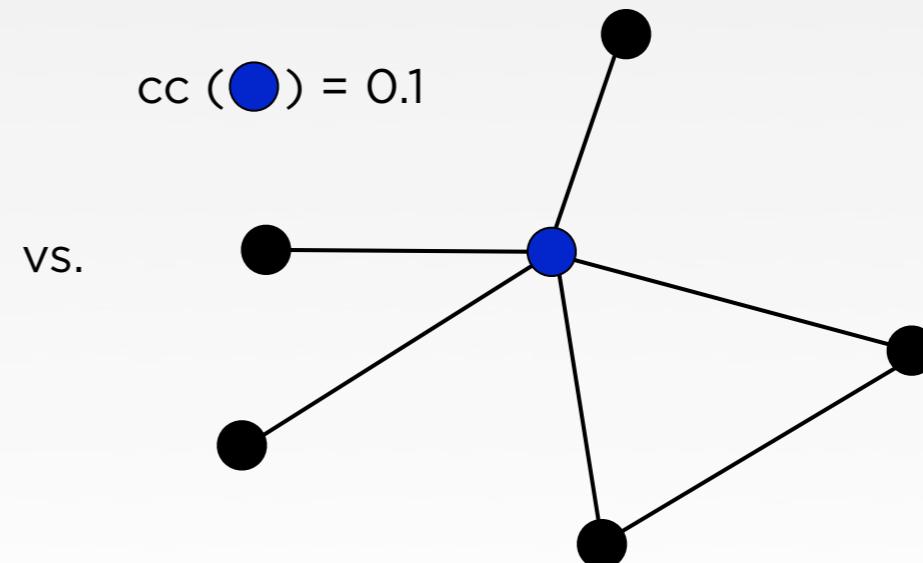
$$cc(\text{green}) = 1$$

$$cc(\text{black}) = 1$$

$$\frac{\#\Delta s \text{ incident on } v}{\binom{d_v}{2}}$$



$$cc(\text{red}) = 0.5$$

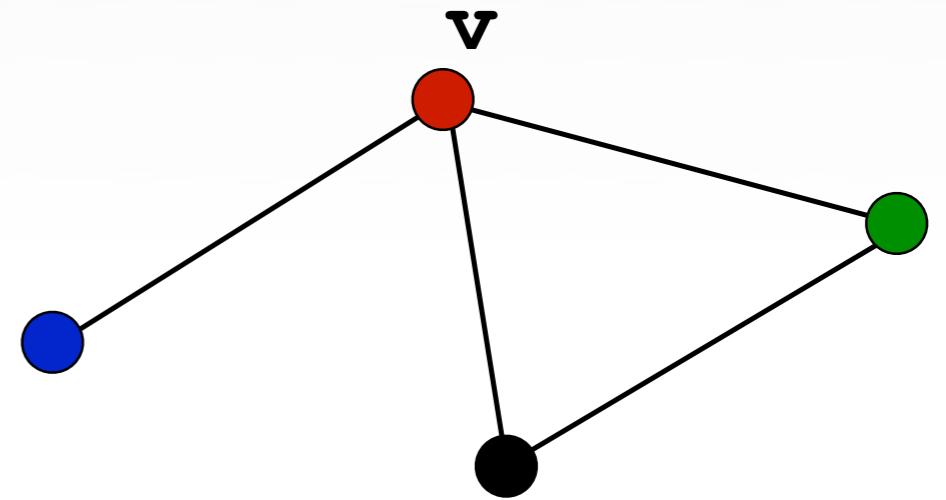


$$cc(\text{blue}) = 0.1$$

vs.

Naive Method

```
foreach v in V
    foreach u,w in Adjacency(v)
        if (u,w) in E
            Triangles[v]++
```



What is its complexity ?

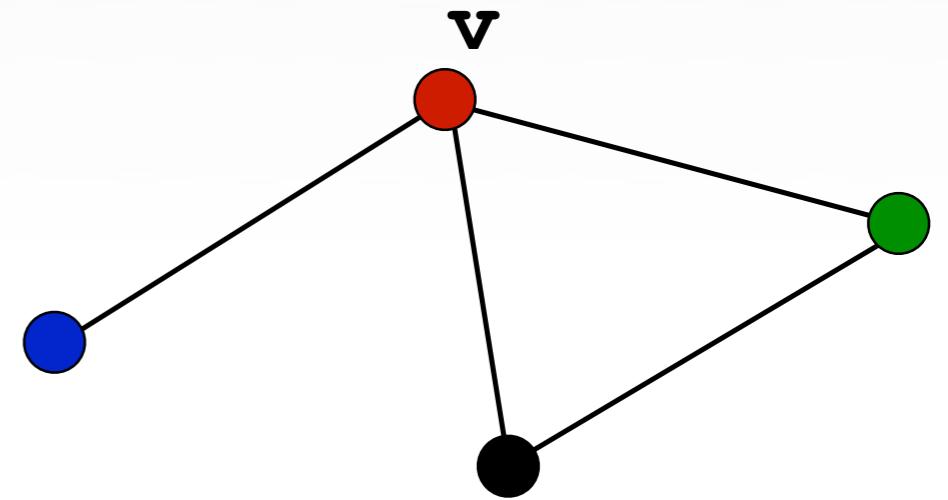
Possible graph inputs: (a) adjacency list , (b) edge list

Preprocessing : Create a hashset of (u,w) pairs
for constant lookup

Preprocessing : Create adjacency list **O(m)**

Naive Method - Complexity

```
foreach v in V
    foreach u,w in Adjacency(v)
        if (u,w) in E
            Triangles[v]++
```



Running time $\sum_{v \in V} d_v^2$

For sparse graphs : quadratic even if a few nodes have large out degree

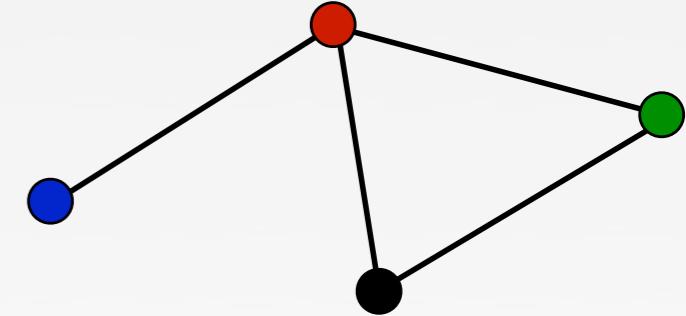
Parallel Method

2-stage Mapreduce

- Map 1: For each v send $(v, \Gamma(v))$ to single machine.
- Reduce 1: Input: $\langle v; \Gamma(v) \rangle$

Output: all 2 paths $\langle (v_1, v_2); u \rangle$ where $v_1, v_2 \in \Gamma(u)$

$(\bullet, \bullet); \bullet$ $(\bullet, \bullet); \bullet$ $(\bullet, \bullet); \bullet$



- Map 2: Send $\langle (v_1, v_2); u \rangle$ and $\langle (v_1, v_2); \$ \rangle$ for $(v_1, v_2) \in E$ to same machine.

- Reduce 2: input: $\langle (v, w); u_1, u_2, \dots, u_k, \$? \rangle$

Output: if $\$$ part of the input, then: $u_i = u_i + 1/3$

$(\bullet, \bullet); \bullet, \$ \rightarrow \bullet + 1/3 \quad \bullet + 1/3 \quad \bullet + 1/3$
 $(\bullet, \bullet); \bullet \rightarrow$

check edges in parallel

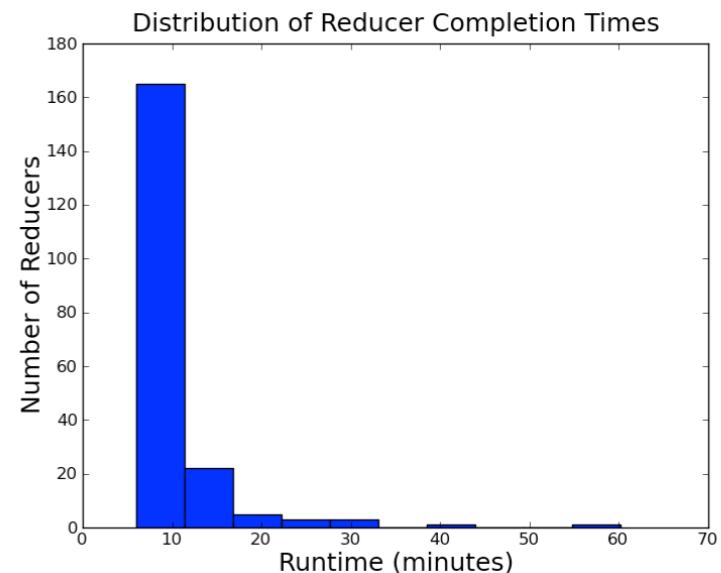
Parallel Method

2-stage Mapreduce

- Map 2: Send $\langle(v_1, v_2); u\rangle$ and $\langle(v_1, v_2); \$\rangle$ for $(v_1, v_2) \in E$ to same machine.
- Reduce 2: input: $\langle(v, w); u_1, u_2, \dots, u_k, \$?\rangle$
Output: if $\$$ part of the input, then: $u_i = u_i + 1/3$

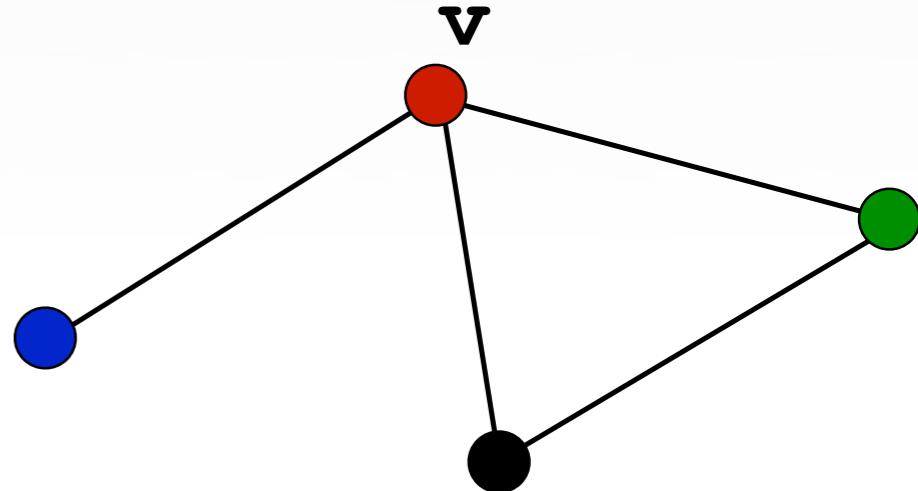
$(\bullet, \bullet); \$ \rightarrow \bullet + 1/3$ $\bullet + 1/3$ $\bullet + 1/3$

$(\bullet, \bullet); \bullet \rightarrow$



- Each triangle counted thrice
- How do we ensure that we count a triangle only once ?
- How do we handle skew ?

Heavy Hitters



- Heavy hitter nodes: Nodes with $\text{out-degree} > m^{1/2}$
- Heavy hitter triangles are triangles with all vertices as heavy hitters
 - How many heavy hitter nodes can there be ?
 - Partition nodes into heavy hitters and non heavy hitters
 - Treat each of these partitions separately

Algorithm

- For Heavy hitter (HH) nodes:
 - Consider all possible combinations of HH nodes and verify if they form triangles
 - Complexity: $O(m^{3/2})$
- For non HH nodes:
 - only consider neighbours with higher degrees or numerical order
 - Complexity: $O(m^{3/2})$
 - Each edge considered only once $O(m)$ and max out-degree is $O(m^{1/2})$

Algorithm

- For Heavy hitter (HH) nodes:

- Consider all possible combinations of HH nodes and verify if they form triangles

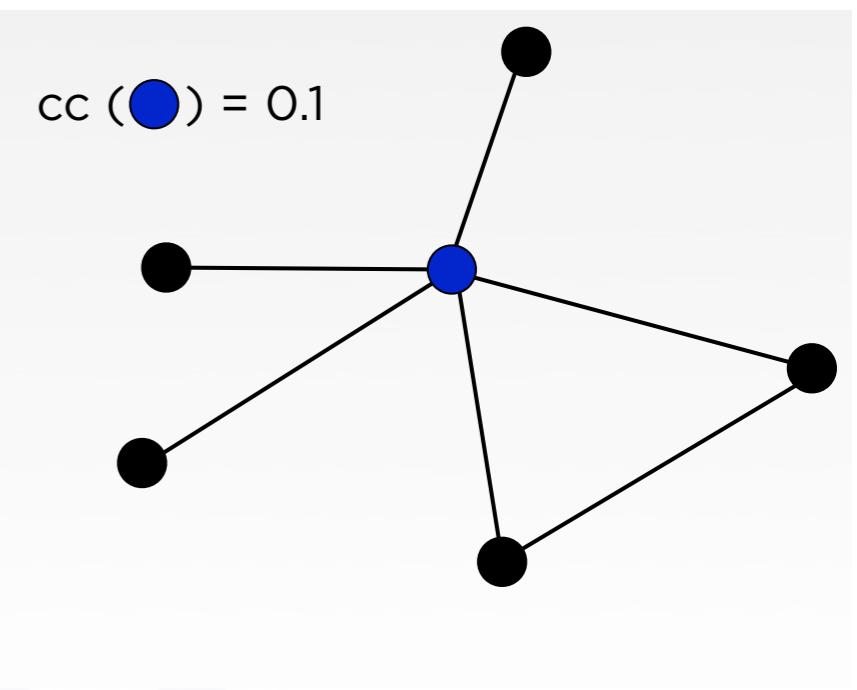
- Complexity: $O(m^{3/2})$

- For non HH nodes:

- only consider neighbours with higher degrees or numerical order

- Complexity: $O(m^{3/2})$

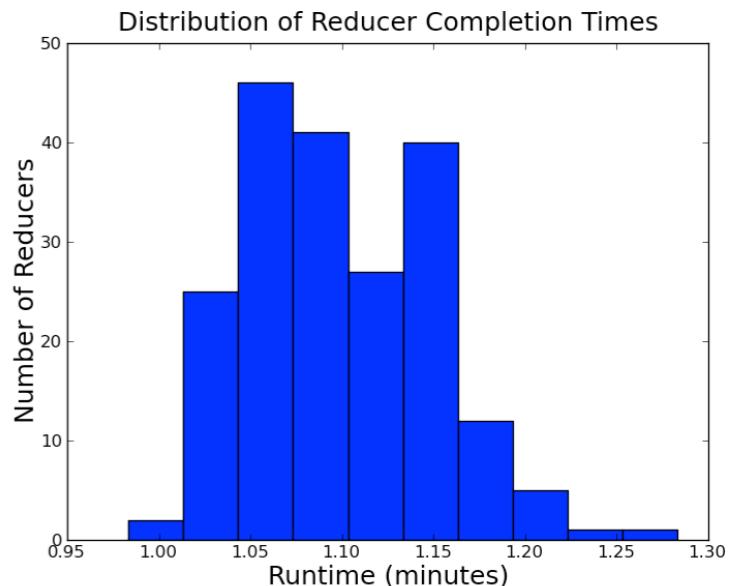
- Each edge considered only once $O(m)$ and max out-degree is $O(m^{1/2})$



```
foreach v in V
    foreach u,w in Adjacency(v)
        if deg(u) > deg(v) && deg(w) > deg(v)
            if (u,w) in E
                Triangles[v]++
```

Map reduce Modification parallelized

- Assume an ordering of the nodes : $u_1 < u_2 < \dots < u_n$
- Minor modification in the 2nd MR step:
 - Map 2: Send $\langle(v_1, v_2); u\rangle$ and $\langle(v_1, v_2); \$\rangle$ for $(v_1, v_2) \in E$ to same machine.
 - **Map:** emit $\langle(v_1, v_2); u\rangle$ only if $u < v_1$ and $u < v_2$
 - Reduce 2: input: $\langle(v, w); u_1, u_2, \dots, u_k, \$?\rangle$
 - Better parallelisation



Summary

- Community detection
 - Overlapping communities have **high edge density** in the overlapping region
 - **AGM** to model overlapping communities and **BIGCLAM** as a relaxed version
 - Efficiently performing MLE using **caching** and **MR**
- Counting triangles
 - **Heavy hitters** and partitioning for skew and runtime improvement
 - **Exploiting order** for better MR performance

More details at...

- [Overlapping Community Detection at Scale: A Nonnegative Matrix Factorization Approach](#) by J. Yang, J. Leskovec. *ACM International Conference on Web Search and Data Mining (WSDM)*, 2013.
- [Detecting Cohesive and 2-mode Communities in Directed and Undirected Networks](#) by J. Yang, J. McAuley, J. Leskovec. *ACM International Conference on Web Search and Data Mining (WSDM)*, 2014.
- [Community Detection in Networks with Node Attributes](#) by J. Yang, J. McAuley, J. Leskovec. *IEEE International Conference On Data Mining (ICDM)*, 2013.