




Management of End-Systems

Network Management

Prof. Dr. Panagiotis Papadimitriou



- End-systems:
 - Desktops, laptops, PDAs, other mobile terminals
 - Run applications (web, email, voice, video, etc.)
- Two service models:
 - Client / server
 - Peer-to-peer





- Network discovery and bootstrapping
 - Joining the network
 - Obtaining a network address
- Interface to network applications
 - Binding network applications to the transport layer
 - Identifying applications from port numbers
 - Port scanning
- Remote Procedure Calls (RPC)
 - XML-RPC




Network Discovery and Bootstrapping



- Network identifiers:
 - Host Name (e.g., `www.uni-hannover.de`)
 - IP address (e.g., `172.23.180.137`)
 - MAC address (e.g., `84-2B-2B-A5-A5-77`)
- Mapping between identifiers:
 - Address Resolution Protocol (ARP)
 - Physical address resolution
 - Domain Name System (DNS)
 - Host name resolution
- Dynamic IP address assignment
 - Dynamic Host Configuration Protocol (DHCP)




- DHCP is used for dynamic network configuration of end systems






- Sending IP packets requires:
 - the IP address of the next-hop (provided by the forwarding table)
 - the physical address of the next-hop



Next-hop interface (IP, **MAC???**)



IP packet encapsulation






- ARP is responsible for resolving physical addresses:
 - ARP allows routers/hosts in the network to build up a table of mappings between IP and physical addresses
 - each entry in the ARP table is removed after a time period (TTL)

IP Address	Physical Address	TTL
172.23.180.137	84-2B-2B-A5-A5-77	13:45:00
172.23.42.42	70-F3-95-3D-14-04	11:12:00

- When a host/router wants to identify a physical address:
 - it performs a lookup in the ARP table and obtains the physical address
 - if there is no match:
 - it broadcasts an ARP query onto the network
 - the host/router that has this IP address sends a message with its physical address (which is subsequently added to the ARP table)



- DNS provides hostname to IP address translation
- DNS uses an hierarchical, distributed database





- Local DNS servers:
 - the first contact point for hostname resolution
 - forward queries to root DNS servers
 - cache hostname mappings locally
 - cache entries expire after a certain period
- Root DNS servers:
 - contacted by local DNS servers that cannot resolve hostname
- Top-level Domain (TLD) servers:
 - responsible for com, org, edu, etc. and all top-level country domains (e.g., de, uk, fr, jp)
- Authoritative DNS servers:
 - belong to organizations (e.g., enterprises, universities)
 - provide hostname resolution for organization's servers (e.g., web, e-mail)



13 worldwide root servers

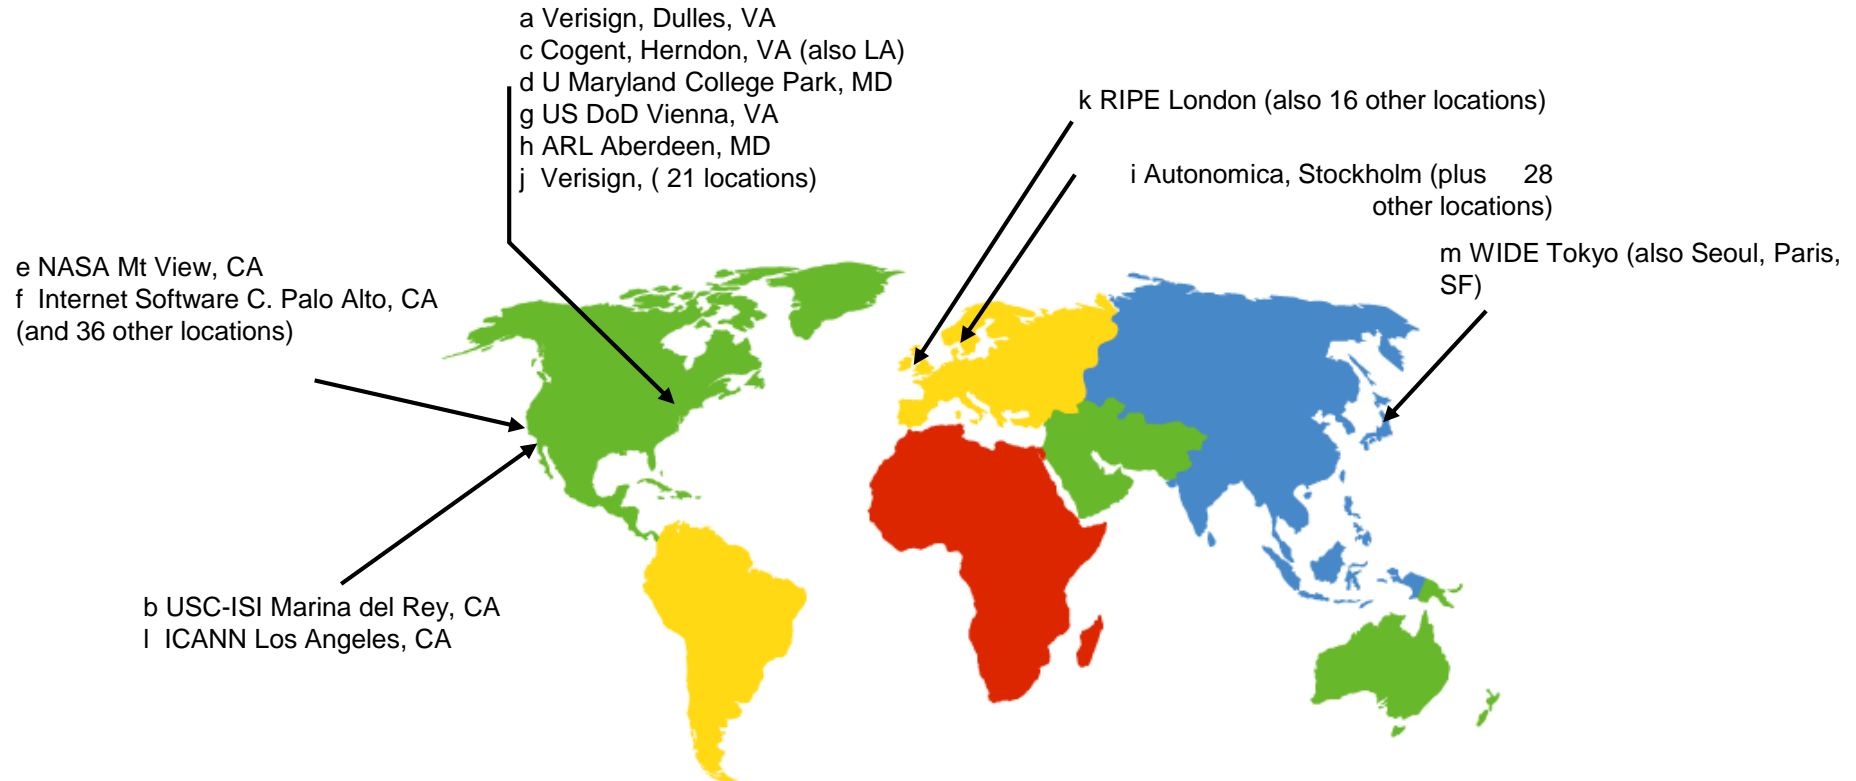
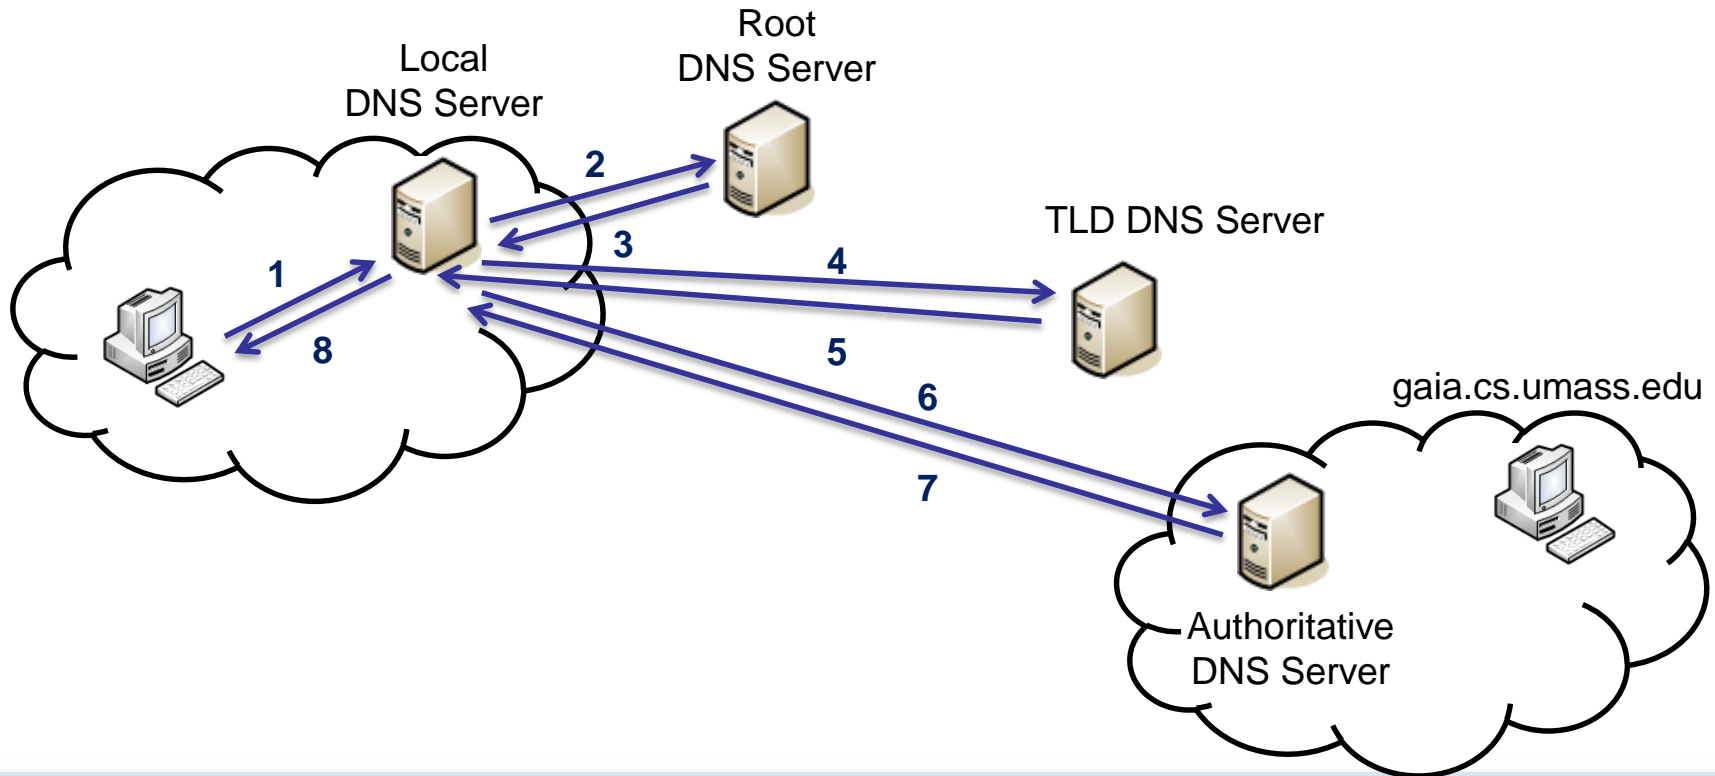


Figure from Computer Networking, A Top-Down Approach

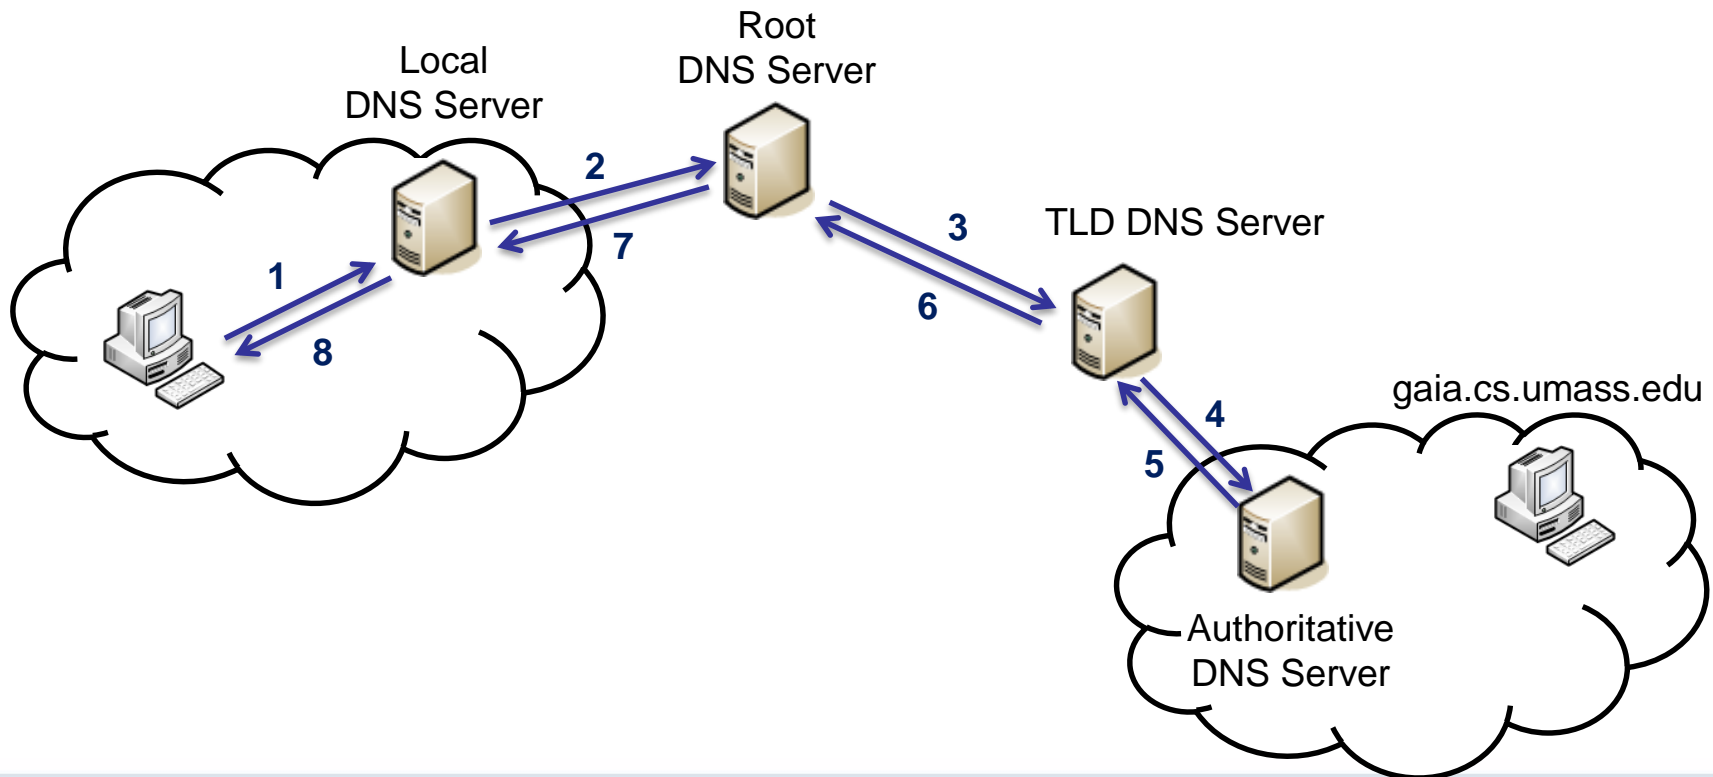


- The host **pc1.ikt.uni-hannover.de** queries the IP address for **gaia.cs.umass.edu**:
 - Iterated query





- The host **pc1.ikt.uni-hannover.de** queries the IP address for **gaia.cs.umass.edu**:
 - Recursive query





- DNS record format: `[name, value, type, TTL]`
- DNS record types:
 - Address (A) record for hosts
 - e.g., `[relay1.bar.example.com, 123.45.120.12, A, TTL]`
 - Mail Exchanger (MX) record for email servers
 - e.g., `[example.com, mail.example.com, MX, TTL]`
 - Name Server (NS) record for authoritative DNS servers of a domain
 - e.g., `[example.com, dns.example.com, NS, TTL]`
 - Canonical Name (CNAME) record:
 - name is an alias for a “canonical” name
 - e.g., `[example.com, relay1.bar.example.com, CNAME, TTL]`
 - allows the use of mnemonic hostnames
 - widely used for the redirection of client requests to a proximate server in content distribution networks (CDNs)

DNS Resolver Configuration



```
$TTL 86400
my-name.com.                IN      SOA      debns1.my-name.com. \
                             joe.my-name.com. {
                             2004011522    ; Serial no., based on date
                             21600         ; Refresh after 6 hours
                             3600          ; Retry after 1 hour
                             604800        ; Expire after 7 days
                             3600          ; Minimum TTL of 1 hour
                             )
;Name servers
debns1                       IN      A        192.168.1.41
debns2.joescuz.com.         IN      A        192.168.1.42

@                             IN      NS       debns1
my-name.com.                 IN      NS       debns2.my-name.com.

;Mail servers
debmail1                     IN      A        192.168.1.51
debmail2.my-name.com.       IN      A        192.168.1.52

@                             IN      MX       10 debmail1
my-name.com.                 IN      MX       20 debmail2.my-name.com.

;Aliased servers
debhp                        IN      A        192.168.1.61
debdell.my-name.com.        IN      A        192.168.1.62

www                          IN      CNAME     debhp
ftp.my-name.com.             IN      CNAME     debdell.my-name.com.
```



- DNSSEC (Domain Name System Security Extensions):
 - provides authentication of DNS data
 - DNS responses contain a DNSSEC certificate
 - deployed on root DNS servers


- edns-client-subnet:
 - includes part of the client IP address in the DNS request
 - allows the selection of a CDN server in proximity to the client (instead of the client's name resolver)



Interface to Network Applications

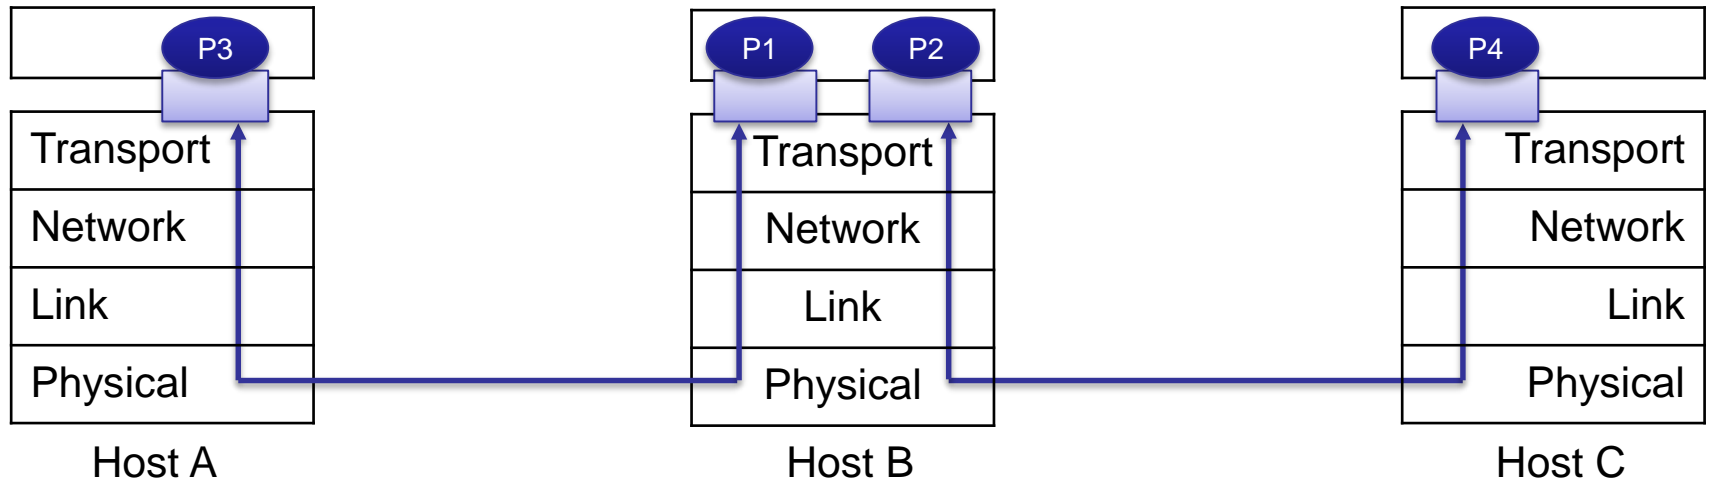


- Sockets provide an interface between applications and transport protocols



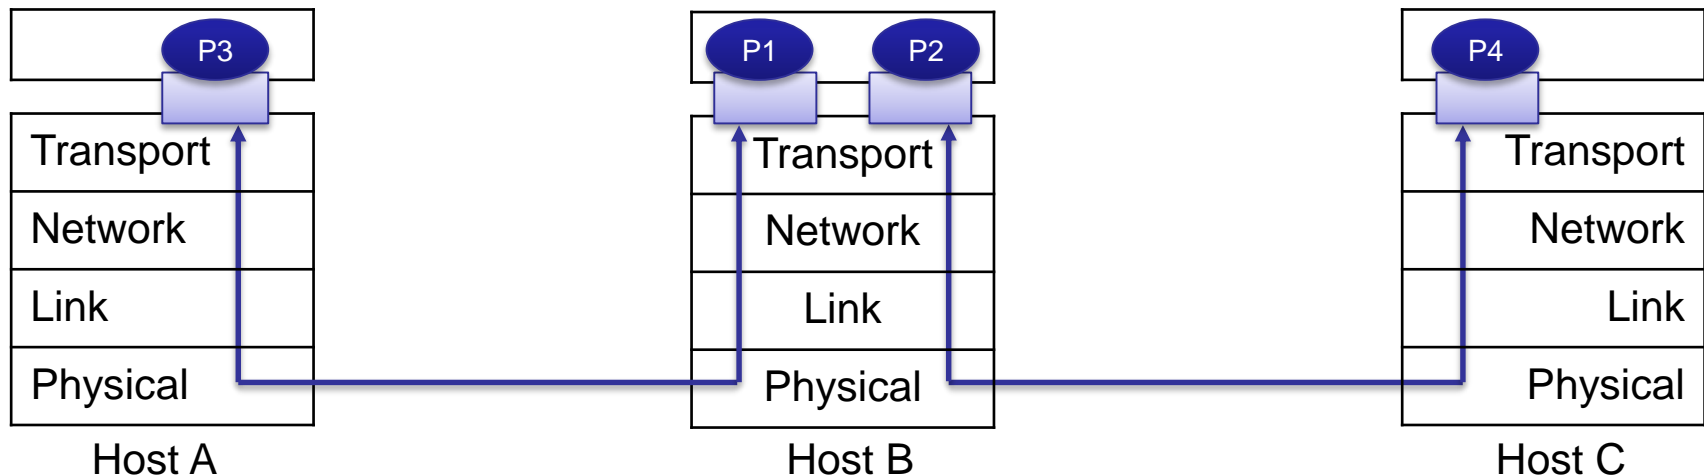


- Host uses IP address and port number to deliver a segment to the appropriate socket



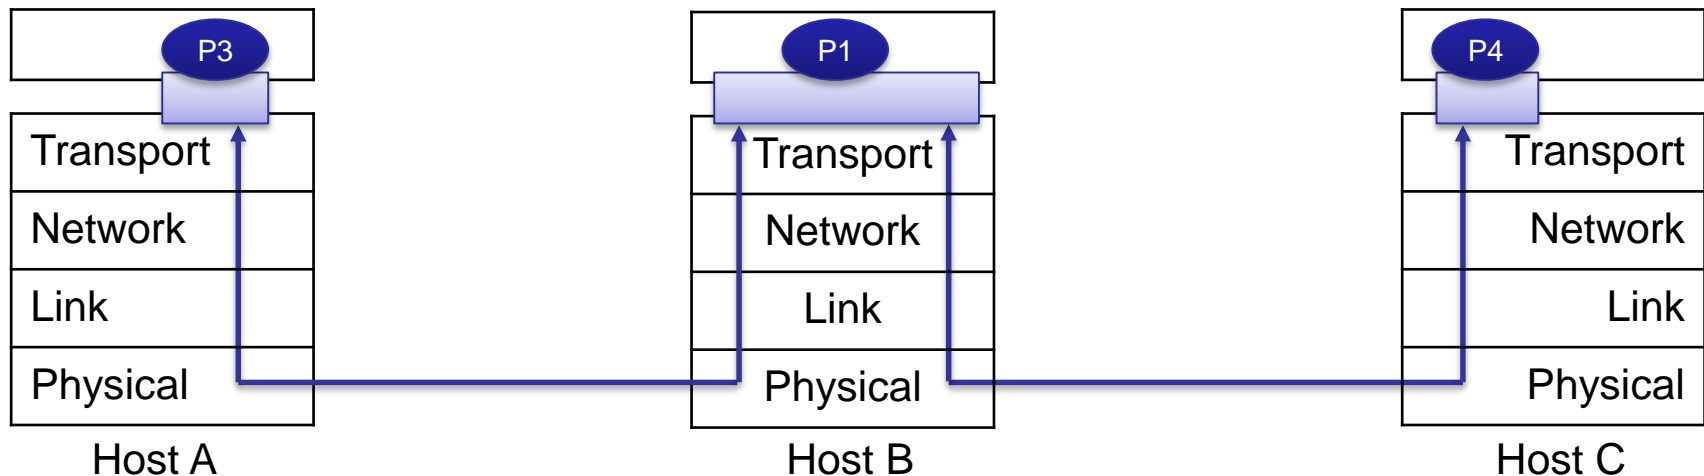


- TCP sockets are identified by a 4-tuple:
 - Source IP address
 - Source port number
 - Destination IP address
 - Destination port number
- A host may support multiple TCP sockets (e.g., host B):
 - Each socket is identified by its own 4-tuple





- UDP sockets are identified by a 2-tuple:
 - Destination IP address
 - Destination port number
- A host supports one UDP socket per port





UDP server socket

```
import socket
server_socket = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
server_socket.bind(("", 5000))

print "UDP server waiting for client on port 5000"

while 1:
    data, address = server_socket.recvfrom(256)
    print "( ", address[0], " ", address[1], " ) said : ", data
```

UDP client socket

```
import socket
client_socket = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
client_socket.sendto("Hello", ("125.10.37.1", 5000))
client_socket.close()
```



- Port numbers represented by 16-bit unsigned integers (0 – 65535)
- Classification of port numbers:
 - Well-known ports:
 - Range: 0 – 1023
 - Only for root-privileged applications (e.g., SSH)
 - Registered ports:
 - Range: 1024 – 49151
 - Dynamic and private ports:
 - Range: 49152 – 65535



Application	Protocol	Port Number
FTP	TCP	20, 21
SSH	TCP / UDP	22
Telnet	TCP	23
SMTP	TCP	25
DNS	UDP / UDP	53
HTTP	TCP / UDP	80
POP3	TCP	110
NTP	UDP	123
IMAP	TCP	143
SNMP	UDP	161
SNMP Trap	TCP	162
DHCP	TCP / UDP	546, 547



- Identifying applications from port numbers may be inaccurate:
 - Some (new) applications may not use well-known port numbers
 - Applications may tunnel their traffic over HTTP to go through firewalls (e.g., tunneling SSH over HTTP)
- Accurate application classification requires further inspection:
 - Packet header inspection
 - Deep-packet inspection for well-known signatures or protocol semantics
- Duration of inspection may vary depending on the classification technique:
 - 1st packet
 - 1st Kbyte
 - Flow

Identifying Applications (2)



Classification	Example Application
BULK	ftp
DATABASE	postgres, sqlnet, oracle, ingres
INTERACTIVE	ssh, klogin, rlogin, telnet
MAIL	imap, pop2/3, smtp
SERVICES	X11, dns, ident, ldap, ntp
WWW	www
P2P	KaZaA, BitTorrent, GnuTella
MALICIOUS	Internet work and virus attacks
GAMES	Half-Life
MULTIMEDIA	Windows Media Player, Real

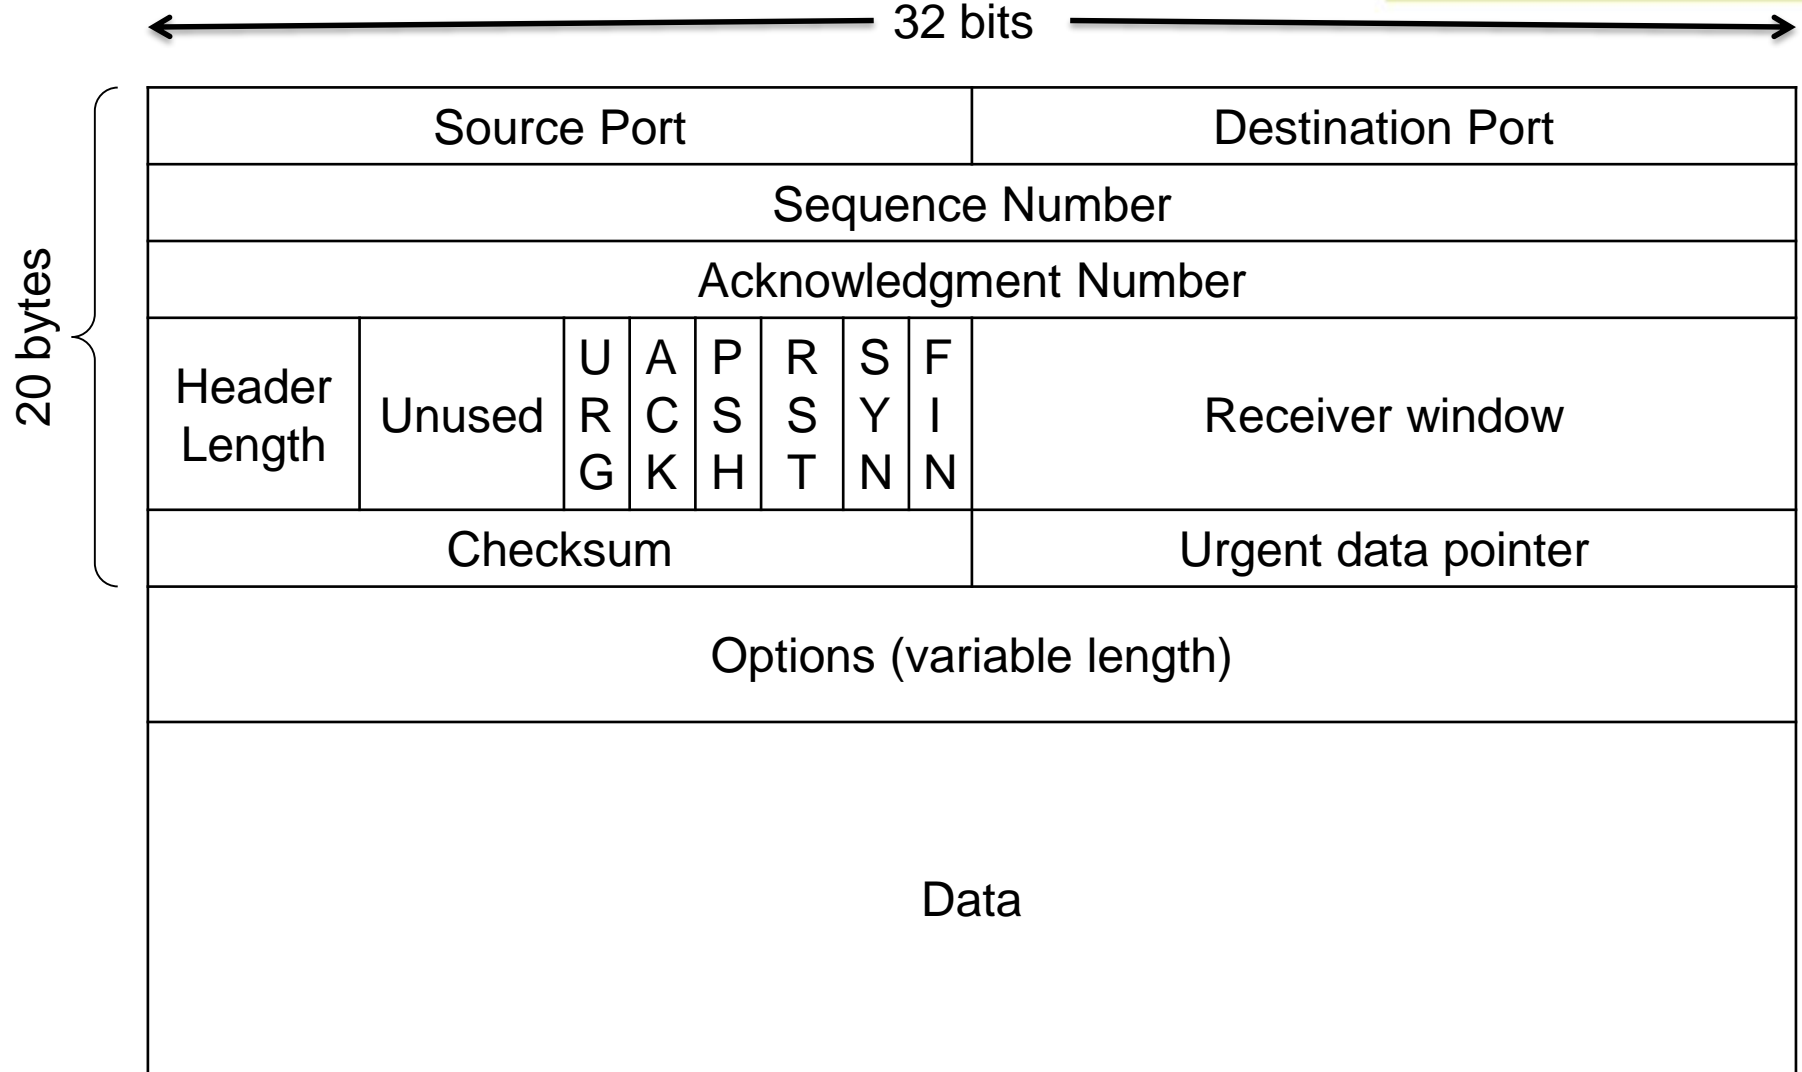
Classification Type	Port-Based		Content-Based	
	Packets	Bytes	Packets	Bytes
As a percentage of total traffic				
BULK	46.97	45.00	65.06	64.54
DATABASE	0.03	0.03	0.84	0.76
GRID	0.03	0.07	0.00	0.00
INTERACTIVE	1.19	0.43	0.75	0.39
MAIL	3.37	3.62	3.37	3.62
SERVICES	0.07	0.02	0.29	0.28
WWW	19.98	20.40	26.49	27.30
UNKNOWN	28.36	30.43	<0.01	<0.01
MALICIOUS	—	—	1.10	1.17
IRC/CHAT	—	—	0.44	0.05
P2P	—	—	1.27	1.50
GAMES	—	—	0.17	0.18
MULTIMEDIA	—	—	0.22	0.21

A. Moore and D. Papagiannaki, "Toward the Accurate Identification of Network", PAM 2005




- Port scanning allows attackers to discover open ports on a certain host:
 - Ports characterize the amount of exposure to external attacks
 - A system administrator should be aware of processes listening on ports and possible weaknesses arising by that

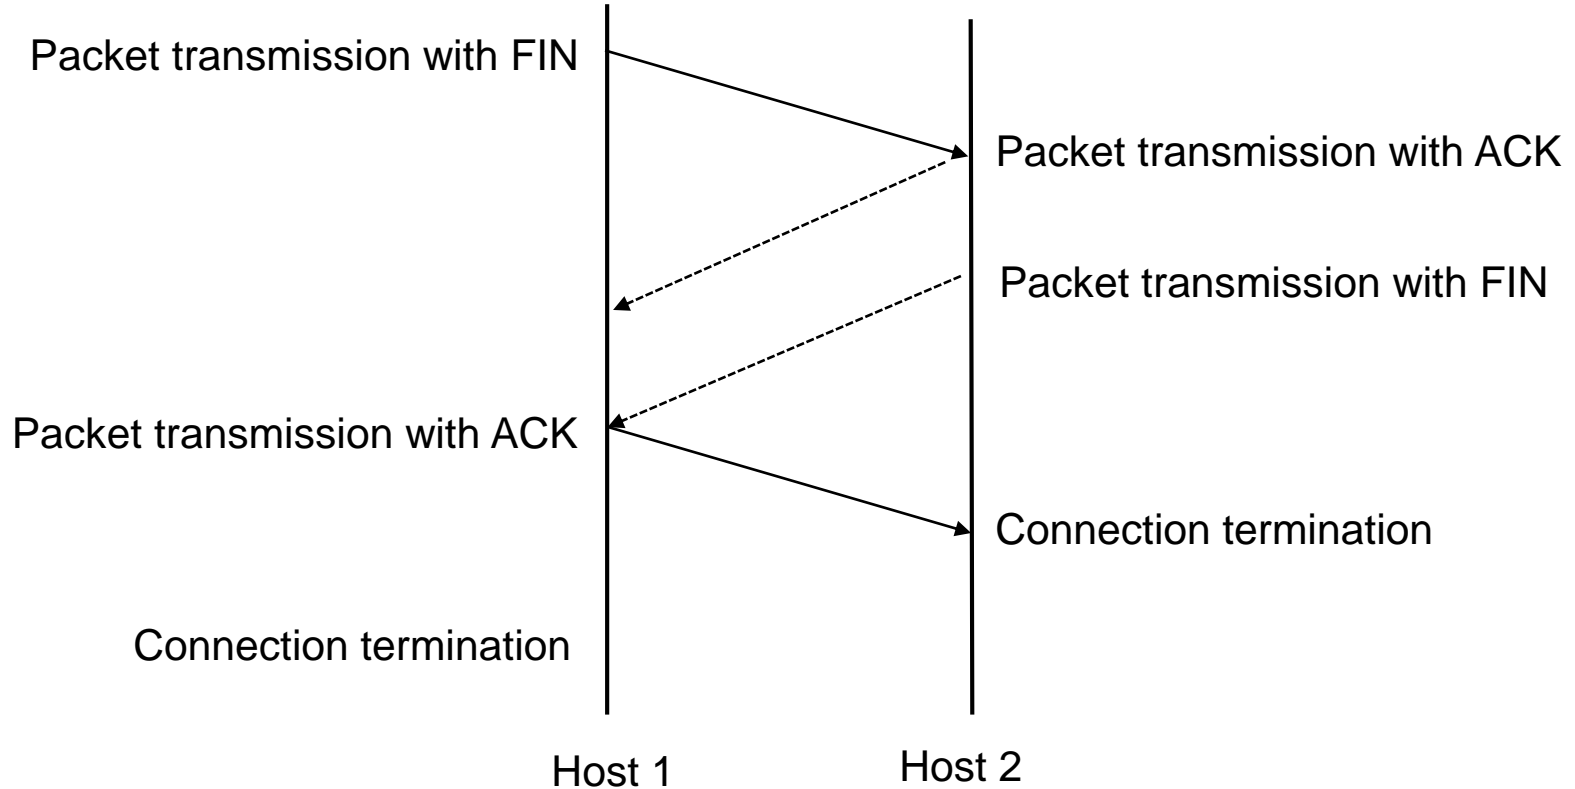
- Various port scanning techniques:
 - Full TCP connection scanning
 - TCP SYN scanning
 - TCP FIN scanning
 - Indirect scanning
 - UDP scanning





- TCP uses a “3-way handshake” to establish connection









- Try to setup connections with chosen ports on the target host:
 - Connection is established for every listening port
 - Otherwise, the port is unreachable

- Full TCP connection scanning:
 - Does not need any root privileges and can be performed by any user
 - Can be easily detected
 - Slow



- Try to setup half-open TCP connections with the target host:
 - The scanning host sends a SYN packet and waits for the response
 - Reception of SYN/ACK indicates that the port is listening
 - The scanning host sends an RST to tear down the connection immediately
 - No response indicates that the port is unreachable

- TCP SYN scanning:
 - Requires root privileges
 - Hosts may not detect this type of port scanning







- Use FIN packets to probe ports on the target host:
 - The scanning host sends a FIN packet and waits for the response
 - If the port is closed, the target host will drop the FIN packet and respond with an RST packet
 - In the case of a listening port, the target host will drop the FIN packet without sending a response

- TCP FIN scanning:
 - Requires root privileges
 - Hosts may not detect this type of port scanning
 - Some hosts may not be vulnerable to TCP FIN scans, since they return RST packets regardless of the port state

TCP FIN Scanning





- Detection of open UDP ports on the target host:
 - The scanning host sends a UDP packet and waits for the response
 - Reception of an ICMP error message indicates that the port is unreachable
 - No response indicates that the port is open
- UDP port scanning:
 - Useful for the detection of phony servers or unauthorized servers (e.g., DNS, NFS, SMTP)




- System administrator may use several applications and loggers to detect port scans:
 - Courtney
 - Gabriel
 - scan_detector
 - TCP Wrapper
 - scanlogd
 - Argus
 - tcplogger



Remote Procedure Calls



- Remote procedure calls (RPC) make a call to a remote system appear as a local call:
 - Whether the server is local or remote is made transparent by RPC
 - RPC allows applications to become distributed transparently
 - RPC makes the architecture of the remote system transparent

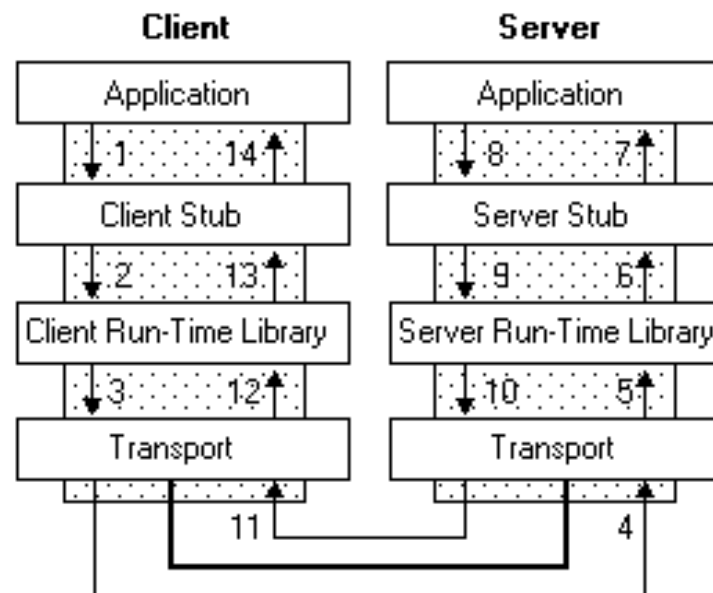




- Stubs run in the client and server carrying out data conversion, as needed:
 - The client and server use different address spaces
 - Pointers to one computer's memory will point to different data in a remote machine
 - The client and server may use different data representations even for simple parameters
 - e.g., big-endian vs. little-endian

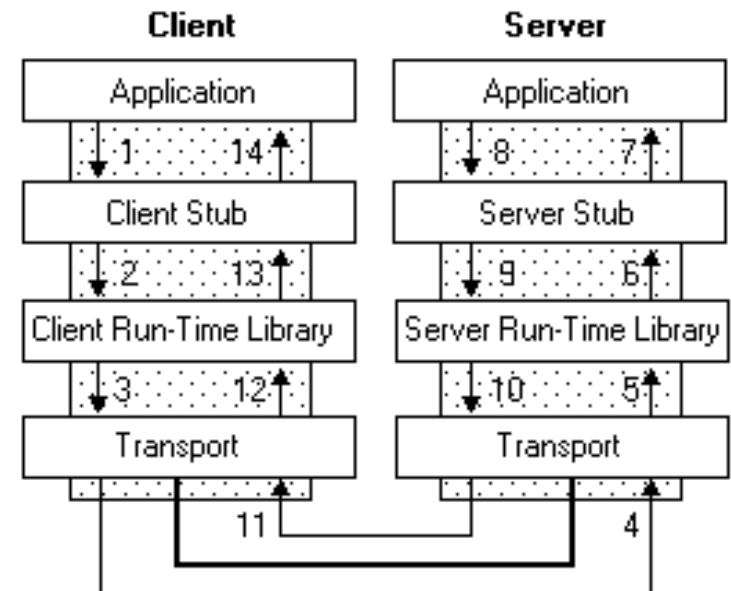


- The client performs the following steps to initiate the RPC at the server:
 1. Retrieves the required parameters from the client address space
 2. Translates the parameters into the required format for transmission over the network
 3. Calls functions in the RPC client run-time library to send the request and its parameters to the server



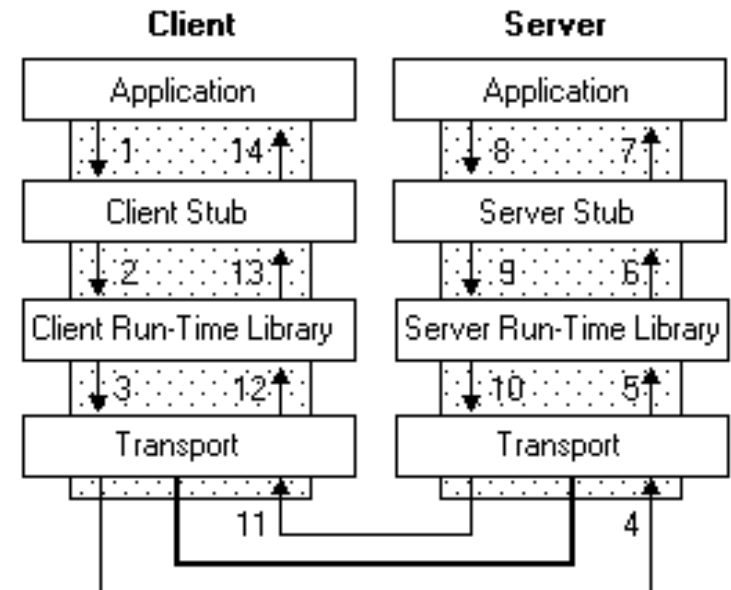


- The server performs the following steps to call the procedure:
 - 5. The server RPC run-time library functions accept the request and call the server stub procedure
 - 6. The server stub retrieves the parameters from the network buffer and converts them from the network transmission format to the format the server needs
 - 7. The server stub calls the actual procedure on the server



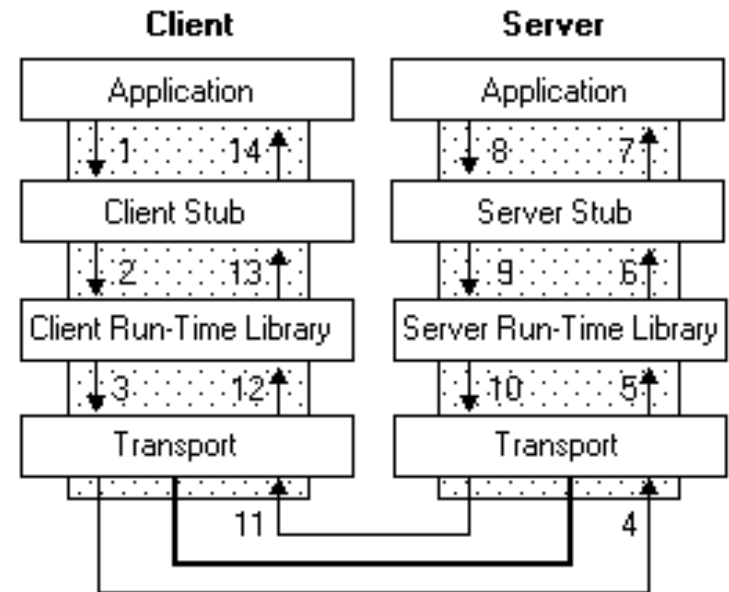


- When the remote procedure is complete, a similar sequence of steps returns the data to the client:
 - 8. The remote procedure returns its data to the server stub
 - 9. The server stub converts output parameters to the format required for transmission over the network and returns them to the RPC run-time library functions
 - 10. The server RPC run-time library functions transmit the data on the network to the client computer





- The client completes the process by accepting the data over the network and returning it to the calling function:
 - 12. The client RPC run-time library receives the remote-procedure return values and returns them to the client stub
 - 13. The client stub converts the data into the format used by the client computer. The stub writes data into the client memory and returns the result to the calling program on the client
 - 14. The calling procedure continues as if the procedure had been called on the same computer






- RPC can facilitate the management of end-systems:
 - Collection of management data from end-systems
 - Store and retrieve management data from a local database
 - Network monitoring
 - Periodically gather operational data
 - Compare previously collected data against the most recent
 - Device configurations
 - Add / modify / delete device configurations (e.g., network interfaces)
 - Fault management and diagnosis



- XML-RPC is an RPC protocol which uses:
 - HTTP for data transport
 - XML for data encoding



Source: JY Stervinou



Data Type	Description
integer (i4)	32-bit signed integer
boolean	boolean (0 or 1)
string	string
double	64-bit signed, floating point number
dateTime.iso8601	pseudo ISO8601 timestamp,(e.g., 19980717T14:08:55). Compared to ISO8601, milliseconds and time zone information is missing
base64	base64 encoded binary data
struct	key-value pair. Keys are strings and values can be any valid data type
array	array of objects. The array elements can be any valid data type



- XML provides a self-describing data format:
 - Application reads data, parses it, and knows exactly what each constituent part of the data means

- Strengths of XML:
 - Easy to understand, parse and debug
 - Can handle complex data
 - Wide range of available tools
 - Extensibility and interoperability



```
<network id="net1">
  <link id="1">
    <description type="link">
      <host1 id="node1" />
      <host2 id="node5" />
      <interface1 id="interface1" />
      <interface2 id="interface0" />
      <bandwidth id="1024" />
    </description>
  </link>
  :
  :
  <link id="10">
    <description type="link">
      <host1 id="node3" />
      <host2 id="node8" />
      <interface1 id="interface1" />
      <interface2 id="interface1" />
      <bandwidth id="850" />
    </description>
  </link>
</network>
```



- A. Moore and D. Papagiannaki, **Toward the Accurate Identification of Network Applications**, PAM 2005
- M. de Vivo, et al., **A Review of Port Scanning Techniques**, ACM CCR, 1999
- **Extensible Markup Language (XML)**, <http://www.w3c.org/xml>
- **XML-RPC**, <http://www.xmlrpc.com/>