

Requirements

Basic idea of concurrent systems

Goals

Understand how to model a concurrent system of multiple independent entities.

Content

- ☐ Examples of agents
- ☐ System environment
- ☐ Agent types
- ☐ Multi-Agent Systems
- ☐ Agreement in MAS
- ☐ Simulation of MAS



- ❑ Organic Computing Systems typically ...
 - consist of **large numbers** of **independent entities**.
 - follow a **system-wide goal** (not necessarily known to individual entity).
- ❑ Each entity ...
 - is **context-aware**: It can “sense” properties of their (physical or virtual) environment.
 - is able to **make changes** to its environment.
 - follows an **individual goal** and makes **local decisions**.
- ❑ Multiple entities ...
 - **co-exist** in the same environment.
 - can act independently and concurrently.
 - sometimes **interact** with each other.
 - can be **cooperative or competing**.

❑ Examples of technical systems

- Groups of mobile robots
- Traffic control systems
- Networked camera systems
- Desktop Grid Computing



❑ Questions that arise

- How can we **model such systems**?
- What are the **important properties**?
- What are **generic problems** that often arise?
For example coordination and agreement.


❑ Modeling by agents or embodied agents

“Entities” = agents

- ❑ System consists of **multiple networked cameras** on a given area.
- ❑ Each camera can be modeled as an autonomous agent.
 - Has a **partial view** of the area (e.g. pattern recognition, object detection).
 - Autonomous **actions**: changes its field of view
 - Communication with other agents possible
- ❑ System goal:
 - Cooperatively observe given area
 - Detection of important events, e.g spatial partitioning, object tracking
 - Communication with users: alarm management



- ❑ System consists of **interconnected routers**.
- ❑ Each router can be modeled as an autonomous agent.
 - (In most cases) partial knowledge of the network topology
 - Autonomous action: routing decision
- ❑ System goal
 - Transport packets from source to destination
 - Transportation cost may influence forwarding decision → Routers may also compete with each other.



- ❑ The term “intelligent agent” refers to a single **computer system** that ...
 - is situated in some **environment**,
 - is aware of its environment (can measure: **sensory input**),
 - can autonomously affect the environment (control: **action output**),
 - (is able to learn and self-optimize)
- ❑ Example 1 (cameras):
 - Environment: physical area
 - Awareness: Camera knows its position/heading; can evaluate images
 - Action output: change field of view
- ❑ Example 2 (routers):
 - Environment: network topology; data packets
 - Awareness: packet rate, size of queues, latencies
 - Action output: routing decision

- ❑ What makes an agent intelligent? (Wooldrige and Jennings, 1995)
- ❑ **Reactivity**: An agent reacts to changes in its environment in a timely fashion (according to its design objectives)
 - In example 1: Follow moving object; compensate failure of other cameras
 - In example 2: Forward incoming packet; change local routing strategy when a link fails
- ❑ **Pro-activeness**: Agents exhibit goal-directed behavior to achieve their goals.
 - Not always easy to achieve: In some systems the precondition and/or the goal may change during execution (e.g. due to **concurrency** or **uncertainty**).
 - In example 1: “search” for objects
 - In example 2: probe link latencies to chose best route to destination
- ❑ Problem in general: Achieve proper **balance between reactivity (fast) and pro-activeness (intelligent)**.

- ❑ “**Social ability**”: Agents can interact with each other.
 - Typically no single agent can achieve the system goal.
 - Agents often need to **cooperate** or **negotiate**.
 - Agents may be **selfish**: They try to maximize their own benefit.
 - In example 1:
 - Perspective of multiple cameras may be necessary.
 - Cameras may avoid cost (energy) of changing their field of view.
 - In example 2:
 - Routers need to cooperatively forward packets.
 - Routers may avoid forwarding cost.

- ❑ To define the system precisely, a (little) more formal wording is helpful 😊

- ❑ An **Agent architecture** consists of defining
 - The environment
 - A perceptual model for the agent
 - The set of actions an agent can take
- ❑ **Agents** can then be divided into several **classes**
 - Purely reactive (stateless)
 - Stateful
 - follows a plan
 - Model-based reactive agent
 - Goal and utility oriented
 - Learning agents: in order to adapt and improve

□ Classification of environments:

- **Accessible** versus **inaccessible**: Whether or not the agent has accurate and up-to-date information about the complete environment (or something in-between)
- **Deterministic** versus **non-deterministic**: Whether or not an agent's action has a predictable effect
- **Static** versus **dynamic**: In a static environment only the agent can modify the environment (multiple agents → dynamic environment).
- **Discrete** versus **continuous**: Discrete environments have a finite (albeit possibly large) number of percepts.

□ In most technical systems: Environment can be regarded as discrete and is defined as

- $E = \{ e_0, e_1, e_2 \dots \}$ with e_i .. discrete sequential states
- For example: all cameras' and tracked objects' positions

- The agent's (finite) **action set** is defined as
 - $A_c = \{ a_0, a_1, \dots \}$ with a_i the available actions (behavioral repertoire)
 - For example: a_1 : forward packet to output o_2
- Interaction between agent and environment can be modeled as a **run r** :
 - $r: e_0 \rightarrow a_0 \rightarrow e_1 \rightarrow a_1 \rightarrow e_2 \rightarrow a_2 \rightarrow \dots \rightarrow a_{u-1} \rightarrow e_u$
 - Where e_0 is the initial state of the environment
 - The set R contains all possible **finite runs**.
 - For example: sequence of forwarding decisions
- R^{A_c} is the subset of R ending with an action.
- R^E is the subset of R ending with an environment state.

□ In non-deterministic environments:

- Same sequence of actions on e_0 **may lead to different runs.**

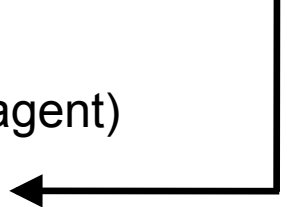
□ State transformer function

- maps a run R^{Ac} to a set of possible environment states.
- $T(r): R^{Ac} \rightarrow 2^E$ *(power set of E)*

$$2^{\{a,b\}} = \{\emptyset, \{a\}, \{b\}, \{a,b\}\}$$


- ❑ Based on the previous definition, the **abstract behavior** of an agent can be modeled as a function Ag mapping environment states to actions Ac .
- ❑ Purely **reactive agent** (stateless):
 - $Ag: E \rightarrow Ac$
 - Chosen action only depends on current state of the environment.
 - Example: thermostat
- ❑ **Stateful agent**:
 - $Ag: R^E \rightarrow Ac$
 - Chosen action Ac depends on history of the system.
 - Example: Navigation system

R^E : Run ending with
an environment state

- ❑ In many systems agents **cannot access** the **complete state** of the environment.
 - ❑ The **limited perception** can be modeled as:
 - see: $E \rightarrow Per$
 - where Per is the **agent's view** of the environment.
 - ❑ Now the action chosen by an agent is based on a (sequence of) perceptions:
 - action: $Per \rightarrow Ac$ (for a purely reactive agent)
 - action: $Per^* \rightarrow Ac$ (for a stateful agent)
- 

- ❑ So far, the state of an agent is composed of the complete sequence of perceptions and actions → may be inconvenient
- ❑ Alternative: Explicit **internal agent state**
 - Goal: smaller number of possible states
 - **State transition** is derived from current state and perception:
next: $I \times Per \rightarrow I$ (where I is the set of possible states)
 - Example: Current field of view of a camera and objects currently in view instead of the sequence of all actions and perceptions
- ❑ Then choosing an action is a function that maps an internal state to one of the actions:
action: $I \rightarrow Ac$

1. Agent starts in some initial internal state $i \leftarrow i_0$.
2. Agent observes its environment state e , and generates a percept $see(e)$.
3. Internal state of the agent is then updated via $next$ function, becoming $i = next(i, see(e))$.
4. The action selected by the agent is $action(next(i, see(e)))$.
5. Goto 2



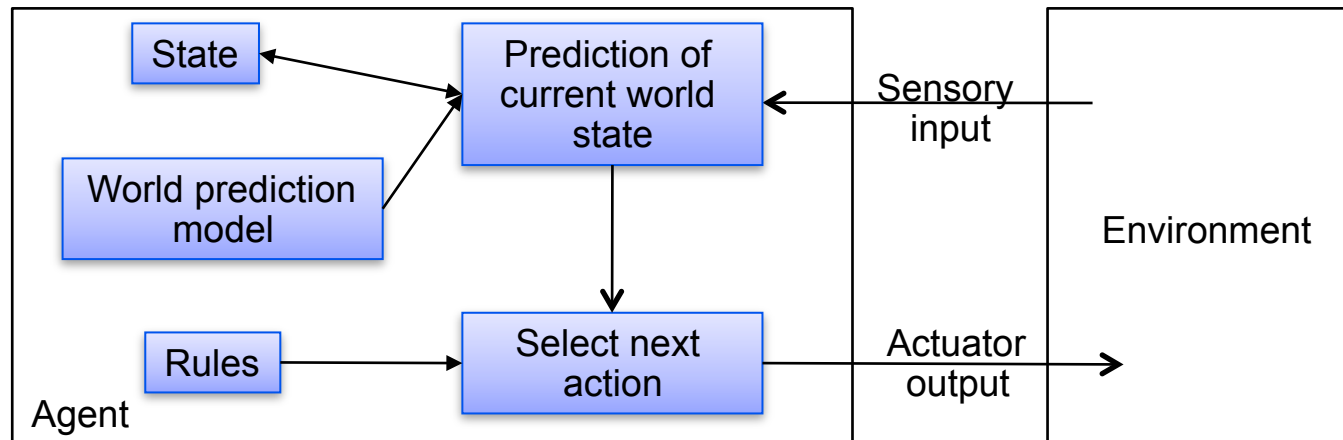
❑ Particular type of stateful agent

❑ Internal environment model

- Model in addition to current sensory input of a purely reactive agent
- Model reflects partial state of the “real” environment.
- Sensed information may be incomplete, inconsistent, or outdated.

❑ Outdated information may be compensated

- by letting the agent know “how the world works”.
- How? A model to predict changes is needed.
- Example:
 - A car started to overtake me: sensory input a while ago
 - Currently I cannot see it: currently out of sensor range
 - Assumption: the car is still there and will pass soon.



Function `Reactive-Agent-with-state(percept)` **returns** action

state: current known state of the world

rules: set of condition/action rules

action: previous action

state \leftarrow update-state(state, action, percept)

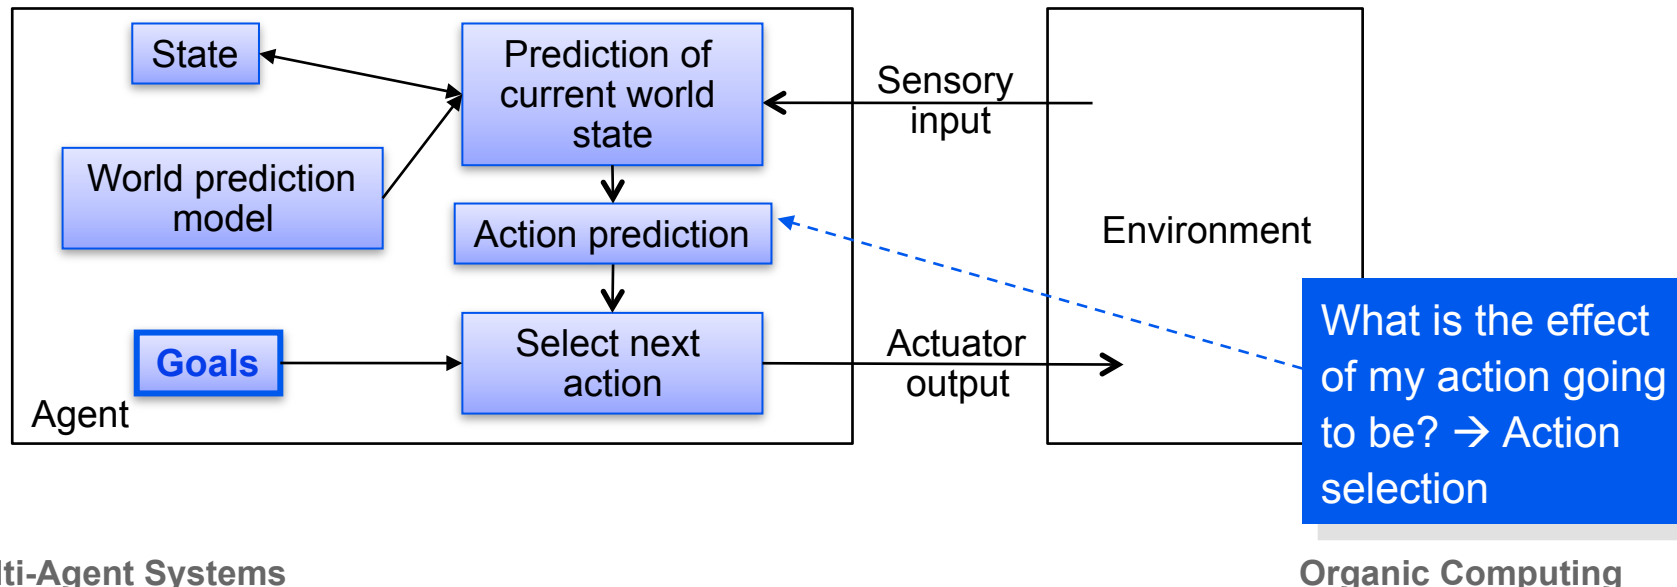
selected rule \leftarrow rule-match(state, rules)

action \leftarrow rule-action(selected rule)



Return action

Rules are fixed here! May be altered by learning algorithms.

- ❑ Additionally to the reactive behavior:
 - Give the agent a description of **what to achieve**.
 - Define a metric for improvement caused by actions → **utility**
- ❑ Example: Additionally to a car driving around (without causing an accident), tell it where to go and what to optimize for (e.g. distance)!
- ❑ Which of the possible **next actions** brings the agent **closest to its goal**?



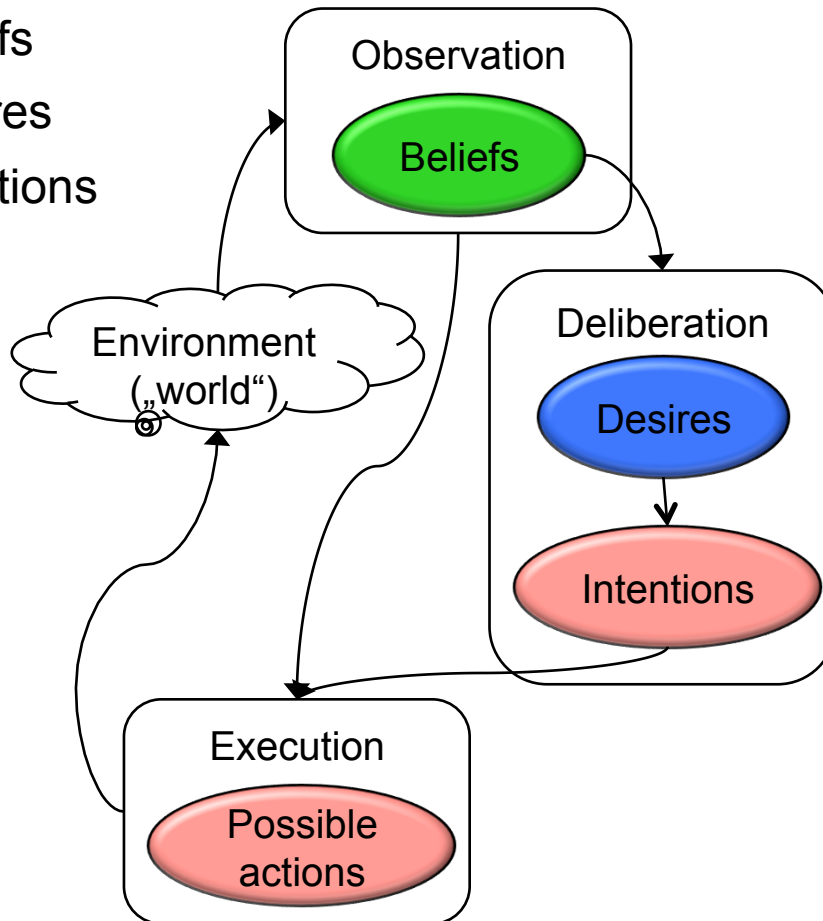
- ❑ Goal: Adjust camera's (agent's) field of view in order to **minimize overlap** with neighbors.
- ❑ Model of an agent:
 - **Environment** consists of the settings of all cameras (position, field of view).
 - Only changes of neighbors can be **perceived**.
 - Other agents' changes are perceived as environmental changes (no explicit interaction among agents).
 - Internal **states** $I = \{ \text{waiting, timeout, neighborChanged} \}$
- ❑ Idea:
 - If change of neighbor is perceived \rightarrow consider changing own settings
 - If nothing happened for a while \rightarrow re-announce own settings



- The BDI software model implements the principal aspects of Michael Bratman's *) theory of human practical reasoning (also referred to as Belief-Desire-Intention, or BDI).
- Psychological model to explain future-directed intention.
- Used in AI and MAS
- Three main components:
 - Beliefs (about the world)
 - Desires (to change the world, = goals)
 - Intentions (how to do it)

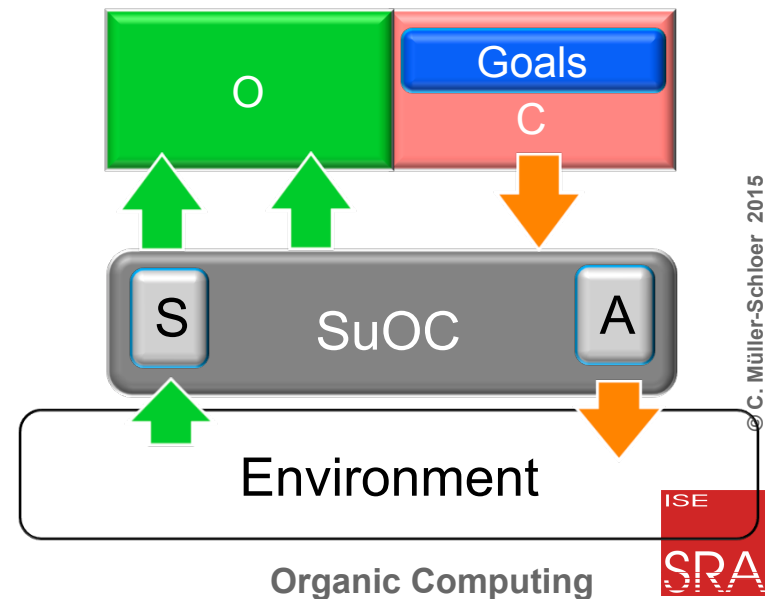

*) **Michael E Bratman** (born 25 July 1945) is Durfee Professor in the School of Humanities & Sciences and Professor of Philosophy at Stanford University. His interests include philosophy of action and moral philosophy. His work in those areas led him to the Belief-Desire-Intention model that is used in many areas, including artificial intelligence, today.

- ❑ Beliefs
- ❑ Desires
- ❑ Intentions



Problem: Practical usage in agents requires reasoning (AI), slow.

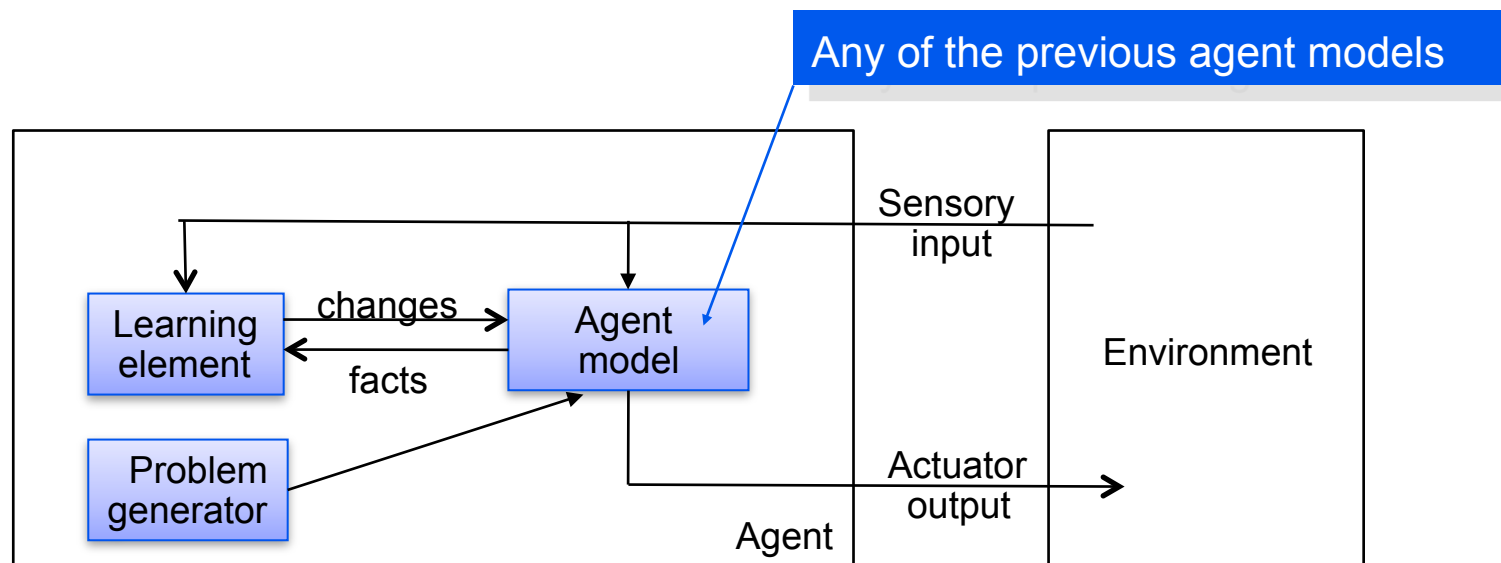
Observer/Controller vs. BDI



- ❑ Agents may improve their capabilities over time.
 - **Adapt** to changing environment.
 - **Improve** quality of performed tasks.
 - Find **new solutions**.
- ❑ Many machine learning algorithms available, e.g.
 - Artificial Neural Networks (ANN)
 - Learning Classifier Systems (LCS)
 - Evolutionary Algorithms (EA)
- ❑ Choice of algorithm class depends on particular problem.
- ❑ We will see some examples later in the lecture.

- ❑ Additional **learning element** ...
 - can modify the agent in operation.
 - gets feedback from the agent's behavior.
- ❑ **Problem generator** encourages new behavior.
- ❑ Online changes may be dangerous and/or time consuming.
- ❑ Online versus offline-learning

Problem generator can be the world itself!



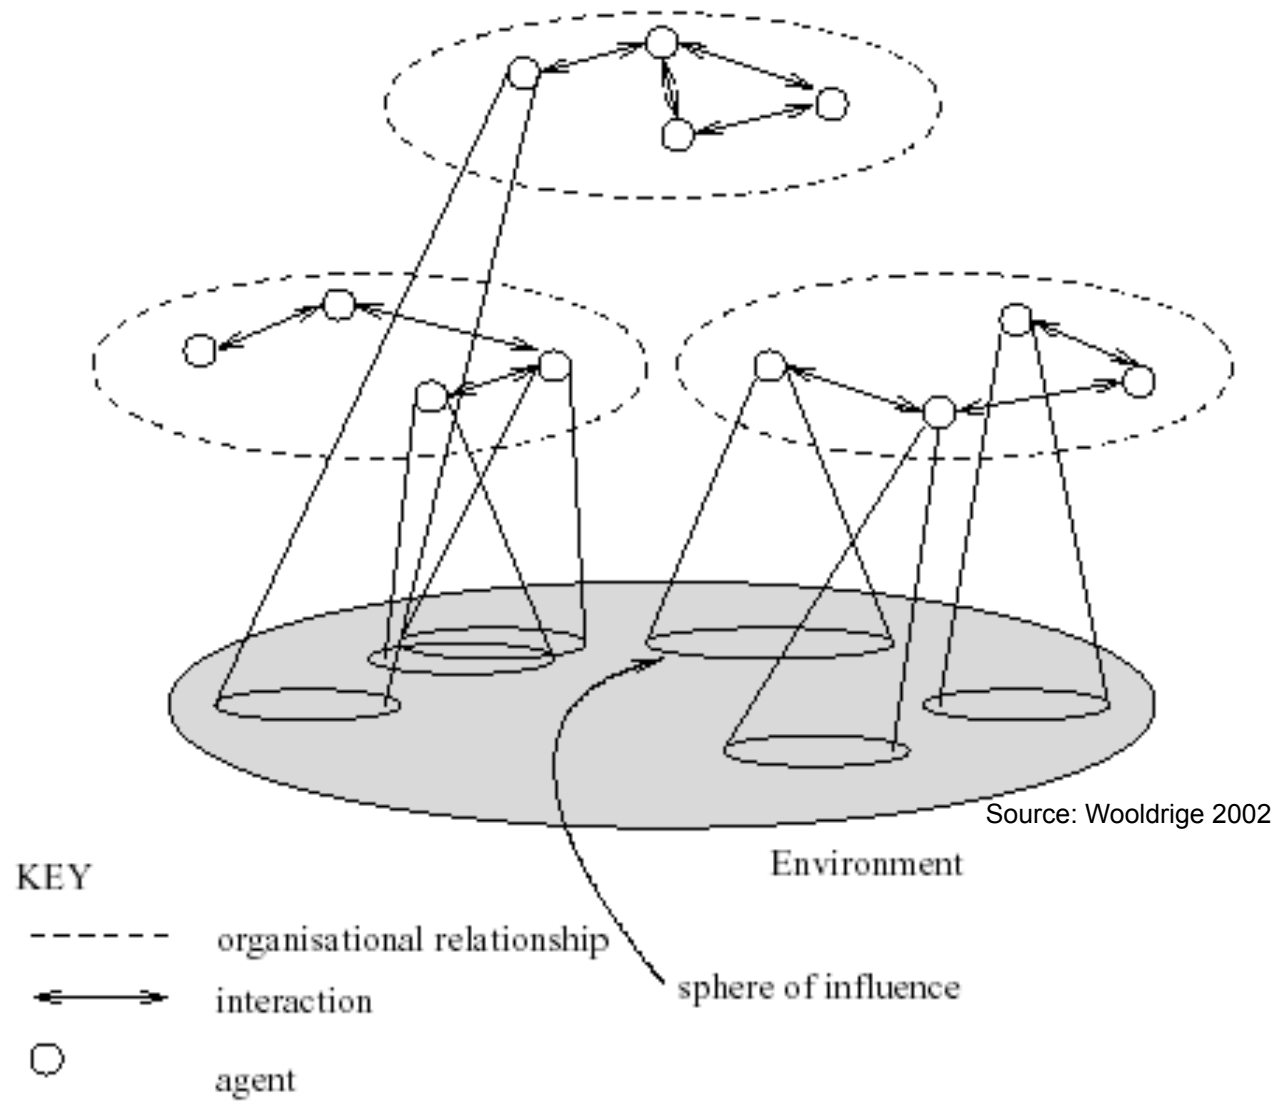
□ So far we have considered individual agents. But:

- Multiple agents share an environment.
- Agents perceive other agents' changes.
- So far no explicit coordination and collaboration among agents

→ What if we have **multiple agents working together** in the same environment?


□ Goal of **Multi-Agent Systems** (MAS)

- Make explicit use of **collaboration** mechanisms among agents!
- Own action may be chosen after **negotiating** with other agents!
- Example: Which camera should track a particular object?



- ❑ In many technical systems all agents “belong to” a single owner.
 - These agents can typically be modeled as **cooperative** agents.
 - Example: Communication network of a single operator
- ❑ Whenever agents “serve” different owners ...
 - agents should be modeled as **competing** or **selfish** agents.
 - Selfish (groups of) agents aim to **maximize** their **utility**.
 - Example: In a multi-operator network, each operator may try to reduce the cost of forwarding packets for others.
- ❑ **Agreement problems** become more difficult to solve, if agents are selfish, e.g.:
 - All the goodies: Resource allocation, election, voting, ...
 - Auctions
 - ...

- ❑ In case of **cooperative agents**: Agreement can be achieved by many algorithms for different purposes. (→ “classical” distributed systems)
- ❑ Example 1: Majority-consensus voting algorithms
 - Goal is to get a majority of agents to agree on doing something.
- ❑ Example 2: Leader Election
 - Goal: All agents must agree electing one single agent among them (e.g. to do a particular task).
 - Works because none of the agents “lies” to avoid getting the job.



Assumptions:

- unidirectional ring
- all processes have unique id
- highest id shall be the leader.

Unless someone
lies!

An idea for leader election on a ring topology:

- ❑ Each process wakes up at some time, at the latest when it receives a message.
- ❑ Start a complete round on the ring: message contains the highest id found so far.
- ❑ When message returns, it contains the highest id on the ring.
- ❑ Highest id becomes the leader, and each agents knows it.

- ❑ Consider an item to be sold in an auction.
- ❑ Goal: Sell item to that agent who really wants it most!
- ❑ Each bidder has a **private valuation** of the item.
- ❑ If everyone was honest, we could use the previous algorithm
 - Collect (honest) bids → highest bid wins
- ❑ Assumption: Agents are rational (as opposed to many humans!)
 - They do not bid more than their valuation.
 - May bid less than their valuation in order to pay less.
- ❑ How is it achieved that bids are truthful in relation to valuation?

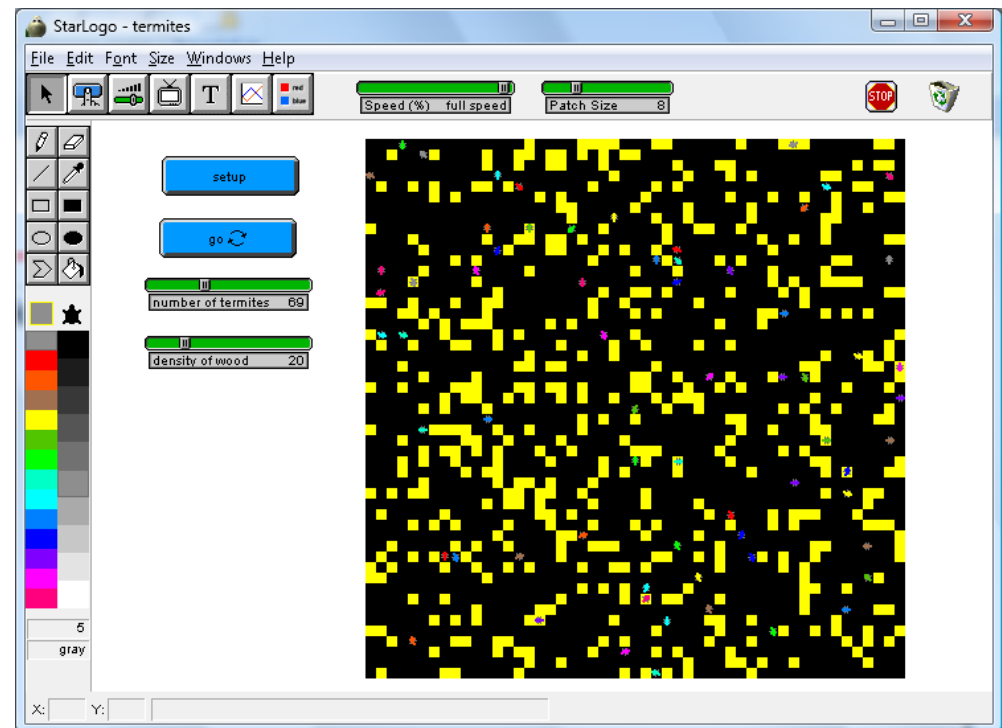
One solution: Vickrey auction


- ❑ Each agent places a sealed bid.
- ❑ When all bids have been placed:
 - The agent who gave the highest bid wins.
 - But only pays the price of the second highest bid.
- ❑ Encourages agents to use real valuations in their bids.
 - Bidding less than valuation → agent may lose auction to someone who wanted to pay a price closer to the real value.
 - Bidding more than valuation → risk to pay over budget
 - Intuitively: Agent does not risk paying “way more” than other bids.
- ❑ There are disadvantages: Think about cooperation between bidders!

Auction theory is part of Game theory.


- ❑ Mathematical analysis of large Multi-Agent Systems is – at least – very difficult in general.
- ❑ Simulation tools are often used for testing and evaluating systems.
- ❑ Many tools are available, e.g.
 - Starlogo (<http://education.mit.edu/starlogo/>)
 - Repast (<http://repast.sourceforge.net/>)
 - MASON (<http://cs.gmu.edu/~eclab/projects/mason/>)
 - NetLogo (<http://ccl.northwestern.edu/netlogo/>)
 - ...
- ❑ These tools offer/use a programming language for modeling environment and agents.

- ☐ 2D environment contains wood chips (yellow patches).
- ☐ Agents model termites.
Termites can ...
 - run around searching for wood chips.
 - pick up wood chips.
 - drop wood chips on a pile.
- ☐ Eventually there will be one single pile of wood chips.
- ☐ Example of a simple **stateful agent**.
- ☐ Included in the Starlogo bundle





☐ Random initialisation



☐ Almost one pile

❑ Implementation of the agent's behaviour

❑ Description of each agent's high-level behaviour:

```
to go
  search-for-chip    ; find a wood chip and pick it up
  find-new-pile      ; find another wood chip
  find-empty-spot    ; find a place to put down wood chip
end
```

❑ Agent walks around and searches for a chip:

```
to search-for-chip
  if pc = yellow      ; if find a wood chip...
    [stamp black      ; remove wood chip from patch
     setshape termite-wood-shape ; orange while carrying chip
     jump 20
     stop]             ; exit procedure
  wiggle
  search-for-chip
end
```

□ Now that the agent carries a chip, search for other chips (pile)

```
to find-new-pile
  if pc = yellow [stop]          ; if find a wood chip, stop
  wiggle
  find-new-pile
end
```

□ Go one step in random direction (hopefully away from the pile):

```
to find-empty-spot
  if pc = black                  ; if find a patch without a wood chip
  [stamp yellow                  ; put down wood chip in patch
   setshape termite-shape        ; set own color back to red
   get-away
   stop]
  seth random 360
  fd 1
  find-empty-spot
end
```

□ Once we found a pile and dropped the chip: go and look for new work:

```
to get-away                                ;leave the pile where you put your chip
  seth random 360
  jump 20
  if pc = black [stop]
  get-away
end
```

□ Make one step and randomly change heading:

```
to wiggle
  fd 1
  rt random 50
  lt random 50
end
```

- ☐ Examples of agents
- ☐ System environment
- ☐ Agent types
- ☐ Multi-Agent Systems
- ☐ Agreement in MAS
- ☐ Simulation of MAS

□ Textbooks

- Michael Wooldridge: An Introduction to MultiAgent Systems, Wiley, 2002
- Stuart Russel, Peter Norvig: Artificial Intelligence – A modern Approach / Künstliche Intelligenz – Ein moderner Ansatz, 2. Edition, Prentice Hall, 2004 (available in german and english)