

Requirements

...

Goals

Insight into an example of machine learning used in OC systems


Content

- ☐ Examples
- ☐ LCS overview
- ☐ ZCS: Zeroth-level classifier system
- ☐ XCS: Accuracy-based classifier system
- ☐ Application example: Organic traffic control (OTC)



- ❑ Machine learning techniques are a **promising approach for self-optimization** in organic computing systems.
- ❑ How can computers be programmed so that **problem solving** capabilities are built up by specifying “**what is to be done**” rather than “how to do it”? (Holland, 1975)
- ❑ Major issues:
 - How can the system react in unforeseen situation?
 - How can the system automatically improve its performance (if possible) at runtime?
 - Overall: Flexible and **autonomous reaction to changes of the environment and/or the system itself are desirable.**
- ❑ One example: Learning Classifier Systems

- ❑ The k-multiplexer is a Boolean function with k variables.
- ❑ The k variables consist of m address bits and 2^m data bits.
- ❑ Number of inputs: $k = m + 2^m$
- ❑ **Function:** Return the data bit (out of 2^m) that is specified by the address (m bits).
- ❑ Example for $m = 2$, $k = 6$
- ❑ Input: 01 1001
- ❑ Address 01b points to bit 1 of the data bit-string (starting to count from the right to the left; bits numbered $0..2^m-1$)
- ❑ Output: 0



© C. Müller-Schloer 2015

- ❑ **Question:** Can we build an agent that – given the MUX as a black-box – learns the underlying Boolean function?
- ❑ One idea to build such an agent:
 - Agent proposes output for a given input.
 - Agent proposal is based on an internal set of rules, e.g. one rule could look like 01 1001 : 0
 - After the proposal, agent receives feedback about the proposal.
 - Feedback is usually called **reward**; high/low reward for good/bad proposals
 - Agent aims at maximizing its reward.

- ❑ Example of an *Animat* problem
- ❑ Basis: Rectangular toroidal regular (n x m)-grid
- ❑ Each grid cell may contain a tree (t), food (F), or it may be empty.
- ❑ Food and trees fixed per instance
- ❑ Animat/agent/robot is initially randomly placed on empty cell.
- ❑ Walks around, looking for food.
- ❑ In each step agent can go to one of the eight neighboring cells but to empty and food cells only.

Example 2: The Woods Scenario

t	t	t	t	t	t	t	t	t	t	t	t	t
t	t				t	t	t	t		t	t	
t		t	t	t		t	t		t		t	
t		t	t	t		t		t	t	t		t
t	F	t	t	t		t	t		t	t	t	t
t	t	t	t	t	t			t	t	t	t	t

Woods14

t	t	F		
t	t	t		
t	t	t		

Woods1

Woods1: Optimal average number of steps to reach food: 1.7 steps (Bull & Hurst, "ZCS redux")

© C. Müller-Schloer 2015


- ❑ **Question:** Can we build an agent that can efficiently find food “in the Woods” **without global knowledge**?
- ❑ One idea to build such an agent:
 - Suppose the agent can “see” the eight surrounding cells
 - Based upon this perception it has to decide where to go next.
 - Reward is paid once the food is found.
- ❑ One approach to building such agents is the use of a **Learning Classifier System**.

- The k-multiplexer and the Woods-scenario are representing two important problem classes:
 - **Single-step problems**, e.g., the k-multiplexer
 - On each input/action there is an immediate feedback: right or wrong
 - **Multi-step problem**, e.g., the Woods scenario
 - Feedback about the quality of actions may be delayed.
 - Multiple steps are necessary to reach the goal.

- ❑ Obviously we could solve the k -multiplexer by letting the agent enumerate all input states (k bits).
- ❑ When k becomes large: **Space and time complexity are too high.**
 - Enumeration is exponential in time w.r.t. k .
 - Enumeration is exponential in space w.r.t. k .
- ❑ Problematic: **Reward may not be received immediately**, if agent interacts with the real world: Enumeration possible?
- ❑ So, we are looking for a heuristic that
 - Produces high quality results (yet possibly suboptimal) in much less time...
 - and requires much less space.

□ Abstract view

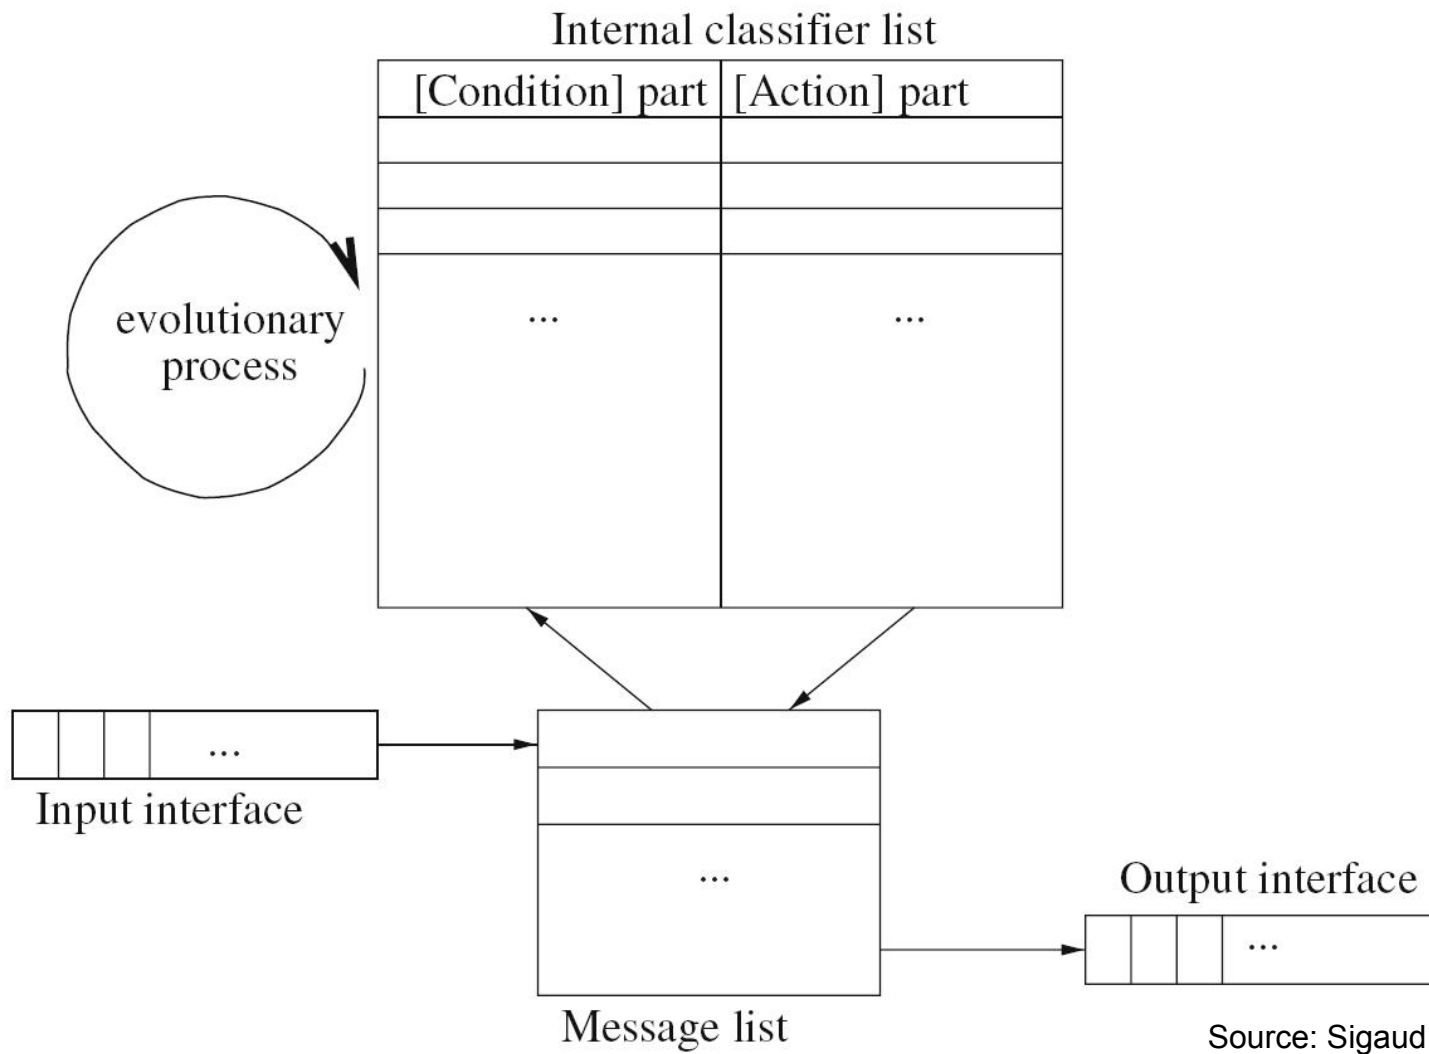
- **Observation**: Obtain (sensory) input about the current state of the environment
- **Action**: After reasoning about the current state, the agent decides on an action that impacts the environment.
- **Evaluation**: The agent observes the effect of its action and evaluates it (good/bad; reward)



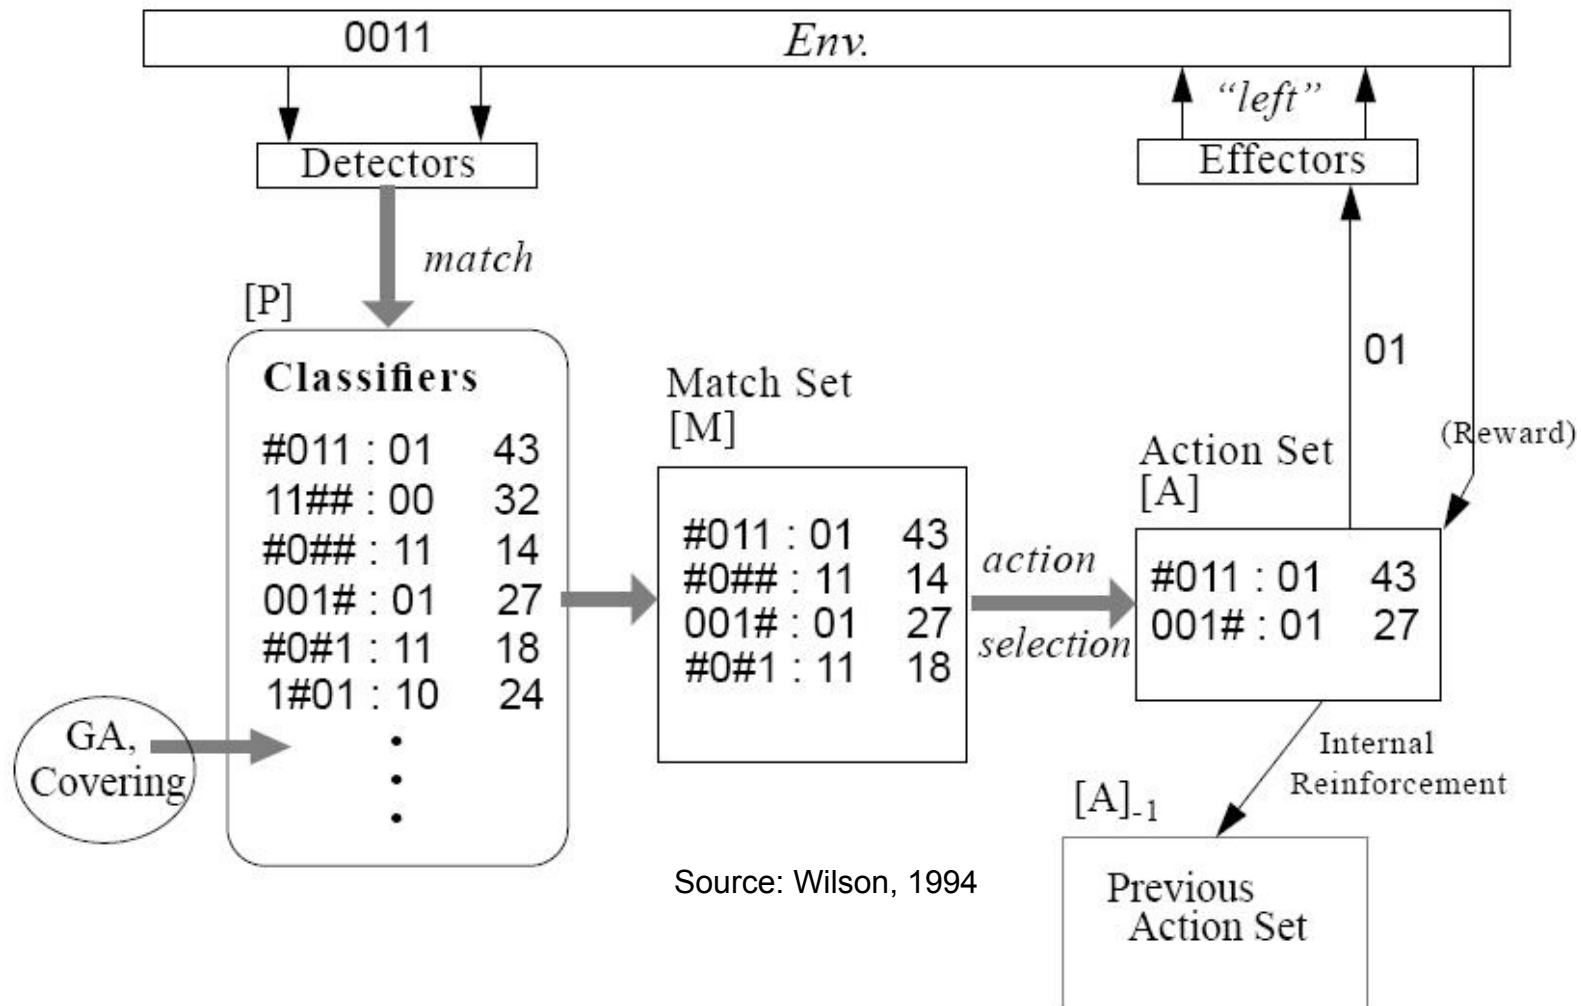
- ❑ Initial LCS was introduced by [John H. Holland](#) in 1975.
- ❑ He was (and still is) interested in complex adaptive systems.
- ❑ How can computers be programmed so that [problem-solving capabilities are built up by specifying “what is to be done”](#) rather than “how to do it”? (Holland, 1975)
- ❑ An important development in LCS was done by [Stewart W. Wilson](#) in 1995.
- ❑ Based on the initial approach by Holland, Wilson proposed a simplified and more efficient classifier system called XCS.
- ❑ XCS is today one of the most studied classifier systems.
- ❑ Many extensions have been proposed.



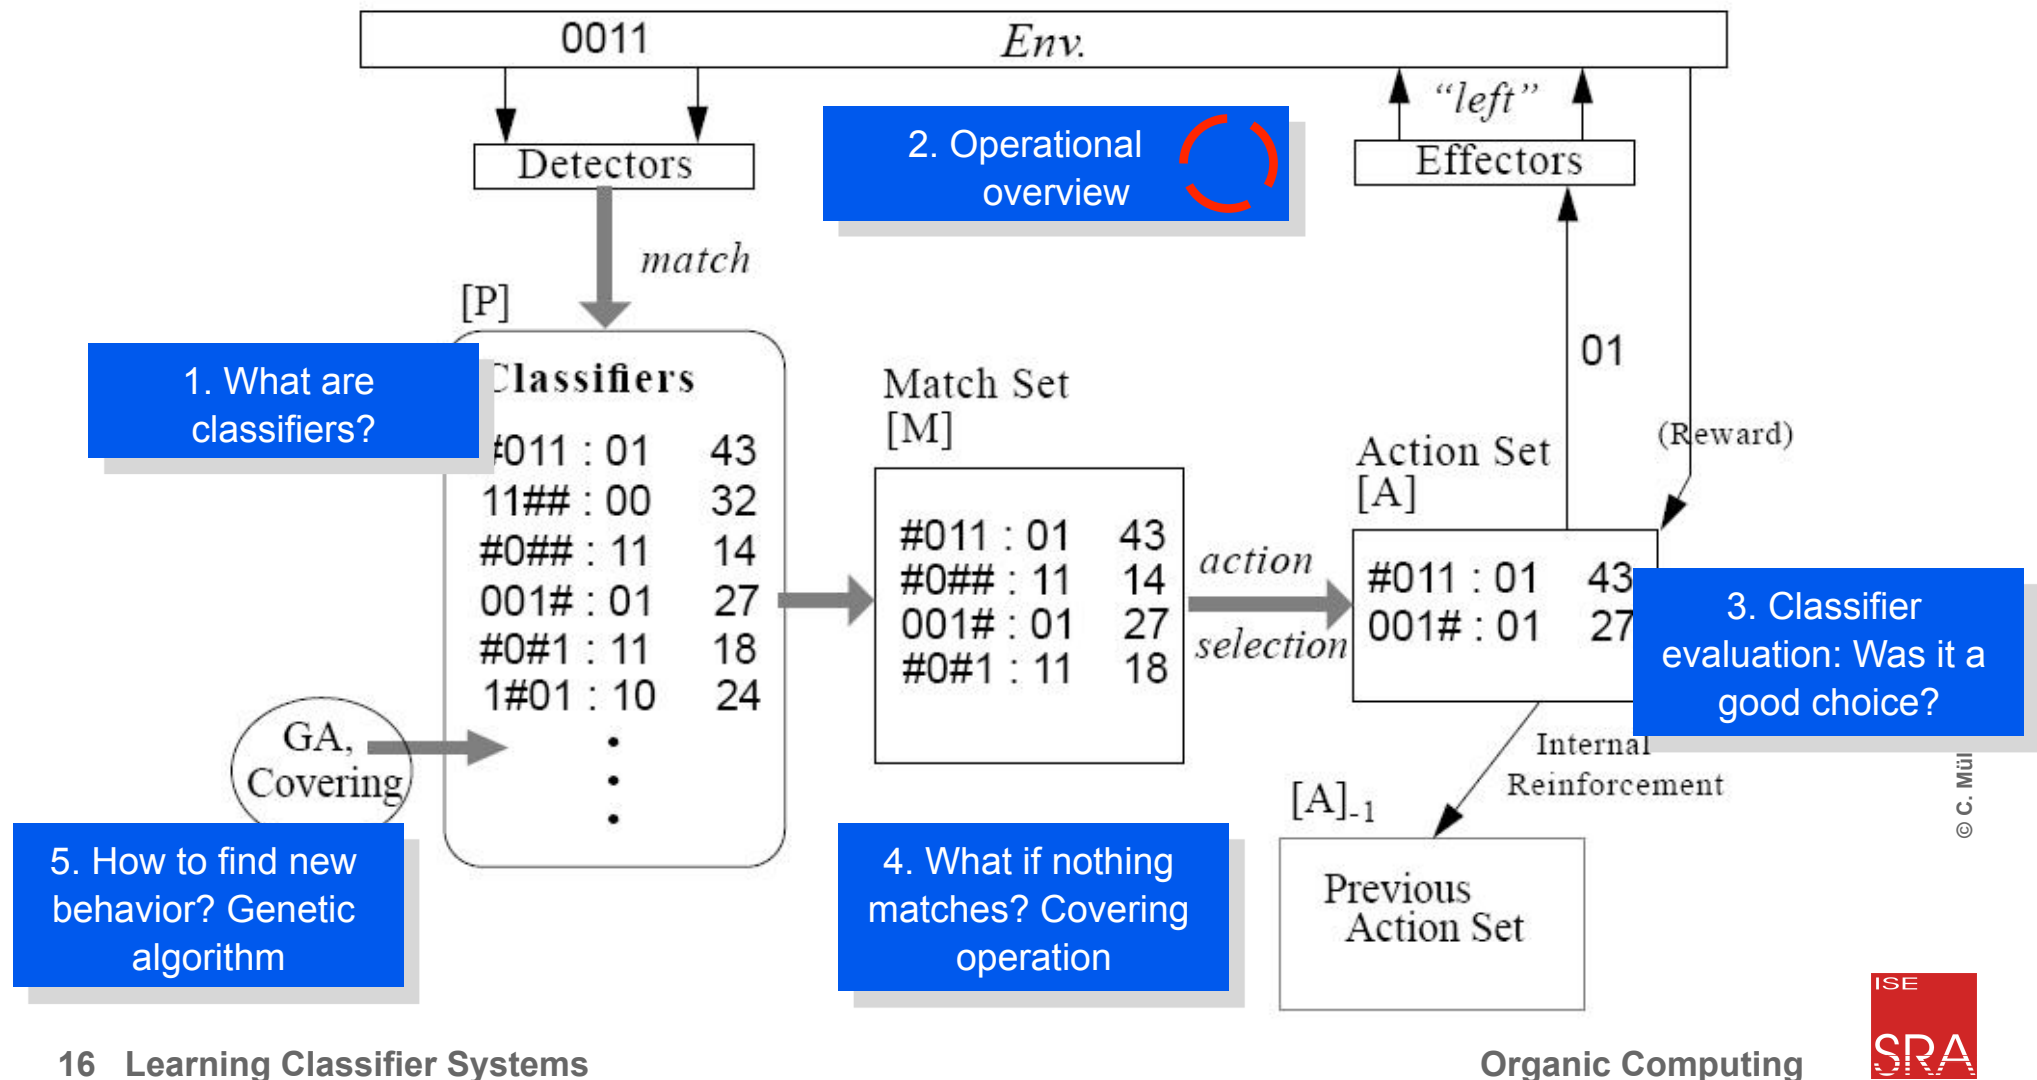
- ❑ Initially, Holland designed a system called CS1 (published in 1978).
- ❑ System contains
 - Set of classifiers (condition/action) pairs
 - Input interface to receive state from the environment
 - Output interface to apply actions to the environment
 - Internal message list as an internal “workspace” for I/O
 - Evolutionary process (genetic algorithm) to generate new classifiers



- ❑ ZCS was introduced by Wilson in 1994.
- ❑ Simplified in comparison to the original LCS proposed by Holland.
- ❑ The message list in Holland's approach was rather complicated to use.
- ❑ Also a mechanism called "rule bidding" was removed.



In the following we will take a closer look at the different components.



1. What are classifiers?

```
if condition then action : fitness
```

- Initially a **classifier is a triple** (Condition, Action, *fitness*).
- **Condition**: Represents a state of the environment
- **Action**: Action for effecting the environment
- **Fitness (= strength)**: to “the best of the agent’s knowledge”, applying the classifier’s action under the given condition will return the expected reward.
- To **match a classifier**: Test condition against environmental state.
- Condition may contain **wildcards** (typically abbreviated by #).
- Initially, only **binary representation** was used.

1. What are classifiers?

```
if condition then action : fitness
```

□ Example:

- (011001, 0, 123): exact match required
- (01#001, 1, 70): matches 010001 and 011001

□ Classifiers have been extended to deal with other representations.

□ For example, real values:

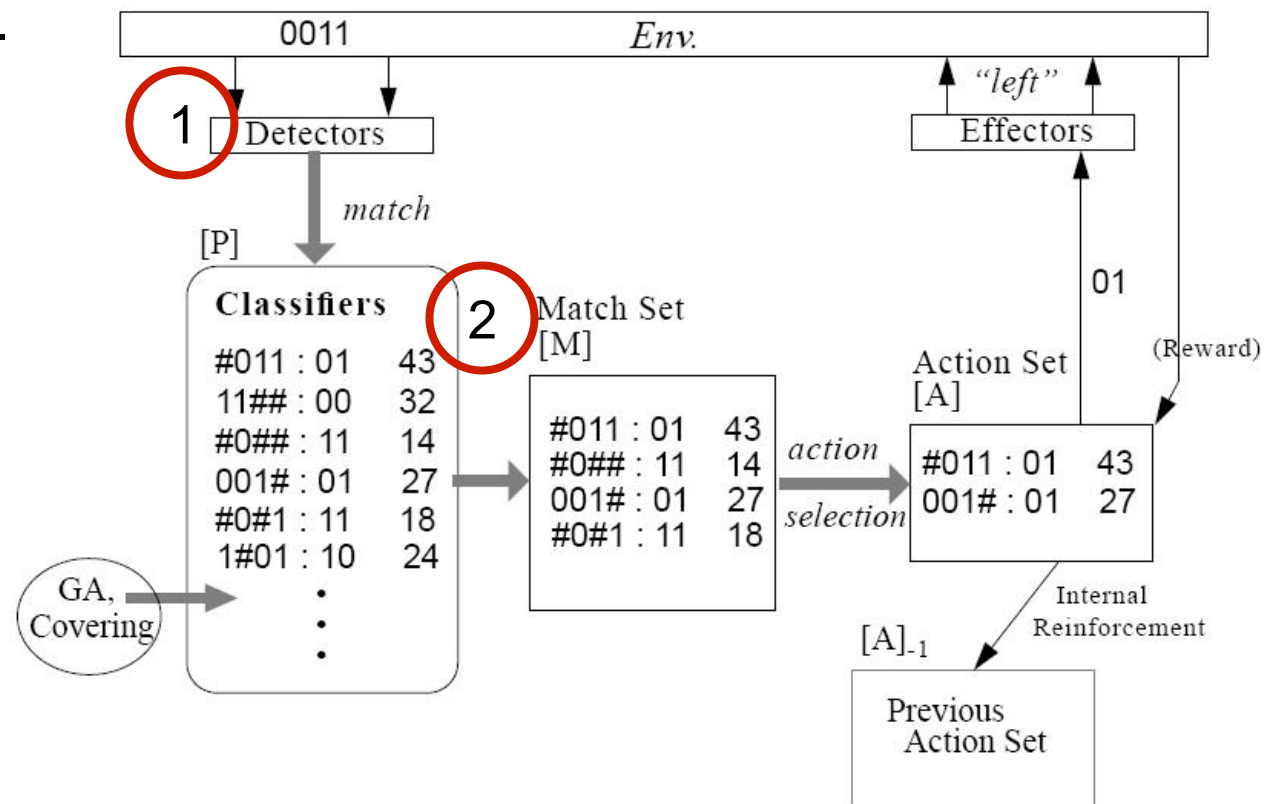
- Use of “interval predicates” instead of binary variables
- Input attribute x_i matches, if $i_i \leq x_i \leq u_i$
- Concatenation of (i_i, u_i) pairs for the i -th entry of the condition

□ Different representation impacts operations of the genetic algorithm used (in a minute).

2. Operational overview

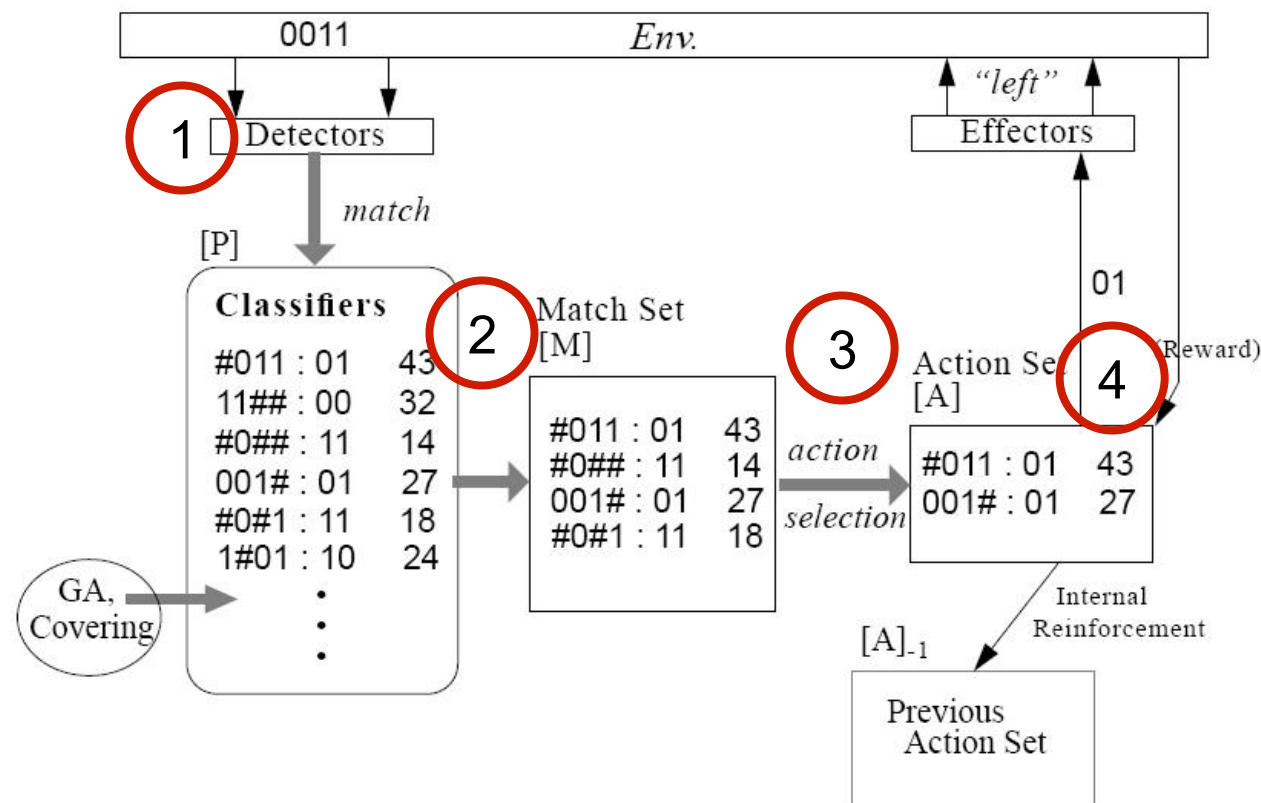
□ First glance: What is happening at runtime?

1. Agent periodically receives information about environment **through detectors**.
2. State is used to calculate the **match set**: Select all classifiers that match the current situation.



2. Operational overview

3. Action selection through roulette-wheel selection according to fitness (high fitness \rightarrow high probability of being selected).
4. All classifiers that propose the same action as the selected one are put into the action set.



3. Classifier evaluation: Was it a good choice?

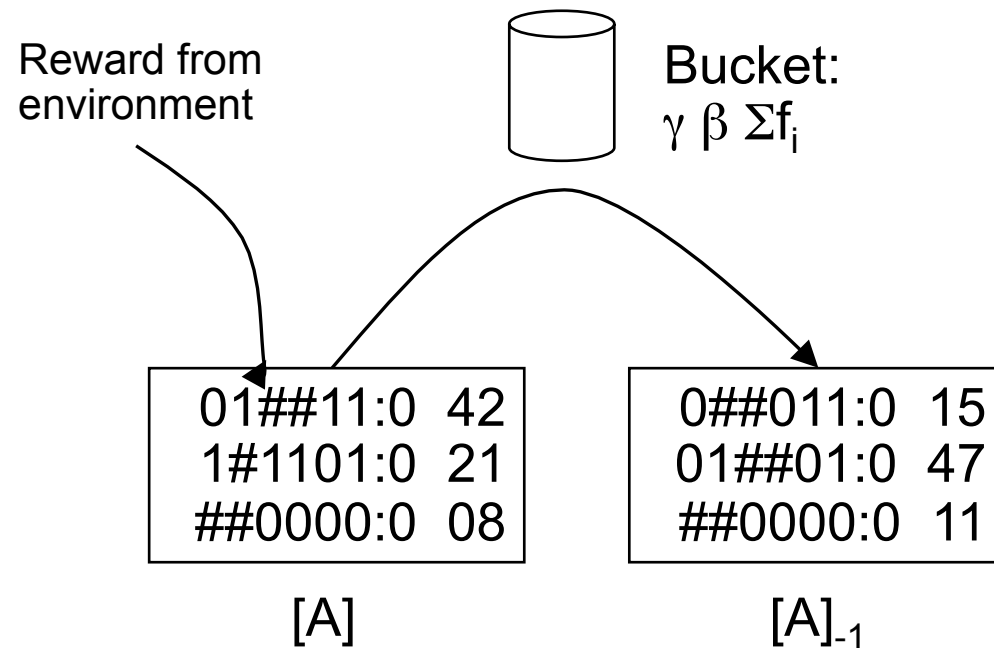
- ❑ As the system is supposed to learn, the usefulness of the classifiers has to be “monitored”: Idea is to **change the fitness according to performance**.
- ❑ **Fitness modification** in current and previous action set $[A]$ and $[A]_{-1}$ as well as $[M]$.
- ❑ Algorithm **parameter**: $\beta \in (0, 1]$
- ❑ Fixed fraction β of the fitness of all classifiers in $[A]$ is put into a **common “bucket”**.
- ❑ cont. ...

3. Classifier evaluation: Was it a good choice?

- ☐ Content of bucket is **discounted by $\gamma \in (0, 1]$** .
- ☐ Bucket content is **distributed evenly among classifiers in $[A]_{-1}$** .
- ☐ **Reward paid by the environment**: Fraction β of reward distributed evenly among members of $[A]$.
- ☐ Remaining members in $[M]$ (and not in $[A]$) are **“taxed” with τ** .


3. Classifier evaluation: Was it a good choice?

- ❑ Reward received for current action is paid to action set.
- ❑ But: Some fitness is passed on to previous action set.
- ❑ Bucket brigade metaphor




3. Classifier evaluation: Was it a good choice?

- ❑ Reward received for current action is paid to action set.
- ❑ But: Some fitness is passed on to previous action set.
- ❑ Bucket brigade metaphor



3. Classifier evaluation: Was it a good choice?

- ❑ Reward received for current action is paid to action set.
- ❑ But: Some fitness is passed on to previous action set.
- ❑ Bucket brigade metaphor



© C. Müller-Schloer 2015

4. What if nothing matches?
Covering operation

□ When input from environment is received:

- There may be **no matching classifier**, or
- the **match set is “poor”**, i.e., combined fitness is less than Φ times the population's average fitness.


□ **Covering** mechanism

- Used to create a **new rule matching the current situation**
- **Action** of new rule is **chosen randomly**.
- **New rule is generalized**: With probability $\frac{1}{3}$ each bit of the condition is replaced by # (wildcard).
- Fitness is initialized with population average.

5. How to find new behavior? Genetic algorithm

- ❑ Inspired by Darwinian evolution: **Survival of the fittest**
- ❑ On each invocation **two classifiers are selected as parents**: roulette-wheel based on fitness
- ❑ With fixed probabilities, these “parents” are subject to **“genetic” operations**
 - Crossover (recombination)
 - Mutation
- ❑ Resulting two classifiers get half of their parents’ fitness (which is deducted from them).
- ❑ **Rate at which GA is invoked** is application dependent.
 - Too frequent: noisy fitness
 - Too seldom: slow development
 - Classifiers should have been evaluated a “couple of times” before becoming a parent.


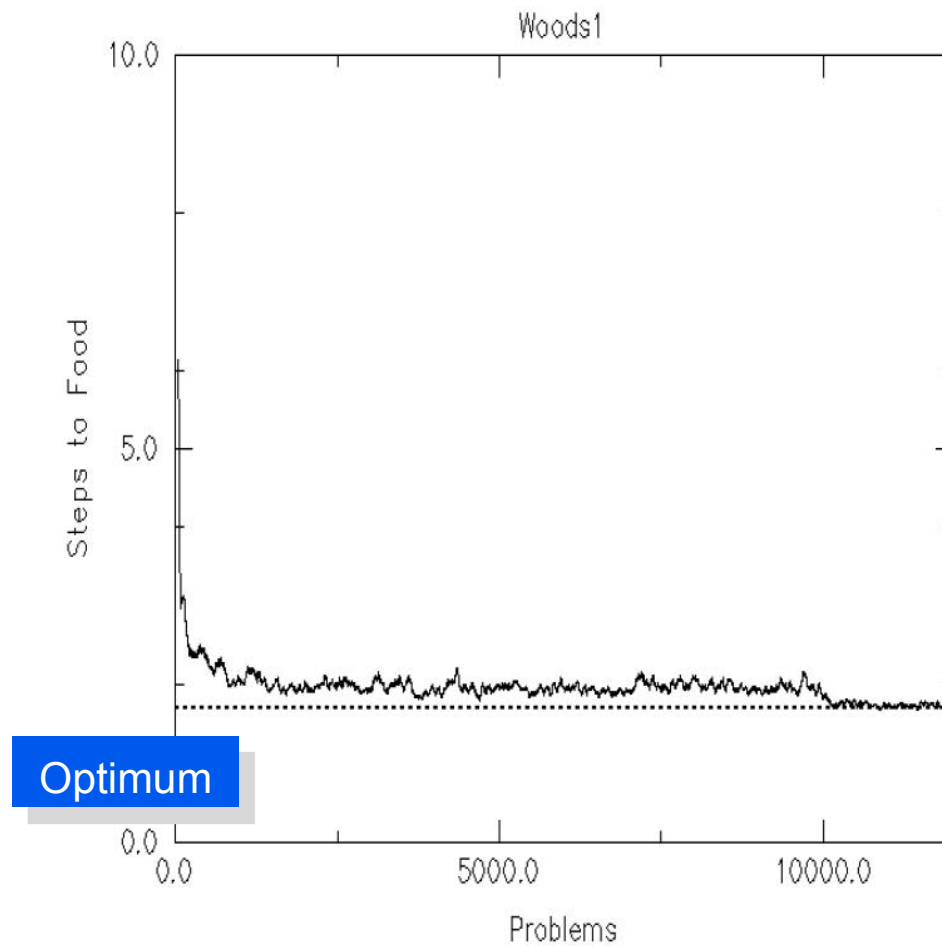
5. How to find new behavior? Genetic algorithm



- ☐ One-point crossover in the example: 3/5 bits
- ☐ Multi-point crossover: More than one “cut point”
- ☐ Mutation: Randomly flips a bit (typically very small probability).

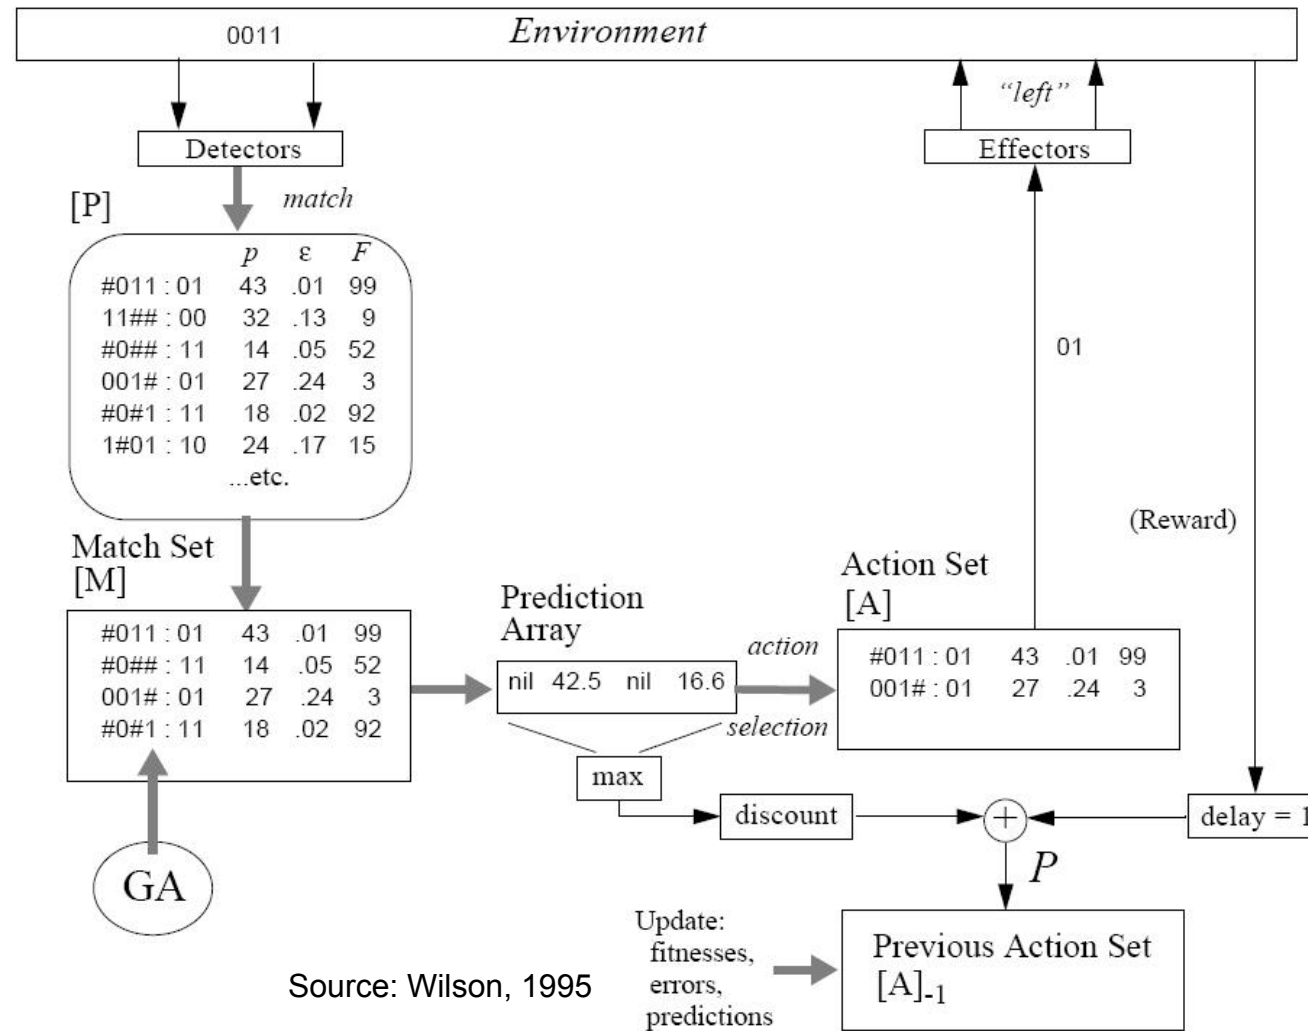
□ Typical parameter values according to Wilson:

- Rule-base (population, [P]) contains 400 classifiers
- Initial fitness of classifiers set to 20
- Discount factor $\gamma = 0.71$
- Tax $\tau = 0.1$
- Covering trigger $\Phi = 0.5$
- GA rate per time step $p = 0.25$
- Crossover rate = 0.5
- Per bit mutation rate $\mu = 0.002$



© C. Müller-Schloer 2015

- ❑ In ZCS (and other LCSs): Fitness of a classifier is measured by its expected reward.
- ❑ Wilson suggested that in some situations a classifier with low pay-off may be best suited.
- ❑ XCS introduces **separate attributes** for expected reward and fitness.
- ❑ In some situations the **best reward may be lower than in other situations**.
- ❑ Fitness of a classifier is consequently measured by the accuracy of the reward *prediction*.
- ❑ Coined the term **accuracy-based** classifiers in contrast to **strength-based (fitness-based)**.



- ❑ Classifiers are extended to a 5-tuple
(Condition, Action, Prediction (p), Prediction error (ε), Fitness (F))
- ❑ Condition, Action: same as in ZCS
- ❑ **Prediction**: A measure for the pay-off received when this classifier's action controlled the system
- ❑ **Prediction error**: A measure for the classifiers miss-prediction
- ❑ **Fitness**: A function of the inverse prediction error
- ❑ **Other parameters** can be used depending on the implementation, e.g., the number of times a classifier has been used (**experience**).
- ❑ Prediction and prediction error are used in action selection.
- ❑ Fitness used by the genetic algorithm (and also in action selection)

- ❑ Match set calculated as before
- ❑ Many policies may be used for action selection, e.g.
 - Deterministic selection: **Highest prediction wins.**
 - Probabilistically based on **fitness-weighted average of predictions** suggesting an action (contained in the architecture's **prediction array**)
- ❑ Selected action is sent to the effectors.
- ❑ Reward may (or may not) be received in return.
- ❑ Classifiers are updated (differently for single- and multi-step problems).

- Reinforcement consists of **updating prediction, prediction error, and fitness**.
- Let's consider the (simpler) update process in single-step problems:
- Update **prediction** for each classifier C_j in $[A]$
 - $p_j \leftarrow p_j + \beta (P - p_j)$
 - P is the current reward.
 - β is called the learning rate.
- Update **prediction error** accordingly
 - $\varepsilon_j \leftarrow \varepsilon_j + \beta (|P - p_j| - \varepsilon_j)$
- Update **fitness** accordingly:
 - $F_j \leftarrow F_j + \beta (\kappa_j' - F_j)$
 - κ_j' is the **relative accuracy** across $[A]$, $\kappa_j \equiv 1 / \varepsilon_j$

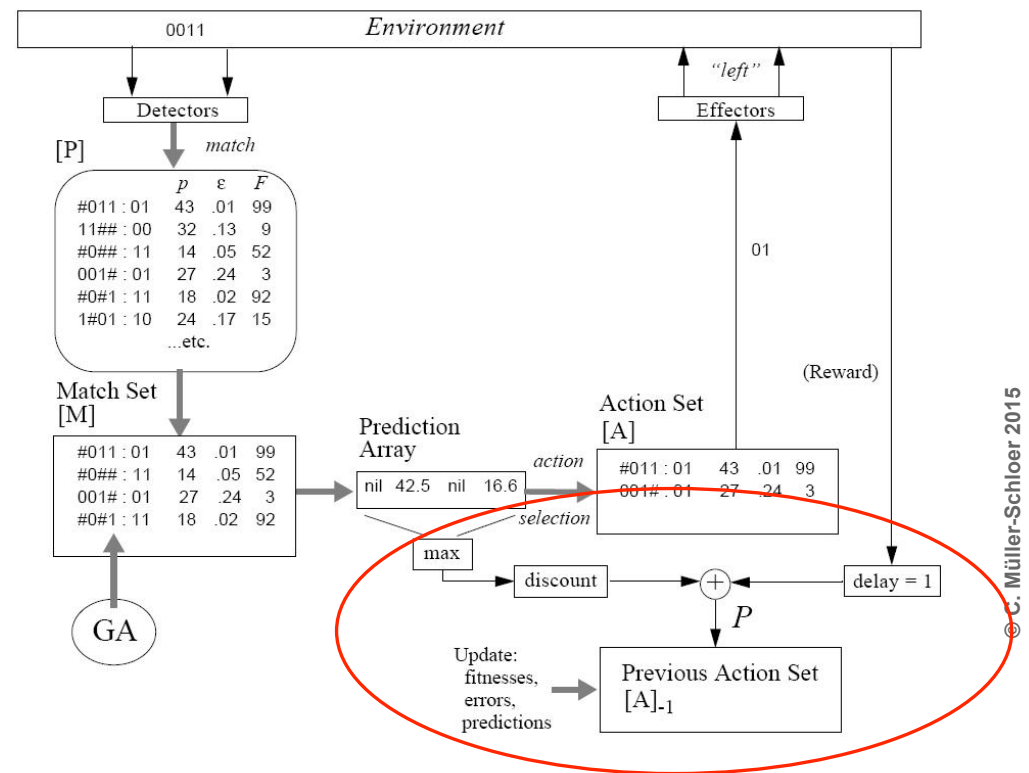
low prediction error \rightarrow high fitness!

- ❑ For multi-step problems updating is done based on $[A]_{-1}$.
- ❑ Delayed update allows to retrieve “information from the future”.
- ❑ Inspired by **Q-Learning** (reinforcement learning technique)
- ❑ Here reward P is calculated differently:


$$P = r + \gamma * \max P(a)$$

Reward
from t-1

Prediction
array of t




- Experiments done on the Woods2 scenario (complicated Woods1)



Source:
Wilson 1995

© C. Müller-Schloer 2015

- Further improvements of XCS suggested by Butz, Goldberg and Lanzi, (2003) show great improvements for multi-step problems





- ❑ Machine learning techniques are applied in OC application.
- ❑ One example is the Organic Traffic Control (OTC) project.
- ❑ Overall goal: **Online adaptation of traffic light controllers** to changing traffic situations
- ❑ Requirements
 - **Adapt autonomously** to the environment
 - **Long term** changes
 - **Short term** fluctuations, incidents
 - **Re-use knowledge**
 - **Safety**: Limit effects of possible errors of learning component!
 - **Comprehensible behavior**
 - Limit necessary manual intervention and effort for setup!



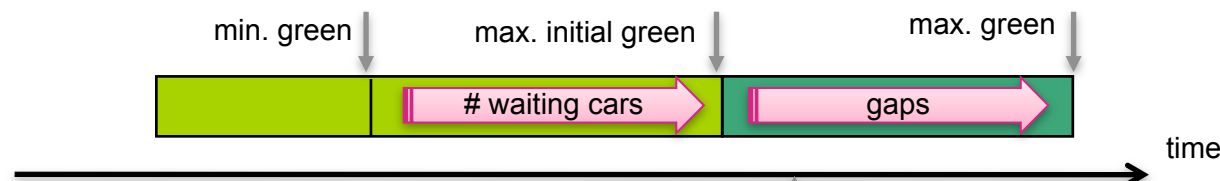
- **Input:** Traffic situation, vehicles per hour (flow) per relation in the junction (turning)
- **Output:** Parameter set modifying the program for traffic light controller (fixed time or traffic responsive control)
- **Objective function:** Level of Service (LoS, average delay time per vehicle); used in Germany (HBS) and the US (HCM)
 - Goal: **Minimize** LoS value
 - Best: LoS = 0 (no delay for anyone)

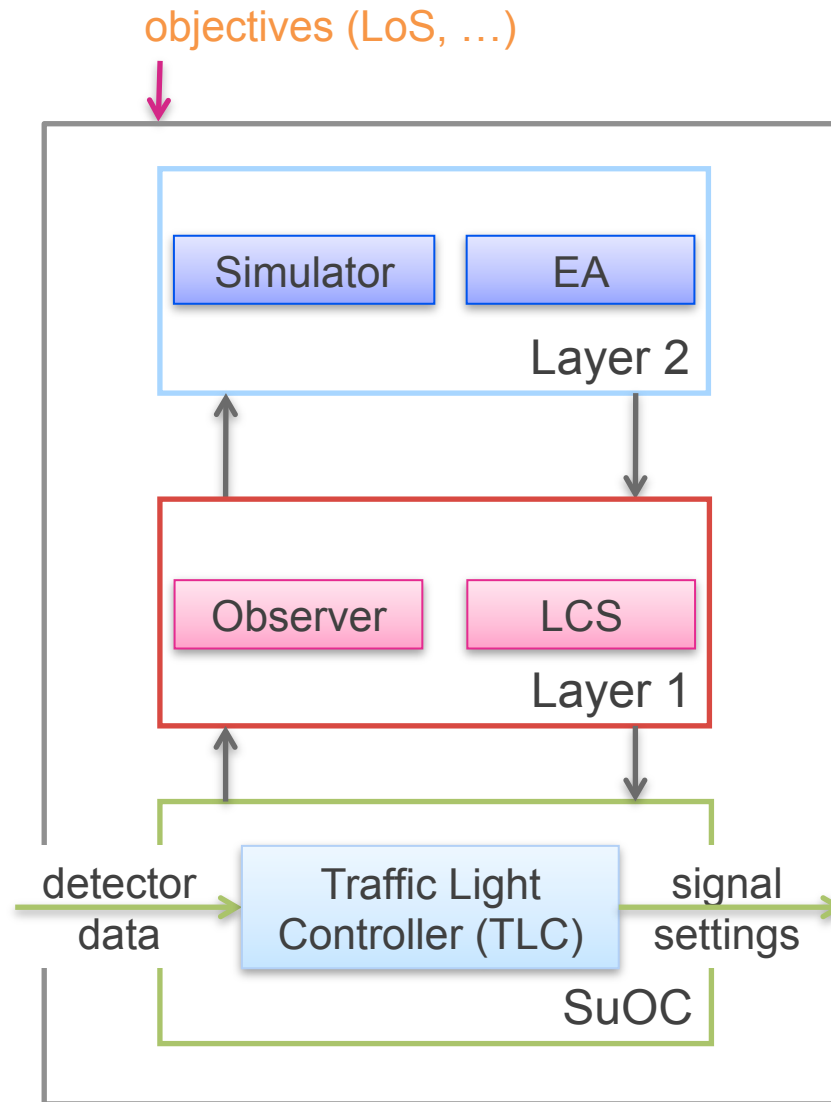
$$LoS = \frac{\sum flow \cdot t_{delay}}{\sum flow}$$

Flows weighted
with delay

Each flow:
vehicles/hour

- ❑ Example: NEMA Controller (traffic responsive controller)
- ❑ Cars in each lane are counted by **detectors**.
- ❑ Detector data represents (typically incomplete) **traffic situation**.
- ❑ Once set to green, it stays there until at **least “min. green”**.
- ❑ Extension up to “max. initial green” **depending on number of waiting cars**
- ❑ After that: When **gap between incoming cars becomes large**: leave green
- ❑ Complete parameter set: 9 parameters
- ❑ **Goal of OTC**: LCS selects TLC controller settings based on current traffic situation.





User interface

- User defines system objectives

Layer 2

- Extend behavioral repertoire of Layer 1
- Offline learning (TLC parameters)

Layer 1

- Adapt SuOC parameters
- Online learning (rule quality)


System under Observation and Control

- Control traffic signals

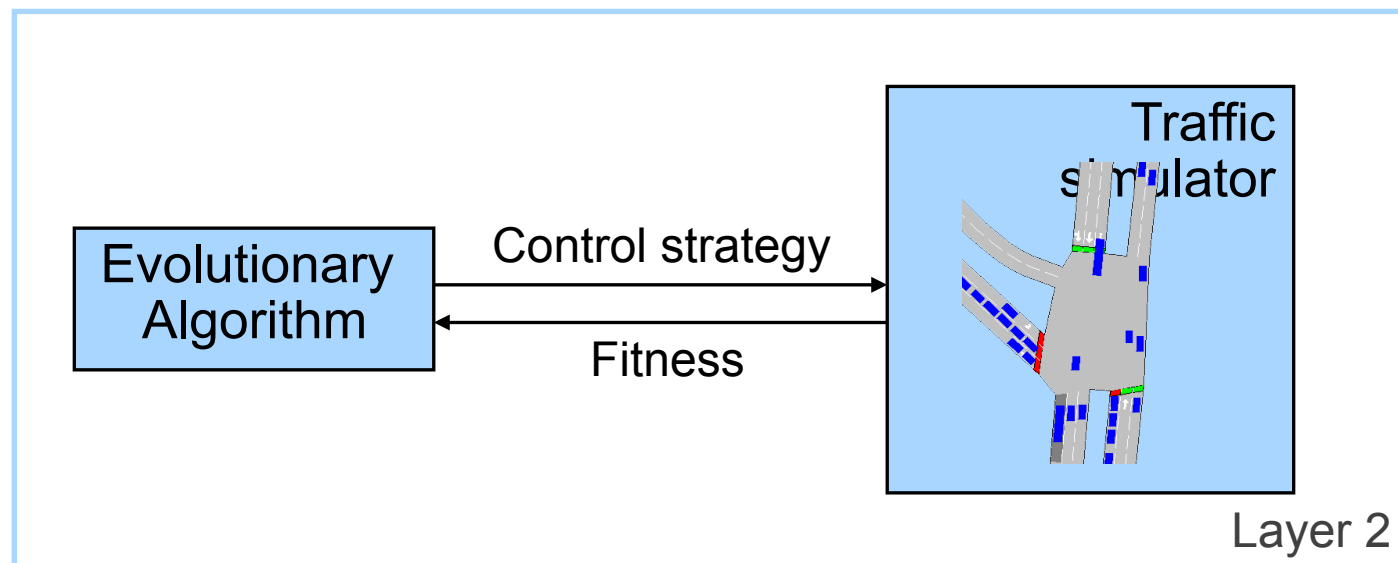
- ❑ Exploring the classifier space **online** using GA **can be dangerous!**
 - System could try the rule: Set all traffic lights to green!
- ❑ Consequence: Generation of new rules done in a separate (isolated) component: “Sandbox”
 - **Learning is done offline** in a simulator (layer 2).
 - LCS updates classifiers, performs covering, applies changes to TLC (layer 1).
- ❑ In classifiers: Representation of input as **real-valued intervals**
- ❑ Building of match set: Trade-off between “use only matching solutions” and creative competition needed for learning

- ❑ In the original XCS: Covering creates new classifier for the current situation randomly.
- ❑ OTC: Application specific **widening of existing classifiers**
- ❑ Select “closest” rule, copy, widen condition
- ❑ Trade-off between “use only tested solutions” and quick reaction time
- ❑ Additionally: Threshold used to **trigger layer 2 GA**



Modified Covering



- ❑ Generates control strategies for traffic signals (layer 2)
- ❑ Quality of strategy is **tested using a traffic simulator**.
- ❑ Fitness metrics: e.g. mean travel time, average number of stops, ...



- ❑ Strategies are evaluated under different traffic conditions.
- ❑ Optimized if-then-rule (condition + strategy) is added to rule set of LCS.



- ❑ Machine learning techniques show promising results for self-adaptation in organic computing systems.
- ❑ Learning classifier systems are a good example of such techniques.
- ❑ The OTC project makes heavy use of LCSs in a practical technical system.
- ❑ Modifications have been proposed by the project
 - Increased safety through offline learning (layer 2)
 - Modified operations adapted to the specific problem

- ❑ Oliver Sigaud, Stewart W. Wilson: Learning classifier systems: a survey. Soft Computing, pp. 1065-1078, 2007
- ❑ John H. Holland: Adaptation in natural and artificial systems. MIT Press, 1992 (contains the text of the 1975 edition with some extensions)
- ❑ Stewart W. Wilson: ZCS: A zeroth-level classifier system. Evolutionary Computing Vol. 2, No. 1, 1994
- ❑ Larry Bull, Jacob Hurst: ZCS Redux. ???
- ❑ Stewart W. Wilson: Classifier Fitness Based on Accuracy. Evolutionary Computing Vol. 3, No. 2, 1995 (original XCS reference)
- ❑ Butz, Goldberg, Lanzi: Gradient Decent Methods in Learning Classifier Systems: Improving XCS Performance in Multistep Problems. IlliGAL Report No. 2003028, 2003