

Teil I: Temporale DBS

1. Einführung

Gegenstand:

in DBMS und DB-Sprache eingebaute Unterstützung
der Speicherung und Anfragebearbeitung
von zeitbehafteten Daten aus Vergangenheit, Gegenwart,
evtl. Zukunft

Applications

- Academic

Transcripts record courses taken in previous and the current semester or term and grades for previous courses

- Accounting

What bills were sent out and when, what payments were received and when?

Delinquent accounts, cash flow over time

Money-management software such as Quicken® can show e.g., account balance over time.

- Budgets

Previous and projected budgets, multi-quarter or multi-year budgets

Applications, cont.

- Data Warehousing

Historical trend analysis for decision support

- Financial

Stock market data

Audit analysis: why were financial decisions made, and with what information available?

Program trading

- Geographic Information Systems (GIS)

Land use over time: boundary of parcels change over time, as parcels get partitioned and merged.

Title searches

Applications, cont.

- Insurance

Which policy was in effect at each point in time, and what time periods did that policy cover?

- Inventory management

Inventory over time, for time-series analysis, accounting

- Legal records

Validity periods for laws

- Medical records

Patient records, drug regimes, lab tests

Tracking course of disease

Longitudinal studies

Applications, cont.

- Payroll

Past employees, employee salary history,
salaries for future months, records of
with-holding requested by employees

- Planning

Use current and previous schedules for
designing new schedules

Planning for future contingencies

Network management

- Process monitoring

Chemical, nuclear power plant
monitoring

Active databases

Applications, cont.

- Project scheduling

Milestones, task assignments

- Reservation systems

E.g., airlines, trains, hotels

Configuring new routes and offers
to ensure high utilization

- Scientific applications

Recording physical experiments

Timestamping weather satellite images

Dating archeological finds

Applications: Conclusion

- It is difficult to identify applications that do *not* involve time aspects and the management of temporal data.
- These applications would benefit from built-in temporal support in the DBMS.

More efficient application development

Potential increase in performance

An SQL Case Study

- University of Arizona's Office of Appointed Personnel has some information in a database.

Schema: Employee(Name, Salary, Title)

- Finding an employee's salary is easy.

```
SELECT Salary  
FROM Employee  
WHERE Name = 'Bob'
```

- The OAP wishes to add the date of birth

Employee(Name, Salary, Title, DateofBirth **DATE**)

An SQL Case Study, cont.

- Finding the employee's date of birth is equally easy.

```
SELECT DateofBirth  
FROM   Employee  
WHERE Name = 'Bob'
```

- Now the OAP wishes to computerize the employment history.

Employee (Name, Salary, Title, DateofBirth,
Start **DATE**, Stop **DATE**)

Store for each tuple the time interval [..) when it was valid.

An SQL Case Study, cont.

- To the data model, these new columns are structurally identical to DateofBirth.

<i>Name</i>	<i>Salary</i>	<i>Title</i>	<i>DateofBirth</i>	<i>Start</i>	<i>Stop</i>
Bob	60000	Assistant Provost	1945-04-19	1993-01-01	1993-06-01
Bob	70000	Assistant Provost	1945-04-19	1993-06-01	1993-10-01
Bob	70000	Provost	1945-04-19	1993-10-01	1994-02-01
Bob	70000	Professor	1945-04-19	1994-02-01	1995-01-01

- To find the employee's current salary, things are more difficult.

```
SELECT Salary
FROM Employee
WHERE Name = 'Bob'
      AND Start <= CURRENT_TIMESTAMP
      AND CURRENT_TIMESTAMP < Stop
```

An SQL Case Study, cont.

- OAP wants to distribute to all employees their salary history.

Output: For each person, maximal intervals for each salary are expected.

<i>Name</i>	<i>Salary</i>	<i>Start</i>	<i>Stop</i>
Bob	60000	1993-01-01	1993-06-01
Bob	70000	1993-06-01	1995-01-01

But: Employee could have arbitrarily many title changes between salary changes.

Thus: Consecutive value-equivalent tuples have to be **coalesced**.

An SQL Case Study, cont.


- Alternative 1

Give the user a printout of Salary and Title information, and have user determine when his/her salary changed.

- Alternative 2

Use SQL as much as possible.

Find those intervals that overlap or are adjacent and thus should be merged.



SQL Program

```
CREATE TABLE Temp(Salary, Start, Stop)
AS SELECT Salary, Start, Stop
FROM Employee
WHERE Name = 'Bob';
```


repeat

```
    UPDATE Temp AS T1
    SET (T1.Stop) = (SELECT MAX(T2.Stop)
                    FROM Temp AS T2
                    WHERE T1.Salary = T2.Salary
                    AND T1.Start < T2.Start
                    AND T1.Stop >= T2.Start
                    AND T1.Stop < T2.Stop)
    WHERE EXISTS (SELECT *
                  FROM Temp AS T2
                  WHERE T1.Salary = T2.Salary
                  AND T1.Start < T2.Start
                  AND T1.Stop >= T2.Start
                  AND T1.Stop < T2.Stop)
```


until *no tuples updated;*

Example


- Initial table



- After first iteration:



- After second iteration:



SQL Program, cont.

- The loop is executed $\log_2 N$ times in the worst case, where N is the number of tuples in a chain of overlapping or adjacent, value-equivalent tuples.
- Then delete extraneous, non-maximal intervals.

```
DELETE FROM Temp T1
WHERE EXISTS (SELECT *
               FROM Temp AS T2
               WHERE T1.Salary = T2.Salary
                  AND ((T1.Start > T2.Start
                      AND T1.Stop <= T2.Stop)
                  OR (T1.Start >= T2.Start
                      AND T1.Stop < T2.Stop))
```

⁰Die SQL-Anweisung ist allgemeiner als hier nötig angelegt.

- Alternative 3: Use pure SQL.

```
CREATE TABLE Temp(Salary, Start, Stop)
AS SELECT Salary, Start, Stop
FROM Employee
WHERE Name = 'Bob';
```

```
SELECT DISTINCT F.Salary, F.Start, L.Stop  // teste First/Last-Paar
FROM Temp AS F, Temp AS L
WHERE F.Start < L.Stop
      AND F.Salary = L.Salary
      AND NOT EXISTS (SELECT *                               // keine Lücken?
                     FROM Temp AS M
                     WHERE M.Salary = F.Salary
                        AND F.Stop < M.Start
                        AND M.Start < L.Stop
                        AND NOT EXISTS (SELECT *
                                       FROM Temp AS T1
                                       WHERE T1.Salary = F.Salary
                                          AND T1.Start < M.Start
                                          AND M.Start <= T1.Stop))
      AND NOT EXISTS (SELECT *                               // maximal?
                     FROM Temp AS T2
                     WHERE T2.Salary = F.Salary
                        AND ((T2.Start < F.Start
                           AND F.Start <= T2.Stop)
                        OR (T2.Start <= L.Stop
                           AND L.Stop < T2.Stop)))
```


Using More Procedural Code

- Alternative 4: Use SQL only to open a cursor on the table.

Check subsequent intervals in an ordered list, whether they can be coalesced. Replace them by the new intervals.

```
DECLARE emp_cursor CURSOR FOR  
    SELECT Salary, Start, Stop FROM Employee  
    WHERE Name='Bob' ORDER BY Salary,Start;  
OPEN emp_cursor;
```

loop:

```
    FETCH emp_cursor INTO :salary,:start,:stop;  
    if no data returned then go to finished;  
    <compare start/stop-values with previous ones ...>  
    go to loop;
```

finished:

```
CLOSE emp_cursor;
```

How the query should look like

in a temporal query language:

```
TEMPORALLY SELECT Salary  
FROM Employee  
WHERE Name = 'Bob'
```

An SQL Case Study, cont.

- Drastic Alternative: Decompose the schema

Separate Salary and Title information;
store coalesced intervals

Employee1 (Name, Salary, Start
DATE, Stop **DATE**)

Employee2 (Name, Title, Start
DATE, Stop **DATE**)

- Getting the salary information is now easy:

```
SELECT Salary, Start, Stop  
FROM Employee1  
WHERE Name = 'Bob'
```

- But what if the OAP wants a table of (salary, title)-intervals?

An Example of a Temporal Join

- Employee1:

<i>Name</i>	<i>Salary</i>	<i>Start</i>	<i>Stop</i>
Bob	60000	1993-01-01	1993-06-01
Bob	70000	1993-06-01	1995-01-01

- Employee2:


<i>Name</i>	<i>Title</i>	<i>Start</i>	<i>Stop</i>
Bob	Assistant Provost	1993-01-01	1993-10-01
Bob	Provost	1993-10-01	1994-02-01
Bob	Full Professor	1994-02-01	1995-01-01

- Result of temporal join:


<i>Name</i>	<i>Salary</i>	<i>Title</i>	<i>Start</i>	<i>Stop</i>
Bob	60000	Assistant Provost	1993-01-01	1993-06-01
Bob	70000	Assistant Provost	1993-06-01	1993-10-01
Bob	70000	Provost	1993-10-01	1994-02-01
Bob	70000	Full Professor	1994-02-01	1995-01-01

Temporal Join in SQL


```
SELECT Employee1.Name, Salary, Title,  
         Employee1.Start, Employee1.Stop  
FROM Employee1, Employee2  
WHERE Employee1.Name = Employee2.Name  
      AND Employee2.Start <= Employee1.Start  
      AND Employee1.Stop <= Employee2.Stop
```




```
UNION  
SELECT Employee1.Name, Salary, Title,  
         Employee1.Start, Employee2.Stop  
FROM Employee1, Employee2  
WHERE Employee1.Name = Employee2.Name  
      AND Employee1.Start > Employee2.Start  
      AND Employee2.Stop < Employee1.Stop  
      AND Employee1.Start < Employee2.Stop
```



```
UNION  
SELECT Employee1.Name, Salary, Title  
         Employee2.Start, Employee1.Stop  
FROM Employee1, Employee2  
WHERE Employee1.Name = Employee2.Name  
      AND Employee2.Start > Employee1.Start  
      AND Employee1.Stop <= Employee2.Stop  
      AND Employee2.Start < Employee1.Stop
```



```
UNION  
SELECT Employee1.Name, Salary, Title  
         Employee2.Start, Employee2.Stop  
FROM Employee1, Employee2  
WHERE Employee1.Name = Employee2 Name  
      AND Employee2.Start => Employee1.Start  
      AND Employee2.Stop <= Employee1.Stop
```



How the query should look like

in a temporal query language:

TEMPORALLY SELECT

Employee1.Name, Salary, Title

FROM Employee1, Employee2

WHERE

Employee1.Name = Employee2.Name


Summary

- Applications manage temporal data.
- If a temporal DBMS is used
 - Schemas are simpler (temporal table)
 - SQL queries are simpler (temporally select)
 - Much less procedural code is needed
- Benefits
 - Application code is less complex
 - Easier to understand
 - Easier to ensure correctness
 - Easier to maintain
 - Performance may be enhanced by delegating functionality to the DBMS.

2. Zeitmodelle

Time Structure


- *Linear*



We will assume **linear** time.

Other Time Structures


- *Possible (Hypothetical) Futures*



- *Tree = Branching Time*
- *or Directed Acyclic Graph*

Other Time Structures, cont.

- *Periodic/Cyclic Time*



Time Density

- *Discrete*

- ◆ Time line is isomorphic to the integers.
- ◆ Time line composed of a sequence of pqp/decomposable time periods of uqo g fixed minimal duration, termed *chronons*.
- ◆ Between each pair of kpuwcpw? chronons is c hpk number of other kpuwcpw.

- *Dense*

- ◆ Time line is isomorphic to the rational numbers.
- ◆ Between gcej r ck qh kpuwcpw is an infinite number of other instants.

Time Density, cont.

- *Continuous*
 - ◆ Time line is isomorphic to the real numbers.
 - ◆ Between each pair of instants is an infinite number of other instants.
- Vgo r qtcnF DU work with **discrete** time.

Clocks

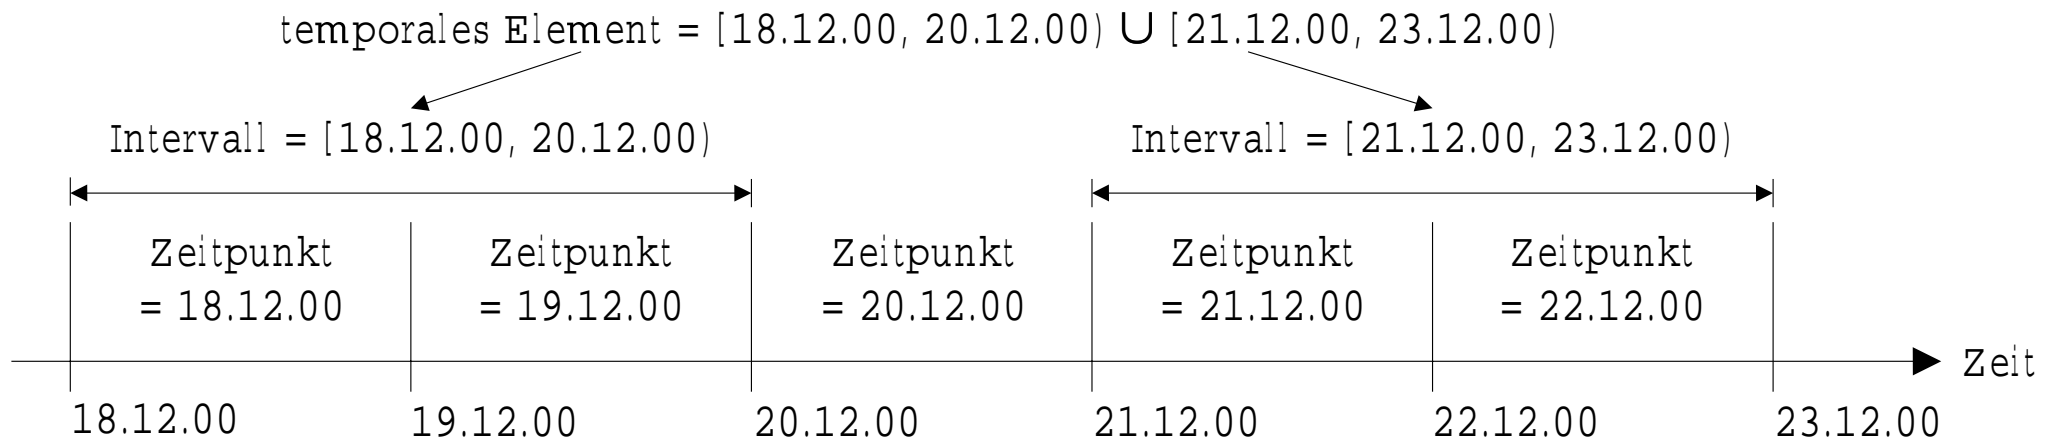

- A *clock* is a physical process coupled with a method of measuring that r t q e g u u 0
- The units of measurement are reported in terms of chronons of the clock.
- Their *granularity* may depend on application, e.g. seconds, days, weeks, months.
- Clocks (or *calendars*) provide the semantics of a timestamp representation.

Temporal Data Types

- A time *instant* t is a point on the time line. [dt. *Zeitpunkt*]
 - ◆ An *event* is an instantaneous fact, i.e., something occurring at an instant. The *event occurrence time* of an event is the instant at which the event occurs in the real world.
- An *instant set* is a set of instants.
- A *time period* $[t_1, t_2)$ is the time between two instants. [dt. *Intervall*]
 - ◆ Also called *interval*, but this differs from the SQL data type **INTERVAL**, e.g. "3 years".

Temporal Data Types, cont.

- A *time span* is a directed duration of time. A *duration* is an amount of time with a known length, but no specific starting or ending instants.
[dt. *Zeitspanne*]
 - ◆ A *positive span* denotes forward motion of time; a *negative span* denotes backwards motion of time.
- A *temporal element* is a finite union of intervals. [dt. *temporales Element*]
The component intervals are usually considered to be disjoint and non-adjacent.



Temporal Dimensions

- *Valid time* of a fact: when the fact is true in the modeled reality
[dt. *Gültigkeitszeit*]
- *Transaction time* of a fact: when the fact is current in the database and may be retrieved
[dt. *Transaktionszeit*]
- Four kinds of tables instead of one kind:
 - ◆ *snapshot*
 - ◆ *valid-time*
 - ◆ *transaction-time*
 - ◆ *bitemporal*

Example: Tom's Employment History

- On January 1, 1984, Tom joined the faculty as an Instructor.
- On December 1, 1984, Tom completed his doctorate, and so was promoted to Assistant Professor retroactively on July 1, 1984.
- On March 1, 1989, Tom was promoted to Associate Professor, proactively on July 1, 1989.

A Snapshot Table

- Can be modified
- Used for static queries, e.g.:

What is Tom's rank?

```
SELECT Rank  
FROM Faculty  
WHERE Name = 'Tom'
```

Analogy: Doorplate

(Jan. 1984)

Instructor Tom

(Dec. 1984)

Assistant
Prof. Tom


(March 1989)

Assistant
Prof. Tom

(July 1989)

Associate
Prof. Tom

A Transaction-time Table

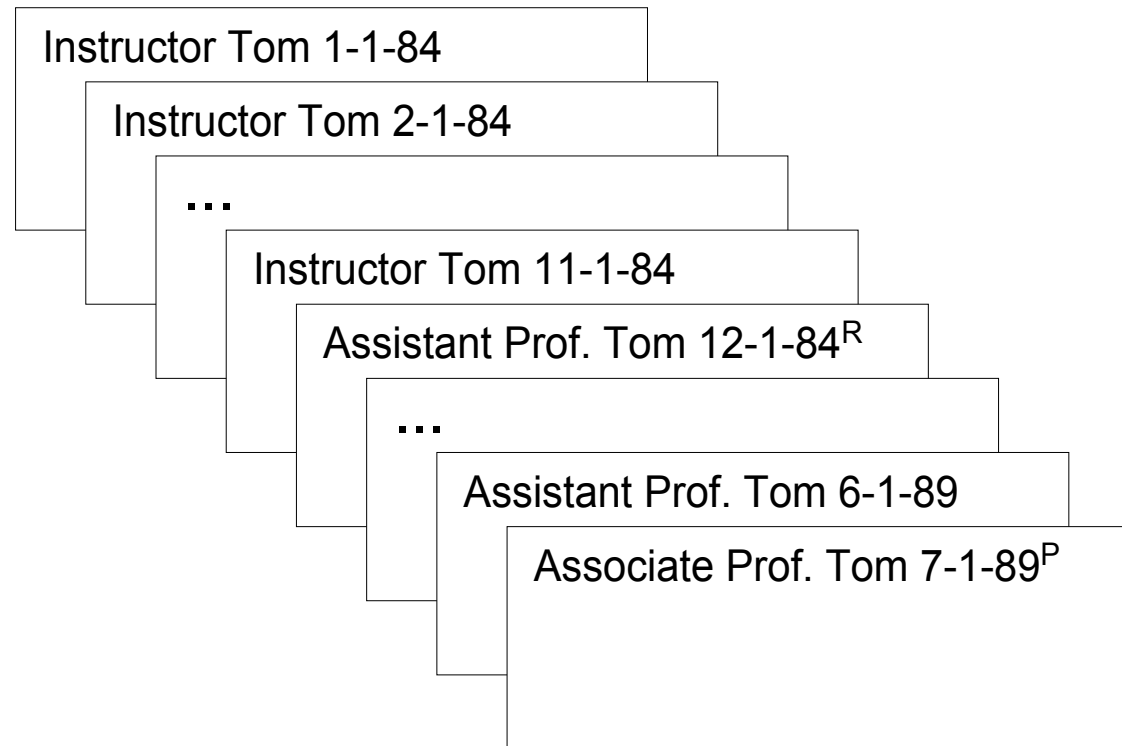


- Append-only: correction to previous snapshot states is not permitted
- Supports transaction time
- Supports rollback queries like:

What did we believe on Oct. 1, 1984
what Tom's rank was?

```
SELECT Rank  
FROM Faculty  
WHERE Name = 'Tom' AND  
      TRANSACTION(Faculty) OVERLAPS DATE '1984-10-01'
```


Analogy: Pay Stubs



R = incl. proof of retroactive payment for months 7...11-84

P = update event reserved externally on 3-1-89 (no future time handling!)

A Valid-Time Table



- May be modified
- Supports valid time
- Supports historical queries like:

What was Tom's rank on October 1, 1984 (as currently best known)?

```
SELECT Rank
FROM Faculty
WHERE Name = 'Tom'
AND VALID(Faculty) OVERLAPS DATE '1984-10-01'
```


Analogy: A Vita


Tom

Employment

Associate Professor	July, 1989
Assistant Professor	July, 1984
Instructor	January, 1984

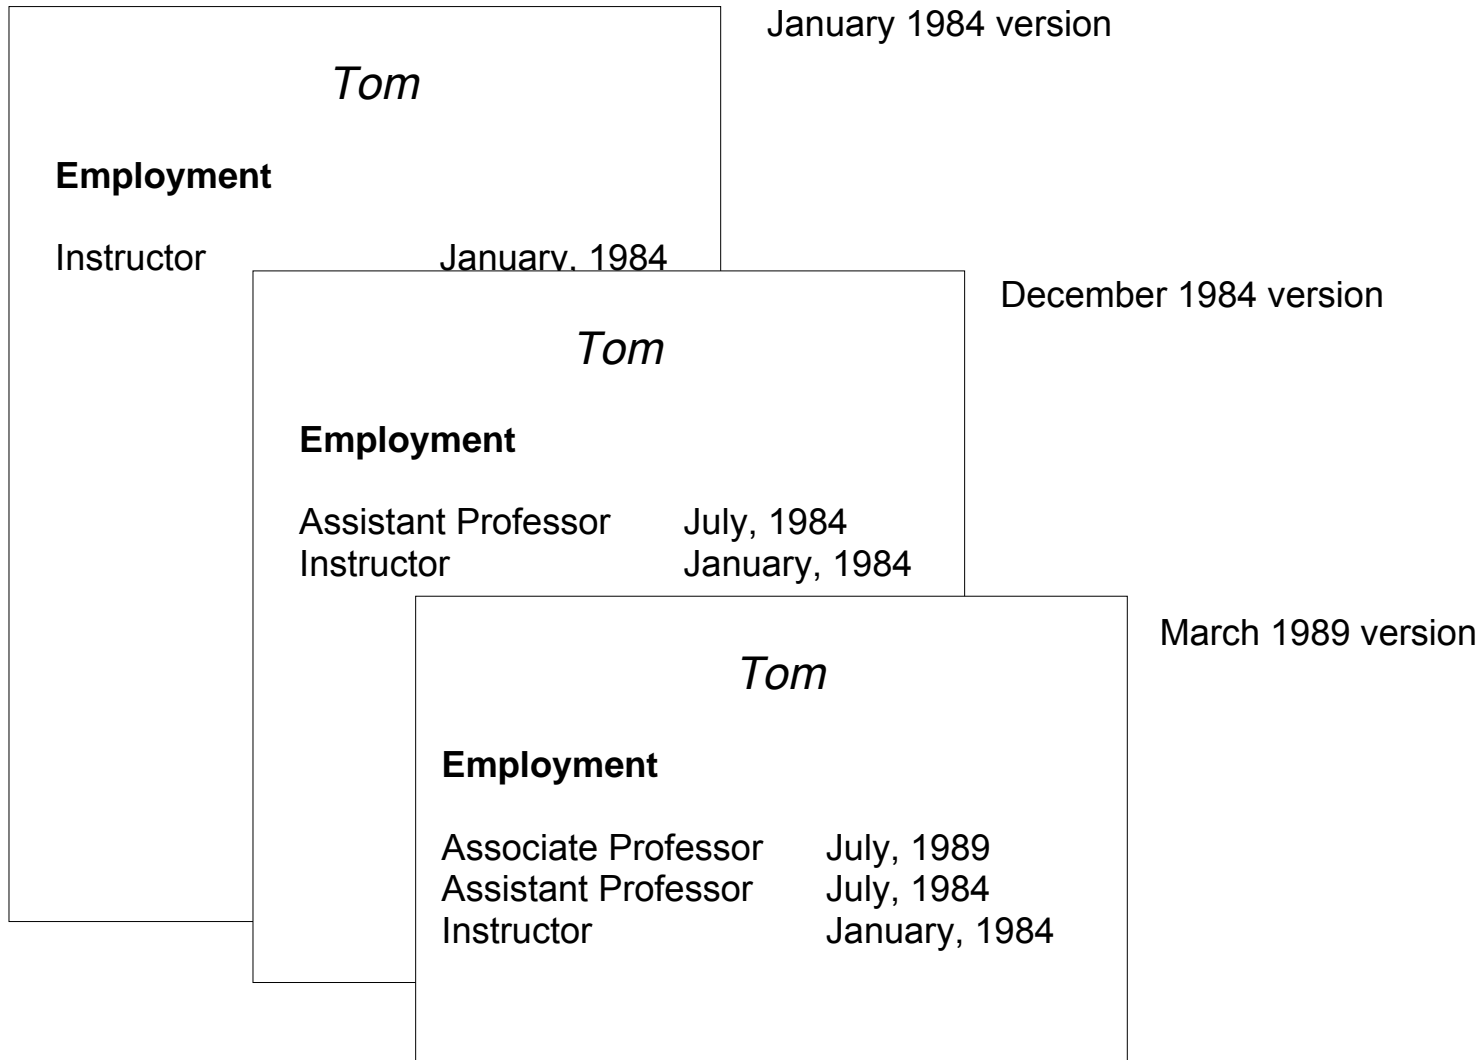
(as of March 1989 or later)

A Bitemporal Table



- Append-only
 - Supports valid and transaction time
 - Supports rollback coupled with historical queries like:
 - On October 1, 1984, what did we think Tom's rank was at that date?
- ```
SELECT Rank
FROM Faculty AS F
WHERE Name = 'Tom'
 AND VALID(F) OVERLAPS DATE '1984-10-01'
 AND TRANSACTION(F) OVERLAPS DATE '1984-10-01'
```

## Analogy: A Stack of Vitæ



Here, only valid/transaction start times are given; corresponding intervals are obvious.

## Summary

- Several different structures of time
  - ◆ Linear is simplest and most common.
- Five fundamental temporal data types
- Many different physical clocks/calendars
- Several dimensions of time
  - ◆ Valid time
  - ◆ Transaction time